

Databases 2

Reliability Control

Topics

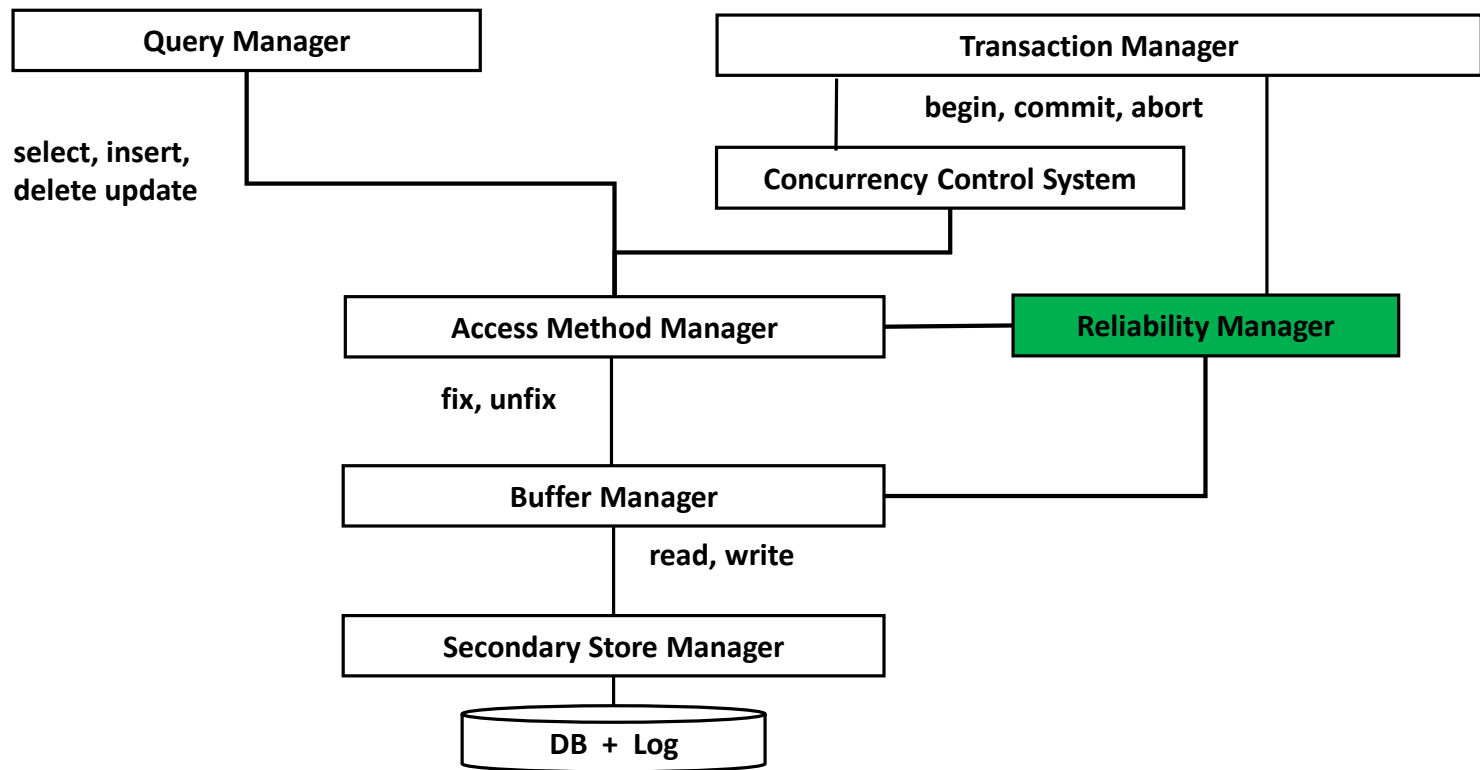
- Reliability definition and DBMS architecture
- Persistence of memory and backup
- Buffer management
- Reliable transaction management
- Log management
- Recovery after failures

Reliability

- *Reliability: The ability of an item to perform a required function under stated conditions for a stated period of time.*
[Reliability Definitions IEEE Dictionary of Electrical and Electronic Terms]
- In databases, reliability control ensures fundamental properties of transactions:
 - Atomicity: all or nothing semantics
 - Durability: persistence of effect
- DBMSs implement a specific architecture for reliability control
- The keys to database reliability are **stable memory** and **log management**

The Reliability Manager

- Realizes the transactional commands commit and abort
- Orchestrates read/write access to pages (both data and log)
- Handles recovery after failures




Persistence of Memory

- Durability implies a memory whose content lasts forever, which is an abstraction procedurally built on top of existing storage technology levels
- Main memory
 - Not persistent
- Mass memory
 - Persistent but can be damaged
- Stable memory
 - Cannot be damaged (clearly, this is an abstraction)
 - In practice stability = probability of failure close to 0
 - Keys to stability: replication and write protocols
- Failure of stable memory is the subject of the specific discipline of disaster recovery


How to Guarantee Stable Memory

- On-line replication: mirroring of two disks
- The popular RAID (Redundant Array of Independent Disks) disk architecture is treated in the COMPUTING INFRASTRUCTURES course



How to Guarantee Stable Memory

- Off-line replication: "tape" units (backup units)



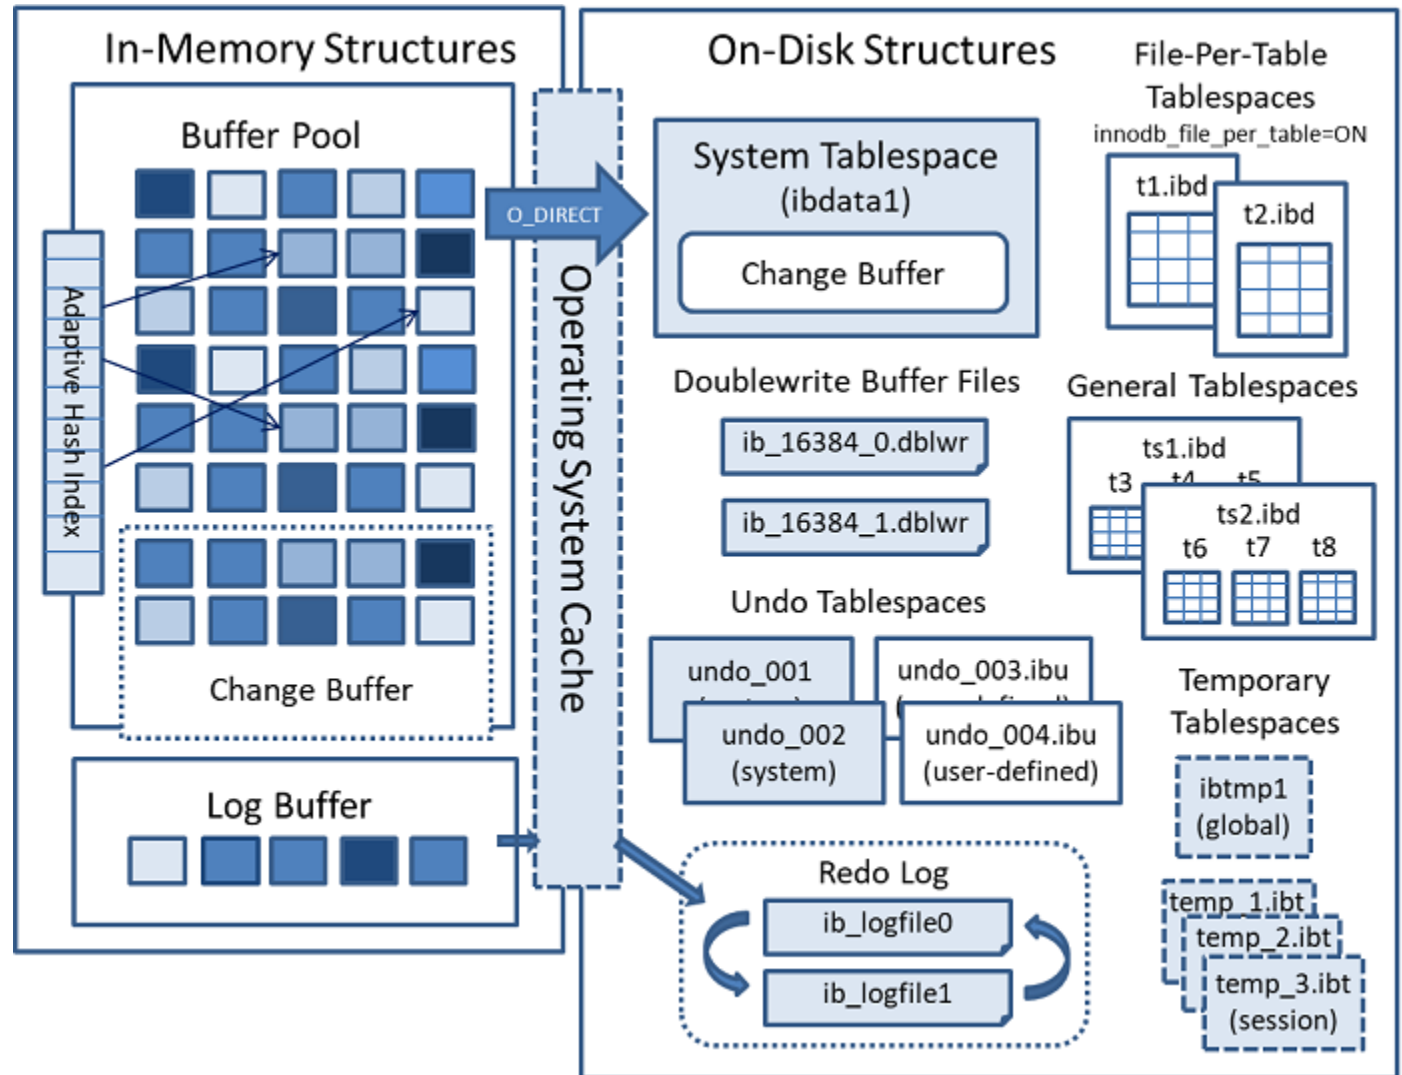
Stability or performance? Main Memory Management

- Rationale: reducing data access time without compromising memory stability
 - Use of buffer to cache data in faster memory
 - Deferred writing onto the secondary storage
- Organization
 - Buffer content accessed by page
 - Page may contain multiple rows and have
 - Transaction counter: how many transactions are using the page
 - Dirty flag: if true, page has been modified and must be aligned to secondary memory
- On dedicated DBMS servers, up to 80% of physical memory is often assigned to the buffer




Real DBMS architecture

- MySQL InnoDB data architecture



Use of Main Memory (buffer)

- Buffer management primitives:
 - fix
 - Responds to request by a transaction of loading of a page into the buffer. Returns a reference to the page and increments usage count
 - unfix
 - De-allocates page from the buffer and decrements usage count
 - force
 - Transfers synchronously a page from buffer to disk
 - setDirty
 - Modifies page status to denote that it has been changed
 - flush
 - Transfers asynchronously pages from buffer to disk, when page no longer needed (usage count == 0)



Execution of a fix primitive

- Searching for the target page
 - Search of the page in the buffer, if already there increment usage counter and return reference
 - Select a free (`usage_counter == 0`) page in the buffer (with FIFO or LRU policy); if found, return reference and increment usage counter. If `dirty_flag = true`, flush current page to disk before loading the new one
 - If no free page found, select page to de-allocate:
 - (**STEAL** policy) grab a victim page from an active transaction and flush it to disk; load new page increment usage counter and return reference
 - (**NO STEAL** policy) put the transaction in a wait list and manage the request when a page becomes free

Buffer Management Policies

- Write Policies: page writing to disk is normally asynchronous w.r.t. to transaction write operations
 - FORCE: pages are always transferred at commit
 - NO FORCE: transfer of pages can be delayed by the buffer manager
 - Normally buffer default configuration is:
 - NO STEAL, NO FORCE
- PRE-FETCHING (read-ahead in MySQL InnoDB)
 - anticipates loading of pages that are likely to be read
 - particularly effective in case of sequential reads
- PRE-FLUSHING
 - anticipates writing of de-allocated pages
 - effective for speeding up page fixing
- MySQL InnoDB LRU buffer management:
<https://dev.mysql.com/doc/refman/8.0/en/innodb-buffer-pool.html>

Buffer statistics in MySQL InnoDB

MySQL Workbench

Local instance MySQL80 x

File Edit View Query Database Server Tools Scripting Help

Navigator

SCHEMAS

Filter objects

- db2_schema
- db_expense_management
- db_gestione_spese
- dbtest
- jpa_book
- sakila
- sys
- world

SQL File 8* x SQL File 4* SQL File 5* mission SQL File 8*

1 • `SELECT * FROM information_schema.INNODB_BUFFER_POOL_STATS`

Result Grid

	POOL_ID	POOL_SIZE	FREE_BUFFERS	DATABASE_PAGES	OLD_DATABASE_PAGES	MODIFIED_DATABASE_PAGES	PENDING_DECOMPRESS
▶	0	512	251	256	0	0	0


- To understand statistics and fine tune the buffer pool configuration:
 - <https://dev.mysql.com/doc/refman/8.0/en/innodb-buffer-pool.html>
 - <https://dev.mysql.com/doc/refman/8.0/en/innodb-information-schema-buffer-pool-tables.html#innodb-information-schema-buffer-pool-stats-example>

Failure handling

- A transaction is an atomic transformation from an initial state into a final state
- Possible outcomes in absence/presence of failures:
 - **commit**: success
 - **rollback or fault before commit**: undo
 - **fault after commit**: redo
- Implications for recovery after failure
 - If failure occurs between commit and database buffers being flushed to secondary storage then, to ensure durability, Reliability Manager has to redo (rollforward) transaction updates.
 - If transaction had not committed at failure time, the Reliability Manager has to undo (rollback) any effects of that transaction for atomicity


Transactions and recovery

- DBMS starts at time t_0 , but fails at time t_f .
- Assume data for transactions T2 and T3 have been written to secondary storage (committed and permanently stored).
- T1 and T6 have to be undone.
- In absence of information of whether modified pages have been flushed, Reliability Manager has to redo T4 and T5.



Transaction Log

- A sequential file, made of records describing the actions carried out by the various transactions
- Written sequentially up to the top block (top = current instant)



Function of the Log

- It records on stable memory, in the form of state transitions, the actions carried out by the various transactions
 - if an UPDATE(U) operation transforms object O from value O1 to value O2
 - then the log records:
 - BEFORE-STATE(U) = O1
 - AFTER-STATE(U) = O2
- Logging INSERT and DELETE operations is identical to logging UPDATE operations, but
 - INSERT log record have no before state
 - DELETE log record have no after state

Example of MySQL log

```
binlogexample - Notepad
File Edit Format View Help
# original_commit_timestamp=1601194273790298 (2020-09-27 10:11:13.790298 W. Europe Summer Time)
# immediate_commit_timestamp=1601194273790298 (2020-09-27 10:11:13.790298 W. Europe Summer Time)
/*!80001 SET @@session.original_commit_timestamp=1601194273790298*//*!*/;
/*!80014 SET @@session.original_server_version=80019*//*!*/;
/*!80014 SET @@session.immediate_server_version=80019*//*!*/;
SET @@SESSION.GTID_NEXT= 'ANONYMOUS'/*!*/;
# at 1156
#200927 10:11:13 server id 1  end_log_pos 1359 CRC32 0x9c06a82a      Query  thread_id=12  exec_time=0  error_code=0  Xid = 15
SET TIMESTAMP=1601194273/*!*/;
/*!80013 SET @@session.sql_require_primary_key=0*//*!*/;
CREATE TABLE TableA (
  ID INT NOT NULL AUTO_INCREMENT PRIMARY KEY,
  Val CHAR(1),
  IntVal INT
)
/*!*/;
# at 1359
#200927 10:11:13 server id 1  end_log_pos 1438 CRC32 0xcb7cbbbe      Anonymous_GTID  last_committed=5  sequence_number=6  rbr_
/*!50718 SET TRANSACTION ISOLATION LEVEL READ COMMITTED*//*!*/;
# original_commit_timestamp=1601194273810786 (2020-09-27 10:11:13.810786 W. Europe Summer Time)
# immediate_commit_timestamp=1601194273810786 (2020-09-27 10:11:13.810786 W. Europe Summer Time)
/*!80001 SET @@session.original_commit_timestamp=1601194273810786*//*!*/;
/*!80014 SET @@session.original_server_version=80019*//*!*/;
/*!80014 SET @@session.immediate_server_version=80019*//*!*/;
SET @@SESSION.GTID_NEXT= 'ANONYMOUS'/*!*/;
# at 1438
#200927 10:11:13 server id 1  end_log_pos 1519 CRC32 0x13b7c80e      Query  thread_id=12  exec_time=0  error_code=0
SET TIMESTAMP=1601194273/*!*/;
BEGIN
/*!*/;
# at 1519
#200927 10:11:13 server id 1  end log pos 1586 CRC32 0x01dc5fb2      Table map: `db2 schema`.`tablea` mapped to number 84
<
```

This is what we did in the lesson on anomalies in MySQL!

Using the Log

- After: rollback-work or a failure that occurred before commit
 - **UNDO** T1: O = O1
- After: a failure that occurred after commit
 - **REDO** T1: O = O2
- **Idempotency** of UNDO and REDO:
 - $\text{UNDO}(T) = \text{UNDO}(\text{UNDO}(T))$
 - $\text{REDO}(T) = \text{REDO}(\text{REDO}(T))$
- Idempotency is necessary because the Reliability Manager may undo/redo operations twice, if its in doubt about the status of a write (flushed or not). Idempotency ensures that the effect is the same in any case

Types of Log Records

- Records concerning transactional commands:
 - $B(T), C(T), A(T)$
- Records concerning UPDATE, INSERT and DELETE operations
 - $U(T, O, BS, AS), I(T, O, AS), D(T, O, BS)$
- Records concerning recovery actions
 - $DUMP, CKPT(T_1, T_2, \dots, T_n)$
- Record fields:
 - T_i : transaction identifier
 - O : object identifier
 - BS, AS : before state, after state

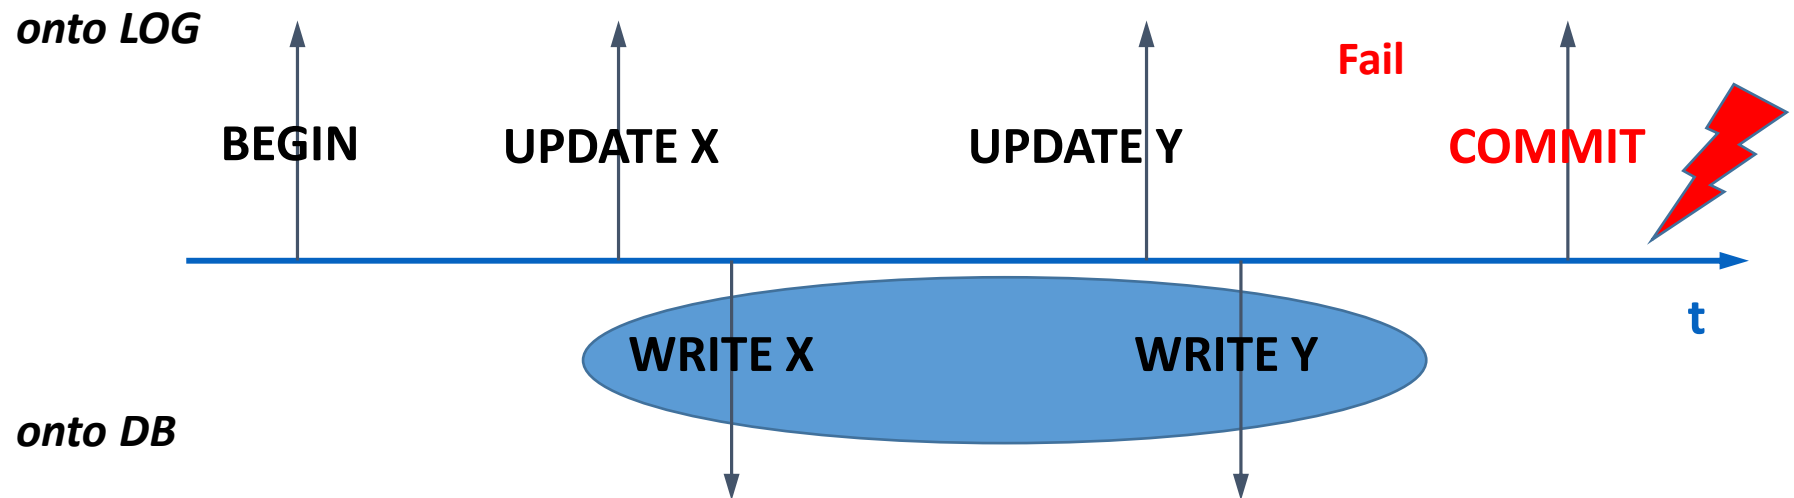
```
BINLOG '
YUlwXxMBAAAAQwAAAB8aAAAAAFgAAAAAAAEACmRiMl9zY2hlbWEABnRhYmx1YQADA/4DAv4EBgEB
AAID/P8A2U7wOA==
YUlwXx8BAAAAOgAAAFkaAAAAAFgAAAAAAAEAAgAD//8AAQAAAFBZAAAAAABAAAAAUGFAAAAFdbV
ng==
'/*!*/;
### UPDATE `db2_schema`.`tablea`
### WHERE
###   @1=1
###   @2='A'
###   @3=100|
### SET
###   @1=1
###   @2='A'
###   @3=133
# at 6745
#200927 10:12:17 server id 1  end_log_pos 6776 CRC32 0x4c5a23c6          Xid = 47
COMMIT/*!*/;
```

Transactional Rules

- Log management rules ensure transactions implement write operations in a way that supports reliability
- A commit log record must be written synchronously (with a force operation)
- **Write-Ahead-Log**
 - Before-state part of the record must be written in the log before actually carrying out the corresponding operation on the database
 - In this way, actions can always be undone
- **Commit Rule**
 - After-state part of the record must be written in the log before carrying out the commit
 - In this way, actions can always be redone
- NOTE: as database writes are asynchronous different implementations are possible, with an impact on recovery


Writing onto Log and Database

- Writing onto the database before commit
 - Redo is not necessary, because the state of the database reflects the state of the log



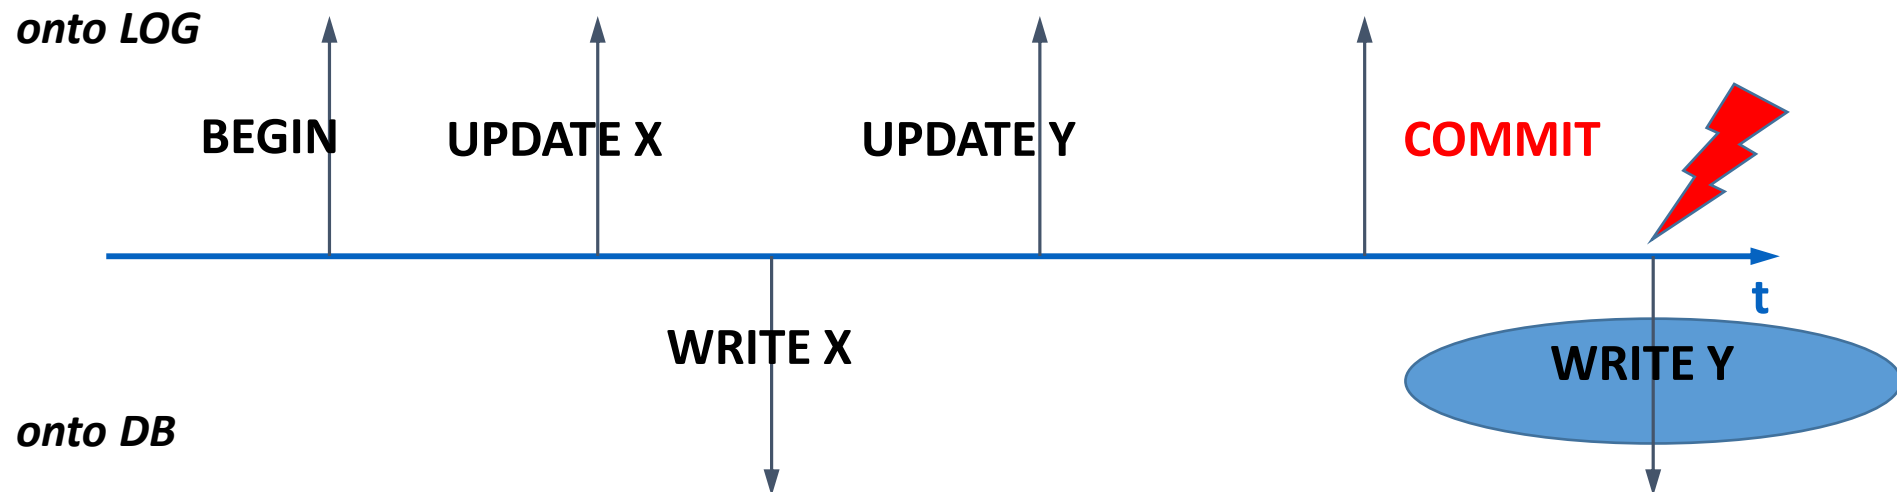
Writing onto Log and Database

- Writing onto the database after commit
 - Undo is not necessary
 - Does not require writing the before states of objects on the DB in order to abort



Writing onto Log and Database

- Writing onto the database in arbitrary points in time
 - Allows optimizing buffer management
 - Requires both undo and redo, in the general case



Recovery after a failure: types of failure

- **Soft failure**

- Loss of the content of (part of) the main memory
- Requires **warm restart**
- Log is exploited to replay transaction operations

- **Hard failure**


- Failure / loss of (part of) secondary memory devices
- Requires **cold restart**
- Dump is exploited to restore database content and log is exploited to replay transaction operations

- **Disaster**

- Loss of stable memory (i.e., of the log and dump)
- Not treated here

Checkpoint

- Performed periodically by the Reliability Manager to identify “consistent” time points
 - all transactions that committed their work flush their data from the buffer to the disk (effects are made persistent)
 - All active transactions (not committed yet) are recorded in the log
 - Aim: to record which transactions are still active at a given point in time (and, dually, to confirm that the others either did not start yet or have already finished)




Checkpoint

- Several possibilities/variants – a simple option is as follows:
 - Acceptance of all commit and abort requests is suspended
 - All “dirty” buffer pages **modified by committed transactions** are transferred to mass storage (synchronously, via force)
 - First the log entries recording the operations on the page data are forced, if not yet done
 - Then the page is flushed
 - The identifiers of the transactions still in progress are recorded in the CKPT record in the log (synchronously, via force); also, no new transaction can start while this recording takes place
 - Then, acceptance of operations is resumed
- In this way, there is guarantee that:
 - for all **committed** transactions, data are now on mass storage
 - transactions that are “half-way” are listed in the checkpoint log record (in stable memory)

Dump

- A time point in which a complete backup copy of the database is created (typically at night or in week-ends)
- The availability of that copy (dump) is recorded in the log
- The content of the dump is stored in stable memory




Warm Restart

- Loss of main memory buffer pages, not of table data on disk
- Requires “replaying” transactional operations and resolving in-doubt situations
- Log records are read starting from the last checkpoint (i.e., last stable point where all changes of committed transactions were flushed to disk)
 - CKPT record identifies the active transactions
- Active transactions are divided in two sets:
 - **UNDO** set: transactions to be undone
 - **REDO** set: transactions to be redone
- **UNDO and REDO** actions are executed

Warm Restart

- Find the most recent checkpoint
- Build the UNDO and REDO sets;
 - `UNDO = active transactions@CKPT; REDO = empty;`
 - For log records from CKPT to TOP
 - `IF B(Ti) then UNDO += Ti // started, may be undone`
 - `IF C(Tj) then UNDO -= Tj; REDO += Tj // ended, to redo`
- For all log records from TOP to 1st action of oldest T in UNDO
 - `undo(Ti) where Ti is in UNDO`
- For all log records from 1st action of oldest T in REDO to TOP
 - `redo(Ti) where Ti is in REDO`

3. Return to the 1st operation of the oldest active transaction while performing UNDO actions (in reverse log order)



4. Execution of REDO actions (in log order) until the top of the log

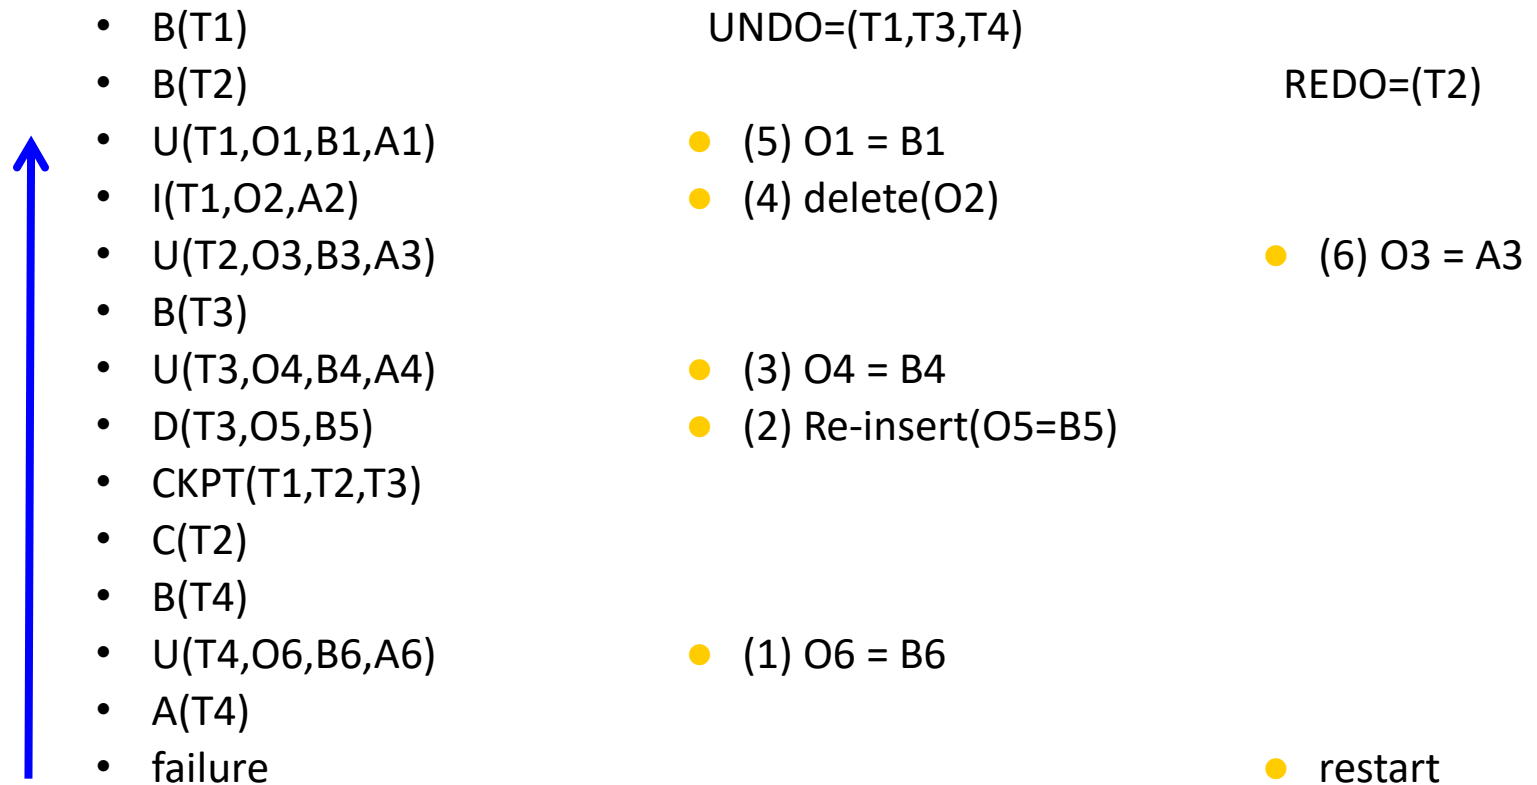
Example of Warm Restart

- B(T1)
- B(T2)
- U(T1, O1, B1, A1)
- I(T1, O2, A2)
- U(T2, O3, B3, A3)
- B(T3)
- U(T3, O4, B4, A4)
- D(T3, O5, B5)
- CKPT(T1, T2, T3)
- C(T2)
- B(T4)
- U(T4, O6, B6, A6)
- A(T4)
- failure



- UNDO=(T1, T2, T3), REDO=()
- UNDO=(T1, T3), REDO=(T2)
- UNDO=(T1, T3, T4), REDO=(T2)

Example of Warm Restart



Cold Restart

- Soft failure
 - Loss of the content of (part of) the main memory
 - Requires warm restart
- Hard failure
 - Failure / loss of (part of) secondary memory devices
 - Requires cold restart
- During cold restart
 - Data are restored starting from the last backup (dump)
 - The operations recorded onto the log until the failure time are executed (in log order)
 - Data on disk are restored in the status existing at the time of failure
 - A warm restart is then executed
 - Uncertain transactions are undone