

Machine Learning
Exercises

Christian Rossi

Academic Year 2023-2024

Abstract

The course topics are:

- Introduction: basic concepts.
- Learning theory:
 - Bias/variance tradeoff. Union and Chernoff/Hoeffding bounds.
 - VC dimension. Worst case (online) learning.
 - Practical advice on how to use learning algorithms.
- Supervised learning:
 - Supervised learning setup. LMS.
 - Logistic regression. Perceptron. Exponential family.
 - Kernel methods: Radial Basis Networks, Gaussian Processes, and Support Vector Machines.
 - Model selection and feature selection.
 - Ensemble methods: Bagging, boosting.
 - Evaluating and debugging learning algorithms.
- Reinforcement learning and control:
 - MDPs. Bellman equations.
 - Value iteration and policy iteration.
 - TD, SARSA, Q-learning.
 - Value function approximation.
 - Policy search. Reinforce. POMDPs.
 - Multi-Armed Bandit.

Contents

1	Exercise session I	1
1.1	Exercise one	1
1.1.1	Solution	1
1.2	Exercise two	2
1.2.1	Solution	2
1.3	Exercise three	2
1.3.1	Solution	3
1.4	Exercise four	4
1.4.1	Solution	4
2	Exercise session II	5
2.1	Exercise one	5
2.2	Exercise two	6
2.3	Exercise three	6
2.4	Exercise four	7
A	Linear regression	9
A.1	Introduction	9
A.2	Linear regression	9
A.2.1	Functions definition	10
A.2.2	Result evaluation	10
A.2.3	Python implementation	11
A.3	Statistical tests on coefficients	11
B	Classification	13
B.1	Introduction	13
B.1.1	Preliminary operations	14
B.2	Perceptron	14
B.2.1	Result evaluation	14
B.2.2	Python implementation	15
B.3	Logistic regression	15
B.3.1	Functions definition	15
B.3.2	Python implementation	16
B.3.3	Logit	16
B.4	Naïve Bayes	16
B.4.1	Functions definition	16
B.4.2	Generative method	17

B.4.3	Python implementation	17
B.5	K-nearest neighbor	17
B.6	Summary	18
B.7	Multiple classes	18

CHAPTER 1

Exercise session I

1.1 Exercise one

Given the relationship:

$$S = f(TV, R, N)$$

where S is the amount of sales revenue, TV , R and N are the amount of money spent on advertisements on TV programs, radio and newspapers, respectively, explain what are the:

1. Response.
2. Independent variables.
3. Features.
4. Model.

Which kind of problem do you think it is trying to solve?

1.1.1 Solution

In the proposed relationship we have:

1. The response (or target or output) is the amount of sales S .
2. The independent variables (or input) are TV , R and N .
3. The features (or input) are TV , R and N .
4. The model is identified by the function $f(\cdot)$.

Since the amount of sales S is a continuous and ordered variable, we are trying to solve a regression problem (supervised learning).

1.2 Exercise two

Why is linear regression important to understand? Select all that apply and justify your choice:

1. The linear model is often correct.
2. Linear regression is extensible and can be used to capture nonlinear effects.
3. Simple methods can outperform more complex ones if the data are noisy.
4. Understanding simpler methods sheds light on more complex ones.
5. A fast way of solving them is available.

1.2.1 Solution

1. False: it rarely happens that the problem we are modeling has linear characteristics.
2. True: it is true that this model is easy to interpret and can be extended to also consider nonlinear relationships among variables, e.g., using basis functions.
3. True: since we are able only to minimize the discrepancy between the considered function and the real one, and we can not reduce the variance introduced by noise, the use of linear model might be a better choice with respect to more complex ones since they usually are prone to overfitting, i.e., they try to model also the noise of the considered process.
4. True: they are easy to interpret and might give suggestions on more sophisticated techniques which can be used to tackle specific problems.
5. True/False: for some loss functions we have a closed form solution for linear model (LS method), thus we can guarantee that we are able to find the parameters minimizing the loss function effectively. That is not always true and depends also on the loss function we want to minimize.

1.3 Exercise three

Consider a linear regression with input x , target y and optimal parameter θ^* .

1. What happens if we consider as input variables x and $2x$?
2. What we expect on the uncertainty about the parameters we get by considering as input variables x and $2x$?
3. Provide a technique to solve the problem.
4. What happens if we consider as input variables x and x^2 ?

Motivate your answers.

1.3.1 Solution

The original formulation is:

$$y = \theta^* x$$

1. In this scenario, the formulation simplifies to:

$$y = \theta_1 x + \theta_2 2x$$

As the two variables are dependent, it can alternatively be expressed as:

$$y = \underbrace{(\theta_1 + 2\theta_2)}_{\theta^*} x$$

Thus, yielding the original formulation:

$$y = \theta^* x$$

Moreover, for computing the closed-form optimization, the formula employed is:

$$\omega = (x^T x)^{-1} x t$$

However, in this particular case, the matrix x takes the form:

$$x = \begin{bmatrix} x_1 & 2x_1 \\ x_2 & 2x_2 \\ \vdots & \vdots \\ x_n & 2x_n \end{bmatrix}$$

Hence, $x^T x$ becomes singular, rendering the previous formula inapplicable.

In general, if x lacks full rank, then $x^T x$ becomes singular, and its inverse cannot be computed.

2. The parameter we get have a high variance, since we have an infinite number of couples of parameters minimizing the loss of the samples in the considered problem. Indeed, if the parameters of the two inputs are w_1 and w_2 we would have that the true relationship would be:

$$y = \theta_1 x + \theta_2 2x = (\theta_1 + 2\theta_2) x$$

Which can be satisfied by an infinite number of solutions.

3. In this case, the use of ridge regression is able to partially cope with the influence of using highly linearly correlated features. Another viable option is to remove the variables which are linearly dependent, for instance by checking if they have correlation equal to 1 or -1 .
4. In this case we do not have a badly conditioned matrix since x and x^2 are not linearly dependent and the corresponding design matrix would not be ill-conditioned. As a result, we can find a closed form optimization.

1.4 Exercise four

After performing Ridge regression on a dataset with $\lambda = 10^{-5}$ we get one of the following one set of eigenvalues for the matrix $(\Phi^T \Phi + \lambda I)$:

1. $\Lambda = \{0.00000000178, 0.014, 12\}$
2. $\Lambda = \{0.0000178, -0.014, 991\}$
3. $\Lambda = \{0.0000178, 0.014, 991\}$
4. $\Lambda = \{0.0000178, 0.0000178, 991\}$

Explain whether these sets are plausible solutions or not.

1.4.1 Solution

Since the matrix $(\Phi^T \Phi + \lambda I)$ is definite positive and its eigenvalues should all be greater than $\lambda = 10^{-5}$, we have:

1. Not plausible: one eigenvalue is smaller than 10^{-5} .
2. Not plausible: one eigenvalue is negative.
3. Plausible: all positive and greater than 10^{-5} .
4. Plausible: all positive and greater than 10^{-5} .

CHAPTER 2

Exercise session II

2.1 Exercise one

Classify the following variable as quantitative or qualitative:

1. Height.
2. Age.
3. Speed.
4. Color.

Provide a technique for transforming qualitative data into quantitative format without imposing additional organization on the data.

Solution

The classification is as follows:

1. Quantitative variable: heights can be ordered and typically belong to a bounded continuous set.
2. Quantitative variable: values are ordered natural numbers.
3. Quantitative variable: takes real number values.
4. Qualitative variable: since there's no inherent order among colors, one-hot encoding can be employed without imposing additional structure on the data.

For a set of all possible colors, $\mathcal{C} = \{c_1, \dots, c_p\}$ one-hot encoding creates a binary variable $z_i \in \{0, 1\}$ for each color c_i . This variable equals one when the color is c_i . Thus, color c_i is represented by a binary vector $z_i = [z_1 \dots z_n]^T$, where $z_i = 1$ and all other $z_j = 0$ for $j \neq i$. This method introduces p new variables without further structuring the data. Notably, any two vectors $z_i \neq z_j$ are equally distant under reasonable metrics like Euclidean distance.

It's important to note that assigning a quantitative variable i to each color would introduce additional structure to the data, which should generally be avoided. While this approach requires only one variable instead of p , it imposes an ordering among the colors. Moreover, it results in color c_i being closer to color $c_i + 1$ than to color $c_i + 2$ in terms of Euclidean distance.

2.2 Exercise two

Consider a dataset comprising workers' attributes such as the number of hours spent working (x_1), the number of completed projects (x_2), and whether they received a bonus (t). After applying logistic regression, we obtain the following coefficients: $w_0 = -6$, $w_1 = 0.05$, and $w_2 = 1$.

1. Determine the likelihood of a worker receiving a bonus given that they worked for 40 hours and completed 3.5 projects.
2. Calculate the number of hours a worker needs to work to have a 50% chance of receiving a bonus.
3. Discuss whether values of z in $\sigma(z)$ lower than -6 are meaningful in this context, and provide reasoning.

Solution

1. The logistic model yields the probability of receiving a bonus as its output, expressed by:

$$P(t = 1|\mathbf{x}) = \sigma(w_0 + w_1x_1 + w_2x_2)$$

Given $x_1 = 40$ and $x_2 = 3.5$:

$$P(t = 1|\mathbf{x}) = \sigma(-6 + 0.05 \cdot 40 + 1 \cdot 3.5) = \sigma(-0.5) = 0.3775$$

2. To ascertain the probability of receiving a bonus with a confidence level $\alpha\%$, we need to invert the sigmoidal function. However, in this instance, a 50% chance corresponds to the sigmoidal argument being zero. Hence:

$$w_0 + w_1\hat{x} + w_2x_2 = 0 \rightarrow -6 + 0.05\hat{x} + 3.5 = 0 \rightarrow \hat{x} = 50$$

3. Considering that all the variables under consideration are positive definite, it is reasonable to regard predictions with values greater than -6 as meaningful.

2.3 Exercise three

Consider a binary classifier trained on a dataset comprising $N = 100$ samples.

1. Given a precision of 0.25 and an F1 score of 0.4, compute the recall.
2. Additionally, with an accuracy of 0.85, calculate the complete confusion matrix.
3. Under what circumstances is accuracy not a dependable metric for evaluating the model's quality?

Solution

1. We have:

$$F1 = \frac{2 \cdot \text{Pre} \cdot \text{Rec}}{\text{Pre} + \text{Rec}} = 0.4$$

$$\text{Pre} = \frac{TP}{TP + FP} = 0.25$$

We seek to find:

$$\text{Rec} = \frac{TP}{TP + FN}$$

Substituting, we obtain:

$$\frac{2 \cdot 0.25 \cdot \text{Rec}}{0.25 + \text{Rec}} = 0.4 \rightarrow \text{Rec} = 1$$

2. Given:

$$F1 = \frac{2 \cdot \text{Pre} \cdot \text{Rec}}{\text{Pre} + \text{Rec}} = 0.4$$

$$\text{Pre} = \frac{TP}{TP + FP} = 0.25$$

$$\text{Rec} = \frac{TP}{TP + FN} = 1$$

$$\text{Acc} = \frac{TP + TN}{N} = 0.85$$



We infer that $FN = 0$ since the recall is unity. Then, from the other formulas:

$$\begin{cases} TP + TN = 85 \\ FN = 0 \\ TP = 0.25(TP + FP) \\ TP + TN + FN + FP = 100 \end{cases} \rightarrow \begin{cases} TP = 5 \\ FN = 0 \\ FP = 15 \\ TN = 80 \end{cases}$$

3. Accuracy is not a reliable indicator of the model's quality primarily under two conditions: when the dataset is imbalanced, and when the consequences of misclassifying positive-class samples differ from misclassifying negative-class samples.

2.4 Exercise four

Consider the provided datasets:



Now, let's analyze whether the learning procedure terminates and the number of steps required for convergence using the online stochastic gradient descent algorithm to train a perceptron.

Solution

The perceptron learning algorithm converges if there exists a linear separation hyperplane. In such a scenario, the classification error can be reduced to zero. If no linear separation exists, the optimization process doesn't halt. The convergence rate isn't assured since it relies on the initial parameterization and the sequence of points used for training. Nonetheless, convergence does occur within a finite number of steps.

In the first dataset (left), convergence isn't guaranteed. However, in the second dataset (right), the online stochastic gradient descent will ultimately converge within a finite number of steps.

Considering that the loss function for logistic regression is convex, the online learning process converges to the global optimum asymptotically, irrespective of the dataset provided.

APPENDIX A

Linear regression

A.1 Introduction

Let's examine the Iris dataset, which comprises the following features for each sample:

1. Sepal length.
2. Sepal width.
3. Petal length.
4. Petal width.
5. Species (including *Iris setosa*, *Iris virginica*, and *Iris versicolor*).

This dataset consists of a total of $N = 150$ samples, with each species contributing 50 samples.

Using linear regression, we can extract valuable insights from the data. From the data, we can infer relationships between variables and make predictions based on them. We can provide forecasts for various quantities using newly observed data. Specifically, we can predict the petal width of a particular type of *Iris setosa* by leveraging the relationship with petal length through linear regression. In this scenario, the target variable is continuous ($t_n \in \mathbb{R}$), indicating a regression problem.

A.2 Linear regression

Preliminary operations The preliminary operations on data are:

- Loading data: load the dataset into memory.
- Inspecting data: examine the dataset to understand its structure and contents.
- Selecting interesting data: identify and choose the relevant features or variables from the dataset.
- Preprocessing: prepare the data for further analysis by performing various preprocessing steps such as:

- Shuffling the data to randomize the order of samples (`shuffle()`).
- Removing inconsistent data points that may contain errors or inconsistencies.
- Removing outliers to ensure data quality.
- Normalizing or standardizing the data to bring all features to a similar scale.
- Filling missing data, for example, by replacing `NaN` values with appropriate values.

These preprocessing steps help ensure that the data is clean, consistent, and ready for analysis.

Data normalization We can normalize a series of samples $\{s_1, \dots, s_N\}$ to a sample s using:

- Z-score:

$$\frac{s - \bar{s}}{S}$$

Here, $\bar{s} = \frac{1}{N} \sum_{n=1}^N s_n$ and $S^2 = \frac{1}{N-1} \sum_{n=1}^N (s_n - \bar{s})^2$.

- Min-max feature scaling:

$$\frac{s - s_{min}}{s_{max} - s_{min}}$$

Here, $s_{max} = \max_{n \in \{1, \dots, N\}} s_n$ and $s_{min} = \min_{n \in \{1, \dots, N\}} s_n$.

A.2.1 Functions definition

To proceed with our analysis, we must define the following elements:

- Hypothesis space: we consider linear models represented by:

$$\hat{t} = y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j x_j = \mathbf{w}^T \mathbf{x}$$

Here, $\mathbf{w} = [w_0, w_1, \dots, w_{M-1}]^T$ and $\mathbf{x} = [x_0, x_1, \dots, x_{M-1}]^T$.

- Loss function: the loss function is defined as the residual sum of squares over the N samples $\{(\mathbf{x}_n, t_n)\}_{n=1}^N$:

$$\text{RSS}(\mathbf{w}) = \sum_{n=1}^N (y(\mathbf{x}_n, \mathbf{w}) - t_n)^2$$

- Optimization method: various optimization methods can be employed, including closed form solutions, gradient descent, and others, to minimize the loss function and find the optimal values for the model parameters.

A.2.2 Result evaluation

Evaluation of the linear regression results involves several metrics:

- Residual sum of squares (RSS) or sum of squared errors (SSE):

$$\text{RSS}(\mathbf{w}) = \sum_{n=1}^N (\hat{t}_n - t_n)^2$$

Here, $\hat{t}_n = y(\mathbf{x}_n, \mathbf{w})$.

- Mean square error (MSE):

$$\text{MSE} = \frac{\text{RSS}(\mathbf{w})}{N}$$

- Root mean square error (RMSE):

$$\text{RMSE} = \sqrt{\frac{\text{RSS}(\mathbf{w})}{N}}$$

- Coefficient of determination (R^2):

$$R^2 = 1 - \frac{\text{RSS}(\mathbf{w})}{\text{TSS}}$$

Here $\text{TSS} = \sum_{n=1}^N (\bar{t} - t_n)^2$ is the total sum of squares and $\bar{t} = \frac{1}{N} \sum_{n=1}^N t_n$.

- Degrees of freedom (dfe):

$$\text{dfe} = N - M$$

- Adjusted coefficient of determination (R_{adj}^2):

$$R_{adj}^2 = 1 - (1 - R^2) \frac{N - 1}{\text{dfe}}$$

A.2.3 Python implementation

We have several solutions available from different libraries:

- Utilizing `sklearn` with `LinearRegression()`.
- Utilizing `statsmodels` with `OLS`.
- Implementing it manually as $\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t}$.

With the `sklearn` option:

- Initialize a linear model (`LinearRegression()`).
- Fit the model to the data (`fit()`).
- Analyze the results.

A.3 Statistical tests on coefficients

Let's consider the scenario where t_n satisfies $t_n = \mathbf{w}\mathbf{x}_n + \epsilon_n$, where ϵ_n is an independent and identically distributed white zero-mean noise with variance σ^2 . In this case, the exact distribution of the statistic is given by:

$$\frac{\hat{w}_j - w_j}{\hat{\sigma} \sqrt{v_j}} \sim t_{N-M}$$

Here:

- w_j represents the true parameters.

- \hat{w}_j denotes the estimated parameter with N samples.
- v_j stands for the j -th diagonal element of the matrix $(\mathbf{X}^T \mathbf{X})^{-1}$.
- t_{N-M} is the Student's t -distribution with $\text{dfe} = N - M$ degrees of freedom.
- $\hat{\sigma}^2$ is the unbiased estimate for the target variance:

$$\hat{\sigma}^2 = \frac{\text{RSS}(\hat{\mathbf{w}})}{N - M}$$

Based on this, we can conduct the following tests:

- Test on single coefficients $j \in \{0, \dots, M-1\}$: we compare $H_0 : w_j = 0$ against $H_1 : w_j \neq 0$. Then, we apply:

$$t_{stat} = \frac{\hat{w}_j - w_j}{\hat{\sigma} \sqrt{v_j}} \sim t_{N-M}$$

Here, t_{N-M} follows the Student's t -distribution with $N - M$ degrees of freedom.

- Test on the overall significance of the model: we compare $H_0 : w_1 = \dots = w_{M-1} = 0$ against $H_1 : \exists j \in \{1, \dots, M-1\}$ such that $w_j \neq 0$. Then, we use:

$$F_{stat} = \frac{N - M}{M - 1} \frac{\text{TSS} - \text{RSS}(\hat{\mathbf{w}})}{\text{RSS}(\hat{\mathbf{w}})} \sim F_{M-1, N-M}$$

Here, $F_{M-1, N-M}$ represents the Fisher-Snedecor distribution with parameters $M - 1$ and $N - M$. This test compares the full linear model $y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j x_j$ with $N - M$ degrees of freedom against the constant model $y(\mathbf{x}, w_0) = w_0$ with $N - 1$ degrees of freedom.

APPENDIX B

Classification

B.1 Introduction

Let's examine the Iris dataset, which comprises the following features for each sample:

1. Sepal length.
2. Sepal width.
3. Petal length.
4. Petal width.
5. Species (including Iris setosa, Iris virginica, and Iris versicolor).

This dataset consists of a total of $N = 150$ samples, with each species contributing 50 samples.

Using petal and sepal length and width variables, it is possible to predict the species of Iris (target). In this context, the target variables are discrete and non-metric, indicating a classification problem. The targets are referred to as classes.

To tackle the problem of predicting the species of Iris based on petal and sepal measurements, several approaches can be considered:

- *Discriminant function approach*: in this method, the model functions as a mapping of inputs to classes, expressed as $f(x) = C_k \in \{C_1, \dots, C_K\}$. The process involves fitting the model to the available data.
- *Probabilistic discriminative approach*: here, the model represents a conditional probability distribution, $P(C_k|\mathbf{x}) \in [0, 1]$. The approach entails fitting the model to the data to establish probabilities.
- *Probabilistic generative approach*: the model incorporates the likelihood $P(\mathbf{x}|C_k) \in [0, 1]$ and the prior $P(C_k) \in [0, 1]$. The process involves fitting the model to the available data. Inference is made using the posterior probability with the Bayes rule. New samples can be generated from the joint distribution $P(C_k|\mathbf{x}) = P(\mathbf{x}|C_k)P(C_k)$.

The classification problem can be addressed through various methods, including linear classification methods like the perceptron and logistic regression, alongside algorithms like naïve Bayes and K -nearest neighbor.

B.1.1 Preliminary operations

As usual before solving the problem we need to perform some preliminary operations:

- Load the data.
- Consistency checks.
- Select and normalize the input.
- Shuffle the data (`shuffle()`).
- Generate the output ($t_n \in \{0, 1\}$ or $t_n \in \{-1, 1\}$).
- Explore the selected data (`scatter`).

B.2 Perceptron

The perceptron operates as a discriminant function approach since it directly assigns elements into classes without providing probabilities.

Functions definition To proceed with our analysis, we must define the following elements:

- Hypothesis space: we consider linear models represented by:

$$y(\mathbf{x}) = \text{sign}(\mathbf{w}^T \mathbf{x}) = \text{sign}(w_0 + x_1 w_1 + x_2 w_2)$$

Here,

$$\text{sign}(z) = \begin{cases} -1 & \text{if } z < 0 \\ +1 & \text{otherwise} \end{cases}$$

- Loss function: distance of misclassified points in $\{(\mathbf{x}_n, t_n)\}_{n=1}^N$ with $t_n \in \{-1, 1\}$:

$$L_P(\mathbf{w}) = - \sum_{n \in \mathcal{M}} \mathbf{w}^T \mathbf{x}_n t_n$$

Here, $\mathcal{M} = \{n \in \{1, \dots, N\} : t_n \neq y(\mathbf{x}_n)\}$

- Optimization method: online gradient descent

B.2.1 Result evaluation

This method converges when the data are linearly separable. To visualize the separating hyperplane or decision boundary (line) we need to plot:

$$\text{sign}(\mathbf{x}^T \mathbf{x}) = 0 \rightarrow \text{sign}(w_0 + x_1 w_1 + x_2 w_2) = 0 \rightarrow x_2 = -\frac{w_1 x_1 + w_0}{w_2}$$

To evaluate the performance of a classifier, we can compute the confusion matrix which tells us the number of points which have been correctly classified and those which have been misclassified

		<i>Actual class</i>	
		1	0
<i>Predicted class</i>	1	True positive	False positive
	0	False negative	True negative

From this table we can compute the following measures:

- Accuracy: fraction of the samples correctly classified in the dataset:

$$\text{Acc} = \frac{TP + TN}{N}$$

- Precision: fraction of samples correctly classified in the positive class among the ones classified in the positive class:

$$\text{Pre} = \frac{TP}{TP + FP}$$

- Recall: fraction of samples correctly classified in the positive class among the ones belonging to the positive class:

$$\text{Rec} = \frac{TP}{TP + FN}$$

- F1 score: harmonic mean of the precision and recall:

$$\text{F1} = \frac{2 \cdot \text{Pre} \cdot \text{Rec}}{\text{Pre} + \text{Rec}}$$

The higher these figures of merits the better the algorithm is performing. These performance measures are not symmetric, but depend on the class we selected as positive. Depending on the application one might switch the classes to have measures which better evaluate the predictive power of the classifier.

B.2.2 Python implementation

Implementation in Python can be accomplished either through the `sklearn` library utilizing the `Perceptron` module, or by manually coding the algorithm.

B.3 Logistic regression

Logistic regression functions as a probabilistic discriminative approach by directly assigning probabilities to elements belonging to certain sets.

B.3.1 Functions definition

To proceed with our analysis, we must define the following elements:

- Hypothesis space:

$$y(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) = \sigma(w_0 + x_1 w_1 + x_2 w_2)$$

Here, $\sigma = \frac{1}{1+e^z}$.

- Loss function: distance of misclassified points in $\{(\mathbf{x}_n, t_n)\}_{n=1}^N$ with $t_n \in \{-1, 1\}$:

$$L_P(\mathbf{w}) = P(\mathbf{y}|\mathbf{X}, \mathbf{w}) = \sum_{n=1}^N t_n \ln y(\mathbf{x}_n) + (1 - t_n) \ln(1 - y(\mathbf{x}_n))$$

- Optimization method: online gradient descent.

Result evaluation To visualize the separating hyperplane or decision boundary (line) we need to plot:

$$\sigma(\mathbf{x}^T \mathbf{x}) = \frac{1}{2} \rightarrow \sigma(w_0 + x_1 w_1 + x_2 w_2) = \frac{1}{2} \rightarrow x_2 = -\frac{w_1 x_1 + w_0}{w_2}$$

B.3.2 Python implementation

Implementation in Python can be accomplished through the `sklearn` library utilizing the `LogisticRegression` module.

B.3.3 Logit

Given the function:

$$\text{logit}(y) = \log\left(\frac{y}{1-y}\right)$$

We can utilize it with the output of logistic regression:

$$\text{logit}(y(\mathbf{x})) = \mathbf{w}^T \mathbf{x} = w_0 + x_1 w_1 + x_2 w_2$$

This exhibits a similar characterization to linear regression. Consequently, we can conduct hypothesis testing to ascertain the significance of the parameters.

B.4 Naïve Bayes

The naïve assumption posits that within the class C_k , each input is conditionally independent of one another. In this scenario the decision function, given a prior $p(C_k)$, maximize the Maximum A Posteriori (MAP) probability:

$$y(\mathbf{x}) = \underset{k}{\operatorname{argmax}} P(C_k) \prod_{i=1}^M P(x_i | C_k)$$

B.4.1 Functions definition

So with this approach we have:

- Hypothesis space:

$$y(\mathbf{x}) = \underset{k}{\operatorname{argmax}} P(C_k) \prod_{i=1}^M P(x_i | C_k)$$

Here, $\sigma = \frac{1}{1+e^z}$.

- Loss function: log likelihood for fitting both the priors $P(C_k)$ and the likelihoods $P(x_j|C_k)$.
- Optimization method: maximum likelihood estimation (MLE).

In our classification problem, we opt for the following:

- Prior: $P(C_k)$ multinomial distribution with parameters (p_1, \dots, p_k) .
- Likelihood: $P(x_j|C_k)$ follows a normal distribution $\mathcal{N}(\mu_{jk}, \sigma_{jk}^2)$, where each feature x_j and each class C_k are represented.

We may select different distributions for the features based on the input.

B.4.2 Generative method

Due to the generative capabilities of the naïve Bayes classifier, we can generate datasets that resemble the original one using the following steps:

1. Select a class $C_{\hat{k}}$ according to the prior multinomial distribution with parameters:

$$\hat{P}(C_1), \dots, \hat{P}(C_k)$$

2. For each feature j , draw a sample x_j from $\mathcal{N}(\hat{\mu}_{j\hat{k}}, \hat{\sigma}_{j\hat{k}}^2)$.
3. Repeat these steps whenever a new sample is desired.

B.4.3 Python implementation

We can implement naïve Bayes in Python using two approaches:

- Utilizing pre-implemented `sklearn GaussianNB`, with a multinomial distribution as prior and Gaussian distributions as likelihood.
- Implementing it by hand:

- Estimate the prior: $\hat{P}(C_k) = \frac{\sum_{i=1}^N I\{\mathbf{x}_n \in C_k\}}{N}$
- Estimate the Maximum Likelihood Estimation (MLE) parameters: $P(x_j|C_k) = \mathcal{N}(x_j; \hat{\mu}_{jk}, \hat{\sigma}_{jk}^2)$, where $\hat{\mu}_{jk}$ and $\hat{\sigma}_{jk}^2$ are computed by maximizing the likelihood.
- Compute $P(C_k) \prod_{j=1}^M P(x_j|C_k)$ for each class C_k and choose the maximum one.

Notice that naïve Bayes isn't strictly a Bayesian method, as the priors are estimated from data and not updated using likelihoods.

B.5 K-nearest neighbor

The k -nearest neighbor is a discriminative function approach.

1-nearest neighbor The concept revolves around leveraging nearby points to predict the target of a new point. Given a dataset $\{\mathbf{x}_n, t_n\}_{n=1}^N$ and a new data point \mathbf{x}_q , we predict the target as:

$$i_q \in \underset{n \in \{1, \dots, N\}}{\operatorname{argmin}} \|\mathbf{x}_q - \mathbf{x}_n\|_2$$

Resulting in $\hat{t}_q = t_{i_q}$. This approach seamlessly caters to both regression and classification tasks without explicit training. The training process is essentially querying the dataset.

Key design choices include the selection of the distance function and the number of neighbors. If $k > 1$, the targets can be combined as follows:

- For classification, predict the mode class (with tie-breaking rules):

$$\hat{t}_q \in \underset{C_k}{\operatorname{argmax}} |\{i \in \mathcal{N}_k(\mathbf{x}_q) : t_i = C_k\}|$$

- For regression, predict the average target:

$$\hat{t}_q = \frac{1}{k} \sum_{i \in \mathcal{N}_k(\mathbf{x}_q)} t_i$$

- Other approaches include providing a probability distribution instead of a class, using weights proportional to the inverse of the distance.

The choice of k introduces varying degrees of regularization, ranging from strong to mild regularization.

B.6 Summary

	Parametric	Frequentist	Category
<i>Perceptron</i>	✓	✓	Discriminative function
<i>Logistic regression</i>	✓	✓	Probabilistic discriminative
<i>Naïve Bayes</i>	✓	✓	Probabilistic generative
<i>K-Nearest neighbor</i>	×	×	Discriminative function

B.7 Multiple classes

If we're dealing with multiple classes, we can utilize the same function by providing a target with more than two labels. This approach involves training k different models, each distinguishing one class from the rest. In this scenario, the parameter vector becomes a matrix W .

Although we can still visualize the separating surfaces, it becomes slightly more challenging compared to the binary classification case. However, extending these methods to handle multiple classes doesn't require any changes.