

Software Engineering II

Theory

Christian Rossi

Academic Year 2023-2024

Abstract

The objective of the course is to teach the principals methods and processes of software engineering needed to develop complex and qualitative software.

The course covers the following arguments:

- Software process and its organization.
- Modelling languages.
- Requirements analysis and definition.
- Software development methods and tools.
- Approaches for verify and validate the software.

Contents

1	Introduction	1
1.1	Definition	1
1.2	History	1
1.3	The process and product	2
1.4	Development process	3
2	Requirements engineering	6
2.1	Definition	6
2.2	Importance and difficulties	6
2.3	Requirement engineering process	7
2.4	World-machine relationship	8
2.5	Elicitation of requirements	9
2.6	Modeling requirements	9
2.7	Use cases and requirements	11
2.8	Requirements-level class diagrams	12
2.9	Dynamic modeling	12
3	Alloy	14
3.1	Introduction	14
3.2	Syntax	14
3.3	Address book example	17
3.4	Family relations example	18
4	Requirement analysis and specification	20
4.1	Structure of a RASD document	20
5	Software design	23
5.1	Software architecture	23
5.2	Software design description and principles	26
5.3	Architectural styles	28
5.4	Interface design and documentation	34

CHAPTER 1

Introduction

1.1 Definition

The realm of software engineering is dedicated to unraveling the multitude of challenges that emerge during the creation of extensive software systems. These systems are inherently intricate due to their substantial scale, the collaboration of individuals from diverse fields, and the necessity for continuous adjustments to meet evolving demands (both during development and post-installation).

Definition

Software engineering is a methodical and managerial discipline that revolves around the systematic creation and upkeep of software products, all of which are crafted and sustained within predefined, controlled timeframes and cost constraints.

In contrast to traditional programming, where a programmer typically crafts an entire piece of software based on known specifications independently, software engineers embark on a distinct path. They identify requirements, shape specifications, design components meant to interconnect with others, and engage in collaborative efforts within a team setting. The key competencies of a software engineer encompass technical acumen, managerial prowess, cognitive aptitude, and organizational skills.

1.2 History

In its early days, software was regarded as an art form. Computers primarily served the purpose of mathematical problem-solving, with the designers themselves acting as the end users. The initial programs were crafted using low-level languages and were subject to stringent resource constraints.

However, as the demand for customized software surged, this artistic endeavor transitioned into a more methodical craft. Developers began creating programs tailored for a broader audience, employing new high-level languages. Towards the culmination of this era, a "software crisis" loomed, marked by the escalating complexity of software and a dearth of effective software development techniques.

With the intention of tackling this growing challenge, the term was coined during a NATO conference in 1968. The primary focus of this pivotal gathering encompassed the following key

areas:

- Development of software and the establishment of standards.
- Strategic planning and proficient management.
- Automation of software development processes.
- Modularization of software design.
- Rigorous quality assurance and verification.

1.3 The process and product

The creation of a software program necessitates a well-defined process. Both the software itself and the procedures employed possess inherent quality, and it is imperative for the software engineer to strive for the highest quality possible. This is essential because the process directly influences the ultimate outcome.

Software distinguishes itself from conventional product types in several ways, as it is:

1. Intangible, making it challenging to precisely describe and evaluate.
2. Adaptable or malleable, capable of being modified to suit evolving requirements.
3. Human-intensive, devoid of straightforward manufacturing processes.

Software quality is influenced by various critical factors, including development technology, process quality, the caliber of individuals involved, cost considerations, and adherence to predefined schedules. Software quality attributes encompass:

- Correctness: ensuring that the software aligns with specified requirements.
- Reliability: minimizing the probability of failure absence over a defined period.
- Robustness: assessing the software's ability to perform reasonably in unforeseen scenarios.
- Performance: evaluating the efficient utilization of resources.
- Usability: focusing on the ease of use for the anticipated user base.
- Maintainability: the capacity to facilitate software maintenance.
- Reusability: extending maintainability to components, promoting their reuse.
- Portability: the adaptability of software to diverse target environments.
- Interoperability: ensuring harmonious coexistence and cooperation with other applications.

Process quality attributes consist of:

- Productivity: assessing the efficiency and output of the development process.
- Unity of effort (person-month): measuring the collective contribution of effort over a month.

- Delivered item (lines of code and function points): quantifying the tangible output of the development process.
- Timeliness: the capability to respond to change requests in a timely fashion.

1.4 Development process

In the initial stages of software development, there was a lack of reference models, leading to a straightforward "code and fix" approach. In response to the previously mentioned software crisis, the necessity for a structured model became evident. The first comprehensive model introduced was the "waterfall" model, characterized by the following fundamental requirements:

1. Identification of distinct phases and associated activities.
2. Enforcement of a linear progression, with no backtracking between phases.
3. Standardization of outputs generated at each phase.
4. Viewing software development as akin to a manufacturing process.

Following the advent of the waterfall model, numerous more flexible methodologies emerged, including iterative models, the agile movement, and DevOps, which aimed to offer more adaptable approaches to software development.

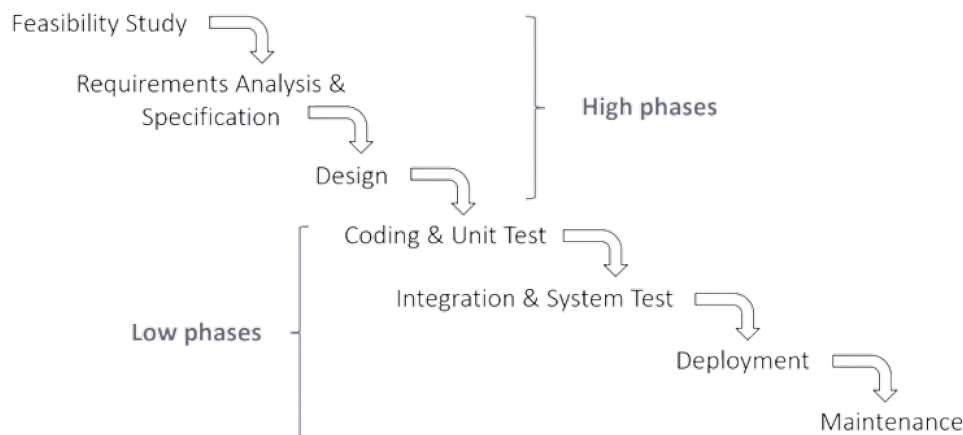


Figure 1.1: Waterfall process model

The key phases depicted in the image encompass the following:

1. Feasibility study and project estimation: this stage determines whether the project should commence, exploring possible alternatives and necessary resources. It yields a "feasibility study document" comprising an initial problem description, scenarios outlining potential solutions, and cost and schedule estimates for various options.
2. Requirement analysis and specification: in this phase, the domain in which the application operates is meticulously analyzed. Requirements are identified, and software specifications are derived, leading to the creation of a "requirement analysis and specification document".

3. Design: this step outlines the software architecture, defining components, their relationships, and interactions. The objective is to facilitate concurrent development and allocate responsibilities. It culminates in the production of a "design document".
4. Coding and unit test: each module is implemented and rigorously tested. Additional quality assurance measures, such as inspections, may be employed. Programs are accompanied by their respective documentation.
5. Integration and system test: modules are integrated into systems, and the integrated systems undergo testing. This phase, as well as the preceding one, can be integrated into an incremental implementation scheme.
6. Deployment.
7. Maintenance; the maintenance phase comprises various aspects, including:
 - Corrective: addressing and rectifying identified faults or defects.
 - Adaptive: adapting the software to changes in its environment.
 - Perfective: accommodating new or modified user requirements.
 - Preventive: involving activities aimed at enhancing the system's maintainability.




Figure 1.2: Effort in each phase

The primary challenges associated with software evolution include:

- It is seldom anticipated and strategically planned.
- Software is highly amenable to change, with modifications often directly affecting the code, leading to inconsistencies in project documentation.

Effectively addressing software evolution necessitates the implementation of sound engineering practices, comprising two key steps: first, modifying the design, and then making corresponding adjustments in the implementation, ensuring the consistent updating of all associated documents. Indeed, one of the central objectives of software engineering is to create software that can be designed to accommodate future changes in a dependable and cost-effective manner.

The waterfall model operates as a black-box system, as the company seeking the software provides initial requirements and remains relatively uninvolved throughout the development phase. If a higher degree of transparency and customer interaction is required, an alternative development model must be adopted, one that permits regular customer feedback. With each customer interaction, two key aspects can be evaluated:

- Validation: ensuring that the product aligns with the customer's specific requests.
- Verification: confirming that the product functions correctly and in the intended manner.

The concept of a flexible development process is centered on its adaptability to changes, particularly in requirements and specifications. This approach involves incremental processes that allow for feedback at various stages. There are various forms of flexible development processes, such as SCRUM, extreme programming, incremental releases, rapid prototyping, DevOps, and more.

CHAPTER 2

Requirements engineering

2.1 Definition

The primary measure of success of a software system is the degree to which it meets the purpose for which it was intended.

Definition

Software systems *requirements engineering* is the process of discovering that purpose, by identifying stakeholders and their needs, and documenting these in a form that is amenable to analysis, communication, and subsequent implementation.

This phase entails pivotal considerations: stakeholder identification, needs assessment, documentation generation, and the analysis, communication, and implementation of requirements. An alternate definition is as follows:

Definition

Requirements engineering is the branch of software engineering concerned with the real-world goals for, function of, and constraints on software systems. It is also concerned with the relationship of these factors to precise specifications of software engineering behavior, and to their evolution over time and across software families.

2.2 Importance and difficulties

Customer-provided requirements can be categorized into three primary types:

- **Functional requirements:** These delineate the system's interactions with its environment, irrespective of implementation details. They represent the core objectives that the software must achieve.
- **Nonfunctional requirements:** These encompass user-visible aspects of the system that aren't directly tied to functional behavior.
- **Constraints:** these restrictions are either set by the client or are inherent to the system's operational environment.

Nonfunctional requirements, often referred to as quality of service attributes, specify how functionality should be delivered to the end user. While they transcend specific application domains, their relevance and prioritization are influenced by the application's context.

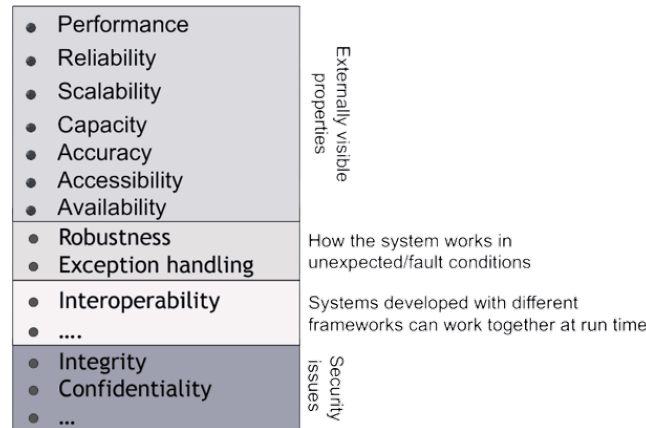


Figure 2.1: Some relevant QoS characteristics

2.3 Requirement engineering process

Inadequate requirements are widespread, and the process of requirement engineering is both challenging and pivotal. This is because issues in the initial stages can potentially escalate costs by a factor of up to two hundred during the final phase.

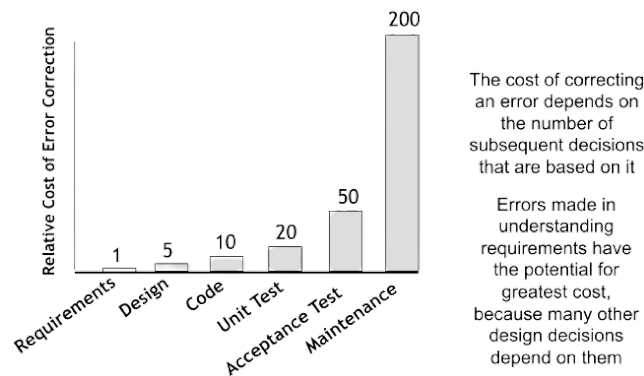


Figure 2.2: Cost of late correction [Boehm, 1981]

The complexity of requirement engineering arises from several factors, including composite systems, the presence of multiple systems, varying abstraction levels, diverse concerns, and involvement of stakeholders with different backgrounds.

Requirement engineers are tasked with various responsibilities, including:

- Eliciting information, which involves gathering details about project objectives, context, scope, domain boundaries, and requirements.
- Modeling and analysis, encompassing the definition of goals, objects, use cases, and scenarios.

- Communicating requirements, which includes providing analysis feedback, documenting in the Requirements Analysis and Specification Document (RASD), and creating system prototypes.
- Negotiating and agreeing upon requirements, which involves resolving conflicts and managing risks, assisting in requirement selection and prioritization.
- Managing and evolving requirements, involves overseeing them across the entire development lifecycle, ensuring backward and forward traceability, and effectively managing changes and their consequences.

2.4 World-machine relationship

The machine signifies the portion of the system to be developed, while the world denotes the segment of the real-world influenced by the machine. It is crucial to note that the machine's purpose always resides within the context of the world.

Requirements engineering is primarily concerned with phenomena occurring in the world, rather than those confined to the machine itself. In essence, requirements models serve as representations of real-world phenomena.

Certain phenomena are shared between the world and the machine. These shared phenomena fall into two categories:

- Phenomena controlled by the world and observed by the machine.
- Phenomena controlled by the machine and observed by the world.




Figure 2.3: Goals, domain assumptions, and requirements

Goals are declarative statements formulated in terms of real-world phenomena, which may or may not be shared. Domain properties and assumptions, on the other hand, are descriptive statements presumed to be valid in the world. Requirements, as prescriptive assertions, are expressed in terms of shared phenomena.

The completeness of requirements, denoted as R , is achieved if:

- R guarantees the satisfaction of goals G within the context of domain properties D , expressed as $R \wedge D \models G$.
- G comprehensively capture all the needs of stakeholders.
- D accurately represent valid properties and assumptions about the world.

2.5 Elicitation of requirements

The complexity in requirement engineering can be managed through various means, including:

- Employing diverse approaches and strategies, and synthesizing the outcomes from each.
- Maintaining proximity to stakeholders.
- Enabling stakeholders to articulate their perspectives.

This scenario can be generalized in terms of:

- Involving participation actors.
- Defining entry conditions.
- Outlining the flow of events.
- Specifying exit conditions.
- Identifying exceptional cases.
- Documenting special requirements.

2.6 Modeling requirements

Definition

A *model* is a depiction, typically in one medium, of an entity in the same or a different medium. It encapsulates the essential facets of the subject while simplifying or omitting extraneous details.

Reality, denoted as R , comprises tangible entities, individuals, processes, and their interconnections. In contrast, a model, symbolized as M , represents an abstraction of these entities, individuals, processes, and the relationships between these abstractions.

To comprehend reality, it necessitates interpretation denoted as I through a mapping function. For a model to be considered effective, it is imperative that the relationships that hold true in reality R also hold true in the model M .




Figure 2.4: Relationship between model M and reality R

Software models serve a variety of purposes, including:

- Capturing and articulating requirements and domain knowledge with precision.
- Facilitating software system design and the generation of practical deliverables.
- Offering a simplified perspective on intricate systems and enabling evaluation and simulation of such complexity.
- Generating prospective system configurations.

The primary modeling challenges revolve around ensuring coherence among different perspectives of the system and addressing the potential for variations in interpretation and ambiguity by establishing clear boundaries for acceptable differences in how the model is understood.

In the domain of requirements engineering, our modeling efforts should primarily focus on:

- Representing the pertinent entities and individuals relevant to the given problem.
- Describing the relevant phenomena.
- Documenting the objectives, requirements, and domain assumptions.

When it comes to choosing the appropriate modeling tools, we have several options at our disposal:

- Natural language (e.g., English, Italian, etc.):
 - Pros: user-friendly and easy to use.
 - Cons: prone to a high level of ambiguity, and it's possible to overlook relevant information.
- Formal language (e.g., FOL, Alloy, Z, etc.):
 - Pros: offers the possibility of employing tools for analysis and validation; this approach compels the user to specify all crucial details.
 - Cons: requires expertise in using the language.
- Semiformal language (e.g., UML):
 - Pros: simpler than a formal language, providing some degree of structure to the models.
 - Cons: not amenable for automated analysis, and still entails some level of ambiguity.
- Mixed approach:
 - Utilizing a semiformal language for foundational aspects.
 - Augmenting semiformal models with explanatory informal text.
 - Leveraging a formal language for the most critical components.

The choice of the modeling tool should align with the specific requirements and characteristics of the project at hand.

2.7 Use cases and requirements

The fundamental steps in formulating use cases encompass the following:

1. Assign a name to the use case.
2. Identify the actors: generalize specific names to encompass participating actors.
3. Focus on the flows of events, entry and exit conditions using natural language.
4. Give due attention to exceptional cases and special requirements.

Each use case can give rise to one or more requirements.

A use case represents a sequence of events within the system, involving interactions with actors. These use cases are initiated by an actor and conclude under specific termination conditions.

Definition

The *use case model* is the set of all use cases specifying the complete functionality of the system.

A *use case association* is a relationship between use cases. The primary types of use case associations include:

- Include, where one use case utilizes another use case.
- Extends, where a use case extends the behavior of another use case.
- Generalization, when an abstract use case has several specialized versions.

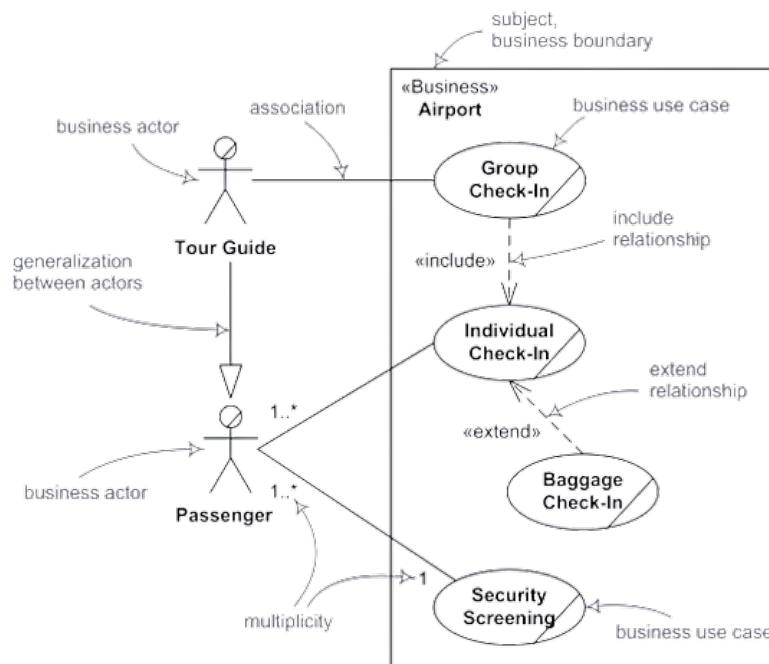


Figure 2.5: Use case model example

2.8 Requirements-level class diagrams

The requirements-level class diagrams are conceptual models for the application domain. They may model objects that will not be represented in the software-to-be. Usually, they do not attach operations to objects: it's best to postpone this kind of decisions until software design.

To find objects and classes we need to:

- Analyze any description of the problem and application domain you may have.
- Analyze your scenarios and use cases descriptions.

Finding objects is the central piece in object modeling. A possible tool to use in the analysis is the Abbott Textual Analysis also called noun-verb analysis: nouns are good candidates for classes and verbs are good candidates for associations and operations.

Example	Grammatical construct	UML component
"Monopoly"	concrete person, thing	object
"toy"	noun	class
"3 years old"	adjective	attribute
"enters"	verb	operation
"depends on"	intransitive verb	operation (event)
"is a", "either", "or", "kind of"	classifying verb	inheritance
"has a", "consists of"	possessive verb	aggregation
"must be", "less than"	modal verb	constraint

Table 2.1: Abbott textual analysis example

2.9 Dynamic modeling

Dynamic modeling serves the purpose of providing methodologies for modeling interactions, the behavior of participants, and workflow within a system. This is achieved through various diagram types, including sequence diagrams, state machine diagrams, and activity diagrams. During the creation of these diagrams, certain objects become apparent.

A sequence diagram is established by following the flow of events outlined in the use case diagram. It represents objects engaged in a use case scenario using a directed acyclic graph notation. The fundamental rules for creating sequence diagrams include:

- Every event involves both a sender and a receiver.
- The representation of the event is sometimes referred to as a message.
- The sender and receiver for each event must be identified.

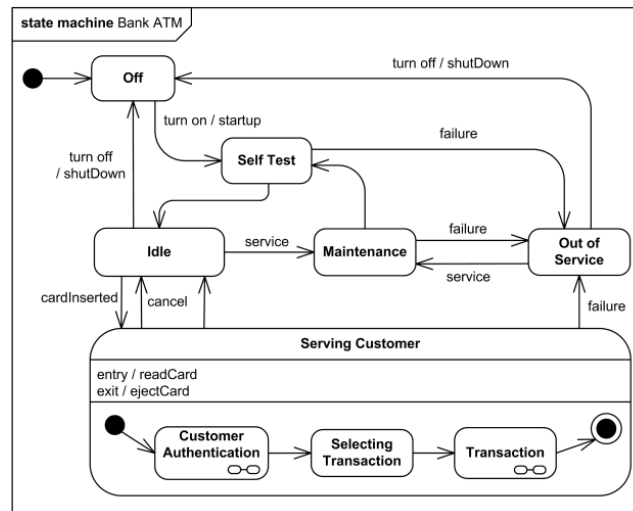


Figure 2.6: Example of a state diagram

For effective dynamic modeling, it is crucial to construct models solely for classes exhibiting substantial dynamic behavior, considering only relevant attributes. Additionally, when deciding on actions and activities, one must account for the granularity of the application and strive to reduce notational complexity.

Alloy

3.1 Introduction

Alloy is a formal notation designed for specifying models of both systems and software. It exhibits characteristics akin to a declarative object-oriented language, underpinned by a robust mathematical foundation. To aid in the practical application of Alloy, a supporting tool is available, facilitating the simulation of specifications and enabling property verification.

Alloy has been crafted with the objective of combining the expressive power found in the Z language with the rigorous automated analysis capabilities offered by the SMW model checker. In the realm of requirement engineering, Alloy finds utility in formally describing the domain, its inherent properties, or the operations that a machine is required to provide. In software design, Alloy serves as a valuable tool for the formal modeling of components and their inter-dependencies.

Alloy represents a fusion of first-order logic and relational calculus. Key features of this formal language include:

- A thoughtfully curated subset of relational algebra, offering a uniform model for individuals, sets, and relations, with an absence of high-order relations.
- Minimal reliance on arithmetic operations.
- Support for modules and hierarchies to enhance organization and scalability.
- Designed for succinct, illustrative specifications.
- Employs a potent and efficient analysis tool.

3.2 Syntax

Alloy presents bounded snapshots of the world that satisfy the given specification.

It employs bounded exhaustive search to uncover counterexamples to asserted properties using SAT.

Definition

Atoms represent Alloy's fundamental entities, characterized by indivisibility, immutability, and a lack of interpretation.

Relations establish connections between atoms, forming sets of tuples, where tuples are sequences of atoms.

In Alloy, relations are inherently typed, with their types determined by the declaration of the relation. The fundamental Alloy relation types include:

- "none" (signifying an empty set).
- "univ" (representing the universal set).
- "iden" (denoting the identity relation).

The logical operators in Alloy encompass:

- Union " \cup ".
- Intersection "&".
- Difference "-".
- Subset "*in*".
- Equality "=".
- Cross product " \rightarrow ", analogous to a natural join.
- Dot join "." or "[]", where the final element of the first relation joins with the corresponding first elements of the second relation, followed by the removal of the combined element from the relation.

Here are the various binary closures applicable to relations in Alloy:

- Transpose " \sim ": inverts the order of the elements in the relation.
- Transitive " \wedge ": it signifies the transitive closure where $\wedge r = r + r.r + r.r.r + \dots$
- Reflexive transitive " $*$ ": it represents the reflexive transitive closure, where $*r = \text{iden} + \wedge r$.

The possible restrictions are:

- Domain restriction "<:", that restricts the elements on the left side to the set on the right side.
- Range restriction ">:", that is same as before, but the relations are inverted.
- Override "++", that removes the tuples on the left that are in the right relations and adds all the remaining relations of the right relation.

Alloy also includes various Boolean operators:

- Negation ("!" or "not").
- Conjunction ("&" or "and").

- Disjunction ("||" or "or").
- Implication (" \implies " or "implies").
- Alternative ("," or "else").
- Bi-Implication (" \iff " or "iff").

Alloy offers logic quantifiers:

- "all": holds for every element.
- "some": holds for at least one element.
- "no": holds for no elements.
- "lone": holds for at most one element.
- "one": holds for exactly one element.

For defining relations with singletons, you can use the following declaration " $x: m\ e$ ", where " x " is the name of the relation, " m " is the multiplicity of the element (e.g., "set", "one", "lone", or "some") and " e " is the name of the element within the relation. When the relation consists of pairs, the declaration appears as " $r: A\ m \rightarrow B\ n$ ", where " r " is the relation name, " A " and " B " denote the element names with multiplicities " m " and " n ", respectively.

Additionally, Alloy includes operators like " $\#$ " (counting the number of tuples in r), integers (0, 1, ...) for defining variable values, arithmetic operators (" $+$ " and " $-$ "), and various comparison operators (" $<$ ", " \leq ", " $=$ ", " \geq ", " $>$ "). There is also the "sum" operator, which adds all elements within a selected tuple.

Other useful keywords are:

- "let:" this keyword is utilized to establish local variables within a formula or expression. These variables serve to store and reuse intermediate results, making it easier to simplify complex expressions.
- "enum": this keyword is employed to define enumerations. Enumerations enable the definition of a finite set of symbolic values, which can be utilized in models to represent various states, options, or categories.
- "var": this keyword is used for declaring variables within an Alloy model or specification.
- "after": this keyword is employed to specify the temporal ordering or sequence of events or states within a model. It is commonly used in temporal logic to define the sequence of events that must occur before or after a specific state or action.
- "always": this keyword is used to express a property that must remain true throughout a system's execution or under certain conditions. It is often utilized to specify invariant properties in Alloy models.
- "eventually": this keyword is used to express temporal logic properties that indicate a particular condition or event will ultimately occur during the execution of a system. It is commonly used to model and verify properties that describe what should happen at some point in the future.

- "historically": keyword is used to express temporal logic properties that describe a condition or event that has remained true for a continuous duration of time leading up to the present or to a specified point in the model's execution. It is often employed to specify properties related to the historical behavior of a system.
- "before".
- "once".

3.3 Address book example

Let's consider the task of creating a basic address book model. This address book contains a collection of addresses associated with their respective names. In this scenario, there are three primary entities:

- "Name": represents names of individuals.
- "Addr": represents addresses.
- "Book": signifies the context of the address book.

The relationship here is established through the entity "Addr", which links "Name" to "Addr" within the context of "Book". We use the "lone" keyword to indicate that each "Name" can be linked to at most one "Addr". To summarize the preceding descriptions:

```
sig Name {}
sig Addr {}
sig Book {
  addr: Name -> lone Addr
}
```

This specification creates the following relationships:

- Sets are unary relations.
- Scalars are sets consisting of a single element.
- The ternary relation involves these three predicates.

A new predicate can be defined using the "pred" keyword.

```
pred show {}
run show for 3 but exactly 1 Book
```

In the second line, it specifies that we should find a maximum of three elements for each "Book". The previously defined predicate "show" is currently empty and always returns true. We can now create a predicate with specific arguments. For instance:

```
pred show [b:Book]{
  # b.addr > 1
}
run show for 3 but exactly 1 Book
```

In the previous example, the consistent predicate imposes a restriction on the quantity of "Address" relations within a given "Book". In the following example, the consistent predicate enforces a constraint on the number of distinct "Address" entries that appear within the "Book"

```
pred show [b:Book]{
  # b.addr > 1
  # Name.(b.addr) > 1
}
run show for 3 but exactly 1 Book
```

In the subsequent example, the inconsistent predicate utilizes the "some" keyword, indicating the existence of an element. In this scenario, there is only one "Book" so the tool will report that no instances can be identified.

```
pred add [b: Book]{
  # b.addr > 1
  some n: Name | # n.(b.addr) > 1
}
run show for 3 but exactly 1 Book
```

All the preceding predicates are considered static as they do not modify the signature. In Alloy, dynamic predicates are used for dynamic analysis. For instance, we can create a predicate that introduces an "Address" and a "Name" to a "Book" as follows:

```
pred add [b, b': Book, n: Name, a: Addr]{
  b'.addr = b.addr + n -> a
}
pred showAdd [b, b': Book, n: Name, a: Addr]{
  add[b, b', n, a]
  #Name.(b'.addr) > 1
}
run showAdd
```

We can now define a predicate for the "Book" deletions.

```
pred del [b, b': Book, n: Name]{
  b'.addr = b.addr - n -> Addr
}
```

We can verify whether performing a delete operation after an add operation restores us to the initial state or not by employing an "assertion".

```
assert delRevertsAdd{
  all b1, b2, b3: Book, n: Name, a: Addr
  add[b1, b2, n, a] and del[b2, b3, n]
  implies b1.addr = b3.addr
}
```

During the assertion verification process, Alloy looks for counterexamples. In this instance, we will encounter a counterexample, causing the assertion to evaluate as false. To rectify the assertion, we should amend it as follows:

```
assert delUndoesAdd{
  all b1, b2, b3: Book, n: Name, a: Addr |
  no n.(b1.addr) and add[b1, b2, n, a] and del[b1, b2, n]
  implies b1.addr = b3.addr
}
```

In some cases, we may need to retrieve specific signatures. To accomplish this, we can leverage Alloy functions. For instance, we can define a function that searches for a particular "Book" and returns a set of "Address":

```
fun lookup[b: Book, n: Name]: set Addr{
  n.(n.addr)
}
```

3.4 Family relations example

Now, let's explore a family relationship tree. To begin, we must define a generic individual, which can be either male or female.

```
abstract sig Person {
  father: lone Man
  mother: lone Woman
}
sig Man extends Person {
```

```

    wife: lone Woman
}
sig Woman extends Person {
    husband: lone Man
}

```

We establish that each "Person" can have at most one father and one mother (indicated by the keyword "lone"), as we require a root for the family tree. The person at the root should have no parents, possibly due to them being unknown. The "Person" signature is declared as "abstract" because it needs to be specialized into one of the subsequent signatures, either "Man" or "Woman". Signatures, denoted by the "sig" keyword, represent sets of elements. Before this keyword, we can specify the number of entities required ("lone", "one", or "some").

Definition

The *fields* of a signature are relations whose domain is a subset of the signature. The keyword *extends* is used to declare a subset of signature.

To obtain the set of grandfathers of a given individual, we can define a function like this:

```

fun grandpas[p: Person]: set Person {
    p.(mother+father).father
}
pred ownGrandpa[p: Person] {
    p in p.grandpas[p]
}

```

We've additionally created a predicate that verifies if a person is within the set of grandfathers returned by the "grandpas" function. However, the issue at hand is that we haven't imposed constraints on relations. To address this, we need to define two new operators for binary relations:

- Transitive closure: $\hat{r} = r + r.r + r.r.r + \dots$
- Reflex transitive closure: $*r = iden + \hat{r}$

We can now specify that no one can be their own father or mother:

```

fact {
    no p: Person | p in p.^(mother+father)
}

```

We must also establish a constraint that if X is the husband of Y, then Y is the wife of X:

```

fact {
    all m: Man, w: Woman | m.wife=w iff w.husband=m
}
fact {
    wife = ~husband
}

```

The two statements are equivalent, with the second being expressed using the transpose operator. A fact can encompass multiple constraints, meaning that the previous constraints can be consolidated into a single fact. It's worth noting that facts are global, while predicates need to be invoked explicitly.

Requirement analysis and specification

4.1 Structure of a RASD document

The Requirements and Specifications Document (RASD) serves several purposes:

- Communication: it conveys an understanding of the requirements, encompassing the application domain and the system under development.
- Contractual: it can be legally binding, serving as a formal agreement between stakeholders.
- Baseline for project planning and estimation: it provides a foundation for project planning and estimation, covering aspects like size, cost, and schedule.
- Baseline for software evaluation: it supports system testing, verification, and validation activities. It contains the information necessary to verify if the delivered system aligns with the requirements.
- Baseline for change control: it establishes a foundation for managing changes in requirements as the software evolves.

The RASD document is utilized by various stakeholders, including:

- Customers and Users: they are interested in a high-level description of system functionalities and requirements.
- System analyst and requirement analysts: these individuals use the RASD to specify how the system interacts with other systems.
- Developers and programmers: they refer to the RASD for implementation details.
- Testers: they use the RASD to check if the system meets its requirements.
- Project managers: they rely on the RASD to control the development process.

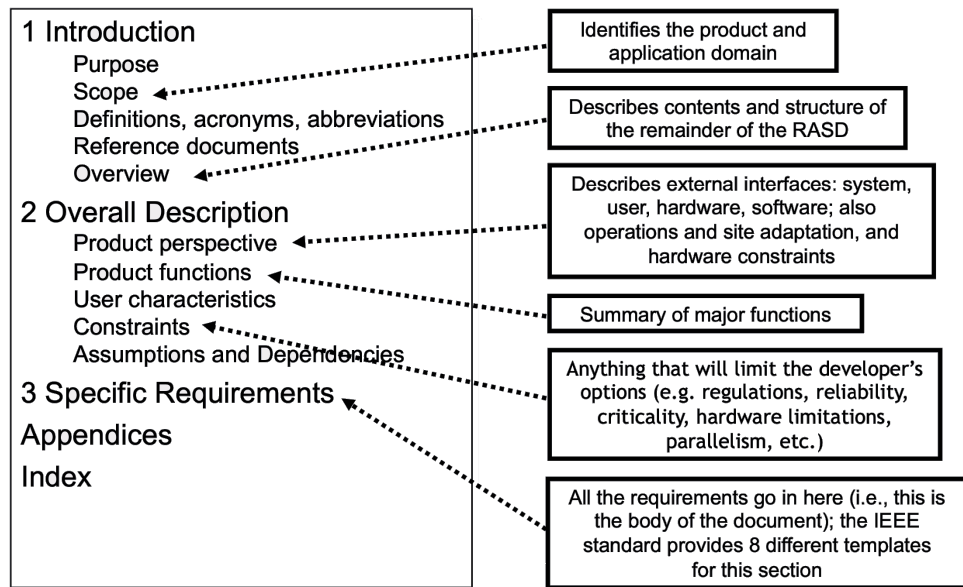


Figure 4.1: IEEE standard for RASD

A well-crafted RASD should possess the following qualities:

- **Completeness:**
 - Regarding goals: all requirements must satisfy the goals within specified domain assumptions.
 - Regarding inputs: software behavior should be specified for all possible inputs.
 - Structural completeness.
- **Pertinence:**
 - Each requirement or domain assumption should be necessary for achieving a goal.
 - Each goal should be genuinely needed by the stakeholders.
 - The RASD should not contain items unrelated to requirement definitions.
- **Consistency:** there should be no contradictions in the formulation of goals, requirements, and assumptions.
- **Unambiguity:**
 - Clear and well-defined vocabulary.
 - Unambiguous assertions.
 - Verifiability of requirements.
 - Clear delineation of responsibilities between the software and its environment.
- **Feasibility:** the goals and requirements must be achievable within the allocated budget and schedules.
- **Comprehensibility:** the RASD should be easily understandable by the target audience.
- **Good structuring:** every item must be defined before it is used.

- Modifiability: the document should be adaptable, and the impact of modifications should be assessable.
- Traceability:
 - Indication of the sources of goals, requirements, and assumptions.
 - Linking requirements and assumptions to underlying goals.
 - Facilitating referencing of requirements in future documentation.

This structured RASD document ensures effective communication and clarity while serving as a contractual and foundational reference throughout the software development process.

Software design

5.1 Software architecture

Definition

The *architecture* of a software system defines the system in terms of computational components and the interactions among these components.

In the context of a software system, the *software architecture* represents the set of structures necessary to reason about the system. These structures include software elements, their relationships, and properties associated with both.

There are three fundamental structures relevant to software architecture:

- Component-and-connector structures: these structures describe how the system is organized as a set of elements with runtime behavior (components) and their interactions (connectors). They allow us to study runtime properties, such as availability and performance, as well as how these structures collaborate.

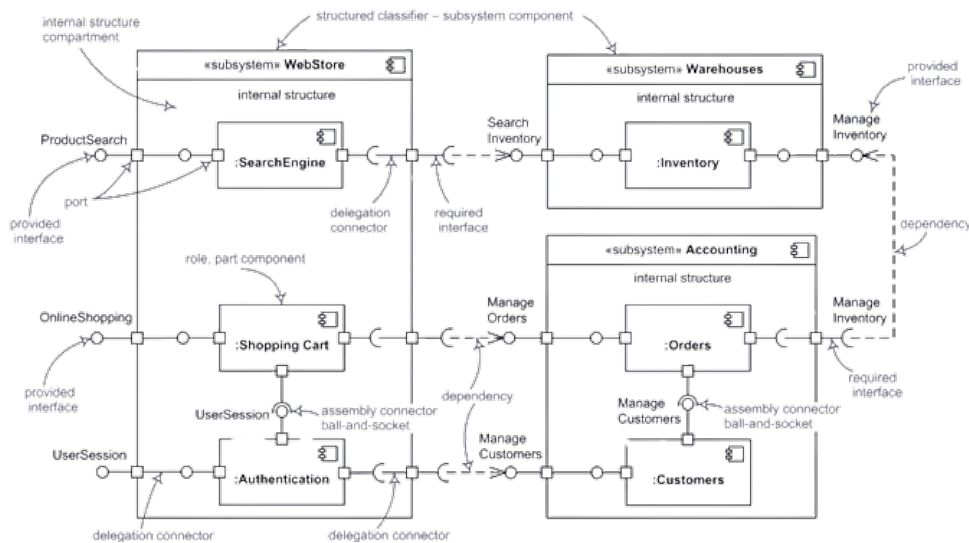


Figure 5.1: UML component diagrams for component and connector structure




Figure 5.2: UML sequence diagrams for component and connector structure

- Module structures: illustrate how a system is structured as a set of code or data units that need to be procured or constructed, along with their relationships. Modules serve as implementation units and form the basis for work division. Typical relationships among these modules include "uses," "is-a" (generalization), and "is-part-of."

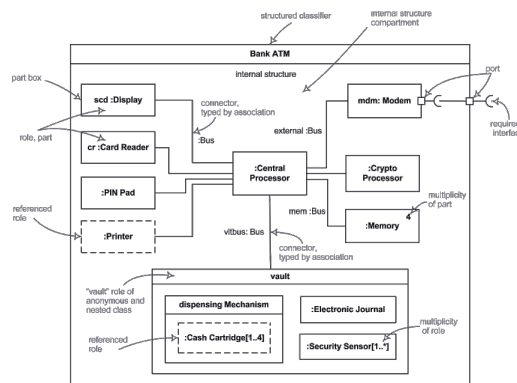


Figure 5.3: UML composite structure diagram for module structure

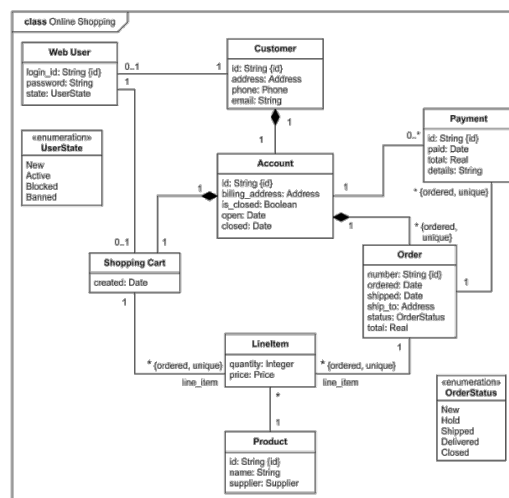


Figure 5.4: UML class diagram for module structure

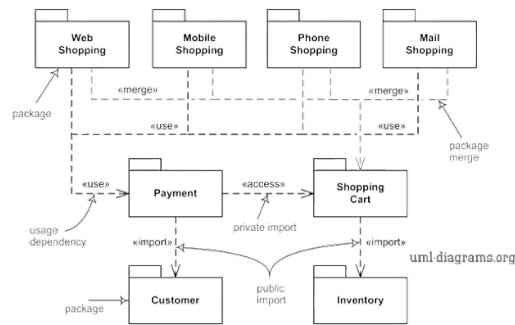


Figure 5.5: UML package diagram for module structure

- Allocation structures: these structures define how the elements from component and connector or module structures map to entities that are not software. Common allocation structures include deployment structure, implementation structure, and work assignment structure. The deployment structure captures the system's hardware topology, specifying the distribution of components and identifying performance bottlenecks.

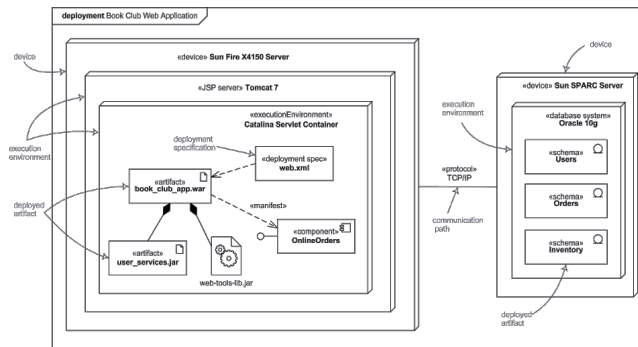


Figure 5.6: UML deployment diagrams and deployment structure




Figure 5.7: UML deployment diagrams and deployment structure

Summary

Structures	Elements	Relations	Useful for
Component diagrams	Components offering services	Provided and required interfaces	Performance analysis Robustness analysis Resource allocation Project planning
Sequence diagrams	Processes Threads	Synchronous and asynchronous messages	Analysis of resource contention Parallelism opportunities

Table 5.1: Component and connectors diagrams and usage

Structures	Elements	Relations	Useful for
Composite structures Package diagrams	Modules Packages	Is a submodule of Uses	Resource allocation Project planning Encapsulation
Class diagrams	Classes	Is-a Is part of	Object-oriented development Planning for extensions
Layered structures	-	Can use Provides abstraction	Incremental development
Data model	Data entities	One-to-one One-to-many Many-to-one Many-to-many Is-a	Engineering global data structures for consistency and performance

Table 5.2: Modular structures diagrams and usage

Structures	Elements	Relations	Useful for
Deployment diagrams	Components Hardware/software execution environment	Allocated to	Performance Security Robustness analysis
Implementation structures	Modules File structures	Stored in	Configuration control Integration Test activities
Work assignment	Modules Organizational units	Assigned to	Project management Development efficiency

Table 5.3: Allocation structures diagrams and usage

5.2 Software design description and principles

The IEEE has established two standards pertaining to architectural models:

- IEEE standard for information technology: software design descriptions.

- IEEE standard for systems and software engineering: architecture description (manner in which architectural descriptions of systems are organized and expressed).

In accordance with IEEE standards, a Software Design Description (SDD) should encompass the following elements:

- Identification of the SDD, which includes information such as the date, authors, and affiliated organization.
- Description of design stakeholders.
- Description of design concerns.
- Selection of design viewpoints.
- Design views.
- Design overlays.
- Design rationale.

Additionally, there are eleven key design principles that should be considered in the development of software systems:

1. Employ the divide and conquer approach: break down complex problems into smaller, more manageable components.
2. Maintain a high level of abstraction in designs to minimize consideration of unnecessary details and reduce complexity.
3. Maximize cohesion where possible to ensure logical and functional consistency within components.
4. Minimize coupling to reduce interdependencies between system elements.
5. Design for reusability, allowing various aspects of the system to be used in different contexts.
6. Promote the reuse of existing designs and code, complementing the design for reusability.
7. Design for flexibility by anticipating and preparing for future changes in system requirements.
8. Anticipate obsolescence by planning for technological or environmental changes to ensure software continuity or ease of adaptation.
9. Design for portability, aiming to make the software compatible with a wide range of platforms.
10. Design for testability, implementing measures that facilitate testing processes.
11. Design defensively, exercising caution when making assumptions about how others will interact with the components you are designing.

5.3 Architectural styles

Definition

An *architectural style* determines the vocabulary of components and connectors that can be used in instances of that style, together with a set of constraints on how they can be combined. These can include topological constraints on architectural descriptions (e.g., no cycles). Other constraints—say, having to do with execution semantics—might also be part of the style definition.

Client-server

The client-server architectural style is primarily employed in distributed applications, where the client initiates requests, and the server delivers responses.


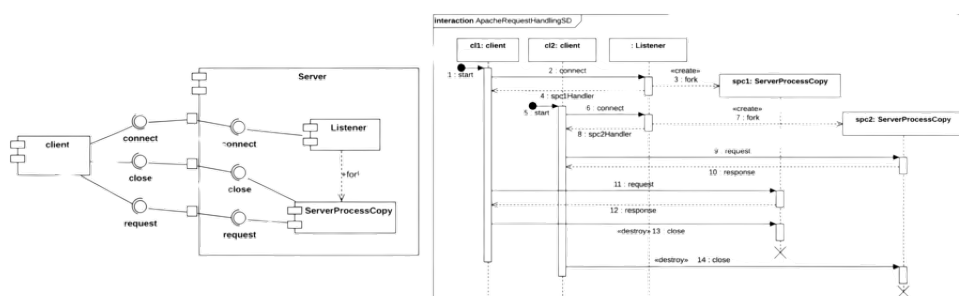


Figure 5.8: Client-server architecture general schema


This approach proves beneficial in situations where multiple users require access to a common resource, when there is a need to remotely access existing software, and when it is advantageous to structure the system around a shared functionality that multiple components utilize. Key technical considerations for this architectural style include the design and documentation of well-defined interfaces for the server and the need to ensure the server's capability to handle concurrent requests efficiently.

In this architectural context, the requirement is to receive and manage requests from multiple clients. Various solutions can be considered to address this challenge:

- Forking (Apache Web Server):
 - Approach: create one process per request or per client.
 - Strengths: simplicity, isolation, and protection provided by the one connection per process model. Effective for up to 2000 users.
 - Issues: limited capacity, uncertainty regarding the number of active processes at a given time, and scalability challenges due to the overhead of fork-kill operations.



- Worker pooling (NGINX Web Server):
 - Approach: specifically designed for high concurrency, addressing scalability issues inherent in the forking model by introducing a new architectural tactic¹.
 - Architectural Tactic: NGINX prioritizes scalability and performance at the expense of availability.
 - Strengths: fixed number of workers for each pool, each worker equipped with a queue (with dropped requests for performance if the queue is full), and a dispatcher for load balancing among workers.



Three-tier architecture

Within three-tier architectures, the intermediate tier is responsible for handling the application logic. The primary advantage lies in the separation of logic from both data and presentation elements. The modular structure allows for enhanced flexibility and maintainability.


With the option of incorporating additional tiers, it becomes feasible to establish an N-tier architecture. This extended architecture enables further refinement and specialization of components, supporting more intricate and scalable system designs.

Event-driven architecture

In an event-driven architecture, components can register to both send and receive events that are broadcasted to all registered components, where neither the sender nor the receiver is known. This model, often referred to as publish-subscribe, facilitates easy addition and deletion of components and is increasingly employed in modern integration strategies. Key characteristics include asynchronous events, reactive computation, message destination determined by the receiver, loose coupling, and flexible communication means. Strengths and challenges of this architecture include its widespread adoption in modern development practices, ease of adding and removing components, potential scalability issues, and complexities in event ordering.

Apache Kafka serves as a framework for the event-driven paradigm, providing primitives for creating event producers and consumers, along with a runtime infrastructure for handling event transfer. Kafka stores events durably and reliably, allowing consumers to process events in real-time or retrospectively. These services are delivered in a distributed, highly scalable, elastic, fault-tolerant, and secure manner.

¹Design decisions that influence the control of one or more quality attributes.



The Kafka architecture involves brokers managing designated topics and partitions, with each partition containing sets of messages related to the respective topic. The partitions operate independently and can be duplicated across multiple brokers to ensure fault tolerance.

Each partition designates one broker as the leader, while the remaining brokers hosting the same partition act as followers. Producers are aware of the leading brokers and direct their messages to them. Messages within a given topic are grouped into batches on the producer's end and transmitted to the broker once the batch size surpasses a specified threshold.

Consumers employ a pull mechanism, retrieving an entire batch of messages associated with a specific partition from a defined offset. Messages persist at the broker level for a predetermined duration and can be accessed multiple times during this period.

The leader broker keeps tabs on the in-sync followers. The correct functioning of the cluster is overseen by ZooKeeper, with all brokers regularly sending heartbeats to it. In the event of a broker failure, ZooKeeper intervenes by appointing a new leader for all partitions previously led by the failing broker. Additionally, ZooKeeper has the capability to initiate or restart brokers as needed.

Brokers finalize message commitment by storing them in their designated partition. The leader includes the message in the followers (replicas) if they are accessible. A potential challenge arises in the event of failure, where the producer may not receive a response, as illustrated by message seven. Consequently, the producer is required to resend the message.

Kafka brokers possess the ability to recognize and eliminate duplicate messages. The synchronization process with replicas can be transactional, allowing for precise control over the message flow. Achieving exactly-once semantics is feasible but comes with the trade-off of longer waiting times. Alternatively, opting for at-least-once semantics involves excluding duplicate management, and at-most-once semantics can be chosen by dispatching messages asynchronously.

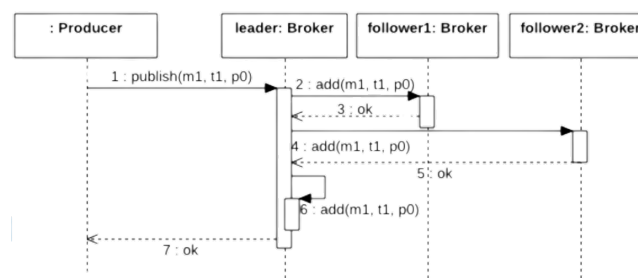


Figure 5.9: Example of a producer

Every consumer has access to a persistent log to maintain the offset, ensuring it is preserved in the event of a failure. If a consumer encounters a failure after processing messages but before recording the new offset in the log, it will retrieve the same messages again, resulting in at-least-once semantics.

The delivery semantics can be altered by storing the new offset before processing. In this case, if a failure occurs after storing the offset, the impact of the received messages does not materialize, leading to at-most-once semantics.

Transactional management of the log introduces the capability for exactly-once semantics, providing a robust mechanism for ensuring that each message is processed only once, even in the face of failures or system disruptions.

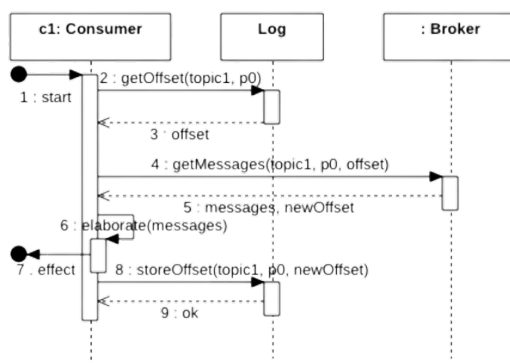


Figure 5.10: Example of a consumer

Apache Kafka's architectural tactics include scalability through multiple partitions and brokers and fault tolerance through persistent partition storage and data replication.

Microservices

In the era preceding microservices, monolithic systems delivered applications as single deployable software artifacts. The microservice architectural style emerged as an approach to developing applications by decomposing monolithic systems into small, specialized services, each running independently in its own process. These services communicate using lightweight mechanisms, often employing HTTP resource APIs.

Microservices focus on addressing a single bounded context within the target domain, offering advantages such as fine-grained scaling strategies, reduced impact of localized issues, improved resilience through decentralized functionality, and enhanced reuse and composability.

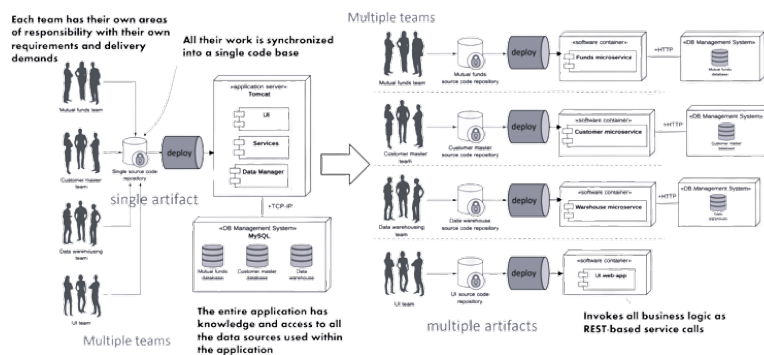


Figure 5.11: Typical microservices architecture process

The process of creating a microservices architecture involves reducing teams' synchronization overhead, organizing small development teams with well-defined responsibilities, and employing technology-neutral protocols like REST APIs for communication, resulting in smaller codebases for easier debugging and more cost-effective maintenance.

The architecture's main elements include a REST API that exposes core service operations, the application or business logic responsible for implementing core operations, and local data storage for each microservice, emphasizing limited data sharing and avoiding global databases.

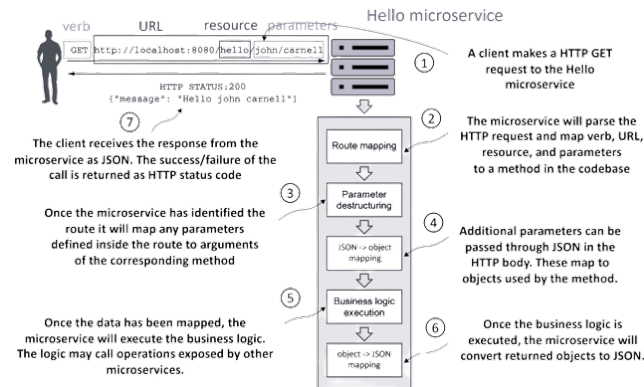


Figure 5.12: Microservices architecture's workflow

To achieve location transparency and scalability in microservices, routing patterns are essential. Service discovery, a critical element, must be highly available, load-balanced, resilient, and fault-tolerant. Resilience patterns, such as the circuit breaker pattern, aim to prevent resource waste and ripple effects by actively monitoring calls to distant services and intervening in case of potential failures.




Figure 5.13: Routing pattern

To allow the microservices to be resilient we need resiliency patterns. The main goal of these patterns are to avoid useless resource consumption and ripple effects. Circuit breaker is a client-side resiliency pattern. The circuit breaker functions as a proxy for a distant service. As it invokes the remote service, the CB actively monitors the call for potential failures, which may include receiving a 5xx error or the call taking an excessively long duration, prompting the CB to terminate the call. In the event of "too many" failures, the circuit breaker intervenes by inhibiting future calls to prevent further issues.

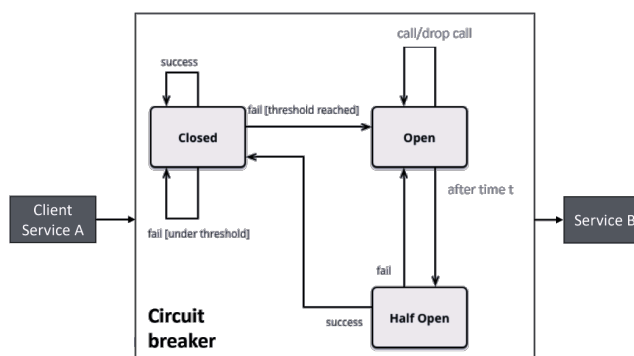


Figure 5.14: Circuit breaker pattern

Security patterns in microservices involve introducing a service or API gateway that acts as a mediator between service clients and invoked services. While this gateway simplifies the implementation of authentication and authorization mechanisms, it can become a single point of failure. This challenge is mitigated by deploying multiple instances of the gateway with a server-side load balancer.

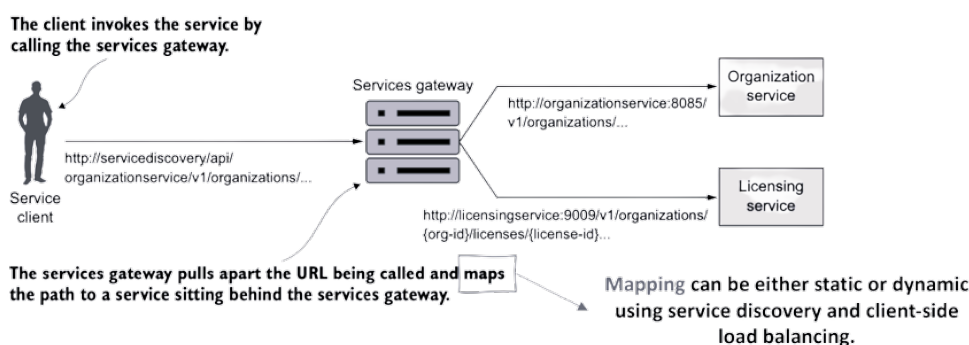


Figure 5.15: Security pattern

Communication patterns play a crucial role in achieving loose coupling in microservices. Despite their strengths in providing flexibility, scalability, and availability, they introduce higher complexity. Event-driven frameworks, supporting multiple communication styles such as notification, request/response, and publish/subscribe, are employed to decouple different parts of the system.




Figure 5.16: Security pattern

5.4 Interface design and documentation

Definition

An *interface* is a boundary where components interact.

The precise definition of interfaces carries significant architectural implications as it directly affects maintainability, usability, testability, performance, and integrability. Two fundamental guiding principles are information hiding and low coupling. The key aspects to consider when designing interfaces are as follows:

- Contract principle: any addition of a resource (operation or data) to an interface signifies a commitment to maintaining it.
- Least surprise principle: interfaces should behave consistently with expectations.
- Small interfaces principle: interfaces should expose the minimum necessary resources.

When designing interfaces, we have also to define:

- Interaction style:
 - Socket: after connection establishment, communication is bidirectional, and both parties must agree on the same protocol.
 - RPC/RMI: resembles a procedure/method call in a centralized setting. It requires stubs and skeletons to transform procedure/method calls into messages and vice versa.
 - REST (REpresentational State Transfer): this is a standardized architectural style for Application Programming Interfaces (APIs) that emphasizes clear separation between distributed, heterogeneous systems and their components.
- Representation and structure of exchanged data: the choice of data representation impacts expressiveness, interoperability, performance, and transparency. Common representations include XML (verbose, requires multiple parsing passes), JSON (more compact than XML, single-pass parsing), and Protocol Buffers (most compact, data passed as binary).
- Error handling: possible reactions to errors include raising an exception, returning an error code, or logging the problem.

A server can offer multiple interfaces concurrently, enabling the separation of concerns, different access rights levels, and support for interface evolution. Interfaces constitute the contract between servers and clients, and sometimes interfaces need to evolve. The strategies used to support this continuity are:

- Deprecation: declare in advance that an interface version will be eliminated by a certain date.
- Versioning: maintain multiple active versions of the interface.
- Extension: a new version extends the previous one.

Representational state transfer

REST APIs are simple and standardized, alleviating developers from concerns about communication protocols (HTTP), data formatting (JSON), and request/response encoding. REST is stateless, which means it doesn't retain states across servers and clients, making it lightweight, scalable, and supportive of caching for high performance.

The available CRUD requests are:

- Create (HTTP POST): used to create a new resource.
- Read (HTTP GET): used to retrieve an existing resource.
- Update (HTTP PUT): used to modify a resource.
- Delete (HTTP DELETE): used to remove a resource.

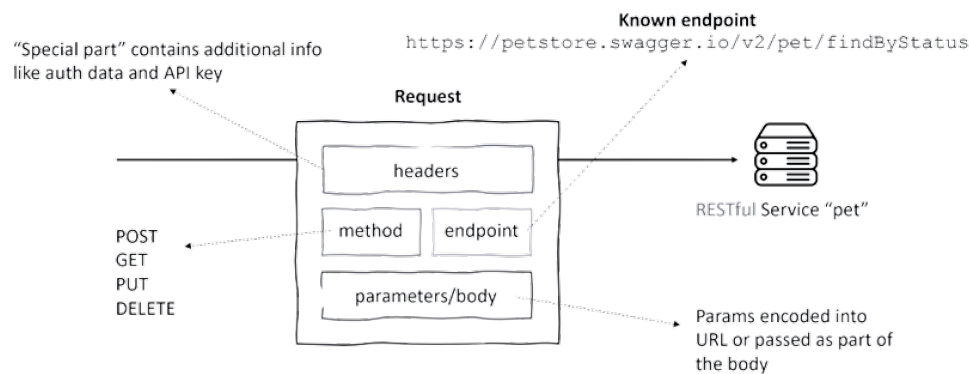


Figure 5.17: Structure of a request




Figure 5.18: Structure of a response

Interface documentation

Interface documentation explains how to use an interface without revealing details of the component's internals. The audience for interface documentation includes developers and maintainers of both the offering and using components, quality assurance teams for system integration and testing, and software architects seeking reusable components.

To document a REST API, the OpenAPI specification is used. It describes a REST API interface through an OpenAPI definition, typically in JSON or YAML format. Benefits of this

type of documentation include its standardized format, public accessibility, and suitability for both human understanding and machine automation tasks like testing and code generation.

The OpenAPI definition covers endpoints, resources, operations, parameters (including data types), and authentication/authorization mechanisms. Support tools for OpenAPI documentation include API validators, documentation generators, and SDK generators for automated client library creation in various programming languages.