

Hardware Architectures For Embedded And Edge AI

Christian Rossi

Academic Year 2024-2025

Abstract

The goal of this course is to provide a comprehensive understanding of embedded and edge Artificial Intelligence by exploring hardware, software, and algorithmic solutions that support the training and inference of machine and deep learning models on embedded systems and edge computing networks, which are constrained by computation, memory, and energy limitations.

The course will begin with an overview of IoT and Smart-X technologies that enable intelligent, autonomous, and distributed embedded and edge applications. It will then dive into a detailed examination of hardware solutions specifically designed for embedded and edge AI. Students will also gain an in-depth understanding of the frameworks and toolchains, such as Tiny Machine Learning, that support the training and inference of machine and deep learning models on these systems. Additionally, the course will cover algorithmic solutions for embedded and edge AI, including techniques like approximate computing and federated learning. A key part of the course will be looking at system-level codesign for machine and deep learning applications.

Contents

1	Introduction	1
1.1	Introduction	1
1.2	Computing continuum	2
1.2.1	Datacenters	2
1.2.2	Edge computing systems	2
1.2.3	Embedded systems	3
1.2.4	Internet of Things	3
2	Hardware	4
2.1	Introduction	4
2.2	Architecture	5
2.3	Sensors and signals	5
2.3.1	Data	6
2.3.2	Sensors	7
2.4	Microprocessor	9
2.4.1	Microcontroller	9
2.4.2	System on chip	10
2.4.3	Comparison	10
3	Software	11
3.1	Introduction	11
3.2	Data preprocessing and feature extraction	12
3.2.1	Data segmentation	12
3.2.2	Data processing	13
3.2.3	Feature extraction	13
3.3	Machine and Deep Learning algorithms	14
3.3.1	Tiny Machine Learning	15
4	Tiny Machine Learning	16
4.1	Machine Learning	16
4.1.1	Supervised Learning	16
4.1.2	Training	17
4.1.3	Validation	18
4.1.4	Neural Networks	19
4.2	Deep Learning	20
4.2.1	Convolutional layer	20
4.2.2	Fully connected layer	21

4.2.3	Memory occupation	22
4.2.4	Tiny Deep Learning	22
4.3	Tiny Deep Learning architectures	23
4.3.1	SqueezeNet	23
4.3.2	MobileNet	24
4.3.3	EfficientNet	26

Introduction

1.1 Introduction

Artificial Intelligence (AI) is a field of computer science focused on developing hardware and software systems capable of performing tasks that typically require human intelligence. These systems can autonomously pursue specific goals by making decisions that were traditionally made by humans.

A key distinction in AI-driven systems lies between smart objects and connected objects. While connected objects primarily send and receive data from the cloud, smart objects analyze data locally, enabling faster decision-making and reducing reliance on constant connectivity.

The definition of AI evolves rapidly, to the point that what was considered AI a decade ago may differ significantly from today's understanding.

AI hardware and software can be categorized similarly to traditional computing environments but are specifically designed to handle AI workloads. In this context, the development environment is often referred to as a framework, platform, or tool, rather than just a conventional programming environment.

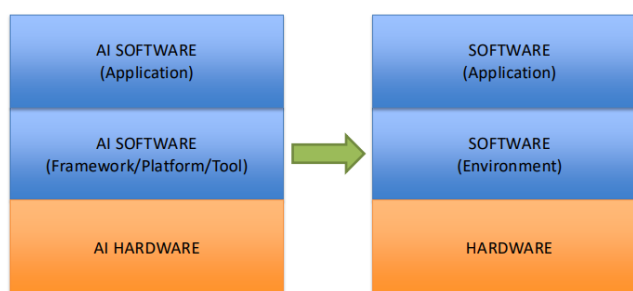


Figure 1.1: Artificial Intelligence stack

The AI stack consists of three main layers:

- *AI software* (application): AI-powered applications running within an IT system.
- *AI software* (framework, platform and tool): programs and libraries that manage physical resources and provide the necessary tools for building AI applications.

- *AI hardware*: the infrastructure supporting AI computation, including data centers, edge computing devices, IoT systems, and specialized processors like CPUs, GPUs, and TPUs.

1.2 Computing continuum

AI hardware ranges from small, low-power devices running on batteries to large-scale, high-performance systems in datacenters. This range represents the computing continuum:

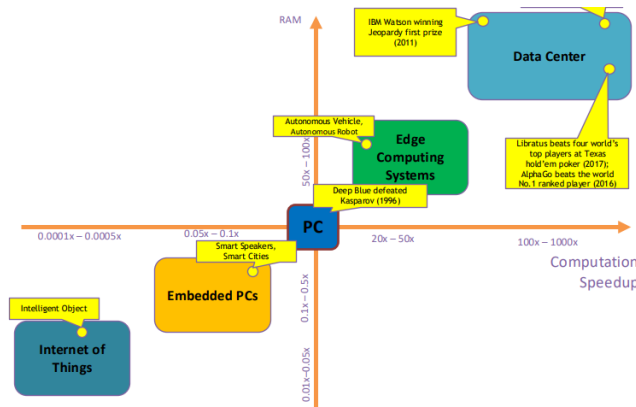


Figure 1.2: Computing continuum

1.2.1 Datacenters

Datacenters sit at the upper end of this spectrum, offering immense computational power for AI workloads. They provide cost-efficient IT infrastructure, high-performance computing capabilities, and instant software updates. Their vast storage capacity ensures data reliability and accessibility, allowing seamless collaboration across devices and locations. Furthermore, by decoupling AI processing from end-user devices, datacenters enable powerful AI applications that are not limited by local hardware constraints.

Datacenters require a constant internet connection, making them less viable in low-bandwidth environments. Their reliance on shared infrastructure can introduce privacy and security concerns, while the lack of direct hardware control may limit customization options. Additionally, the high energy consumption of large-scale AI operations raises both environmental and cost concerns. In latency-sensitive applications, delays in data transmission and processing can further impact real-time decision-making.

1.2.2 Edge computing systems

Edge computing delivers high computational power with the advantage of distributed processing. By bringing computation closer to where data is generated, it enhances privacy and security while significantly reducing latency in decision-making. However, these systems depend on a stable power supply and often integrate with cloud services to extend their processing capabilities.

By processing data locally, edge computing minimizes the need for constant data transmission, optimizing bandwidth and improving energy efficiency. This approach not only strengthens security and privacy but also enables real-time decision-making and adaptive learning across distributed networks. However, edge devices often operate with limited computing resources,

constrained memory, and restricted energy availability. Their design requires careful coordination of hardware, software, and machine learning models, adding complexity to development and deployment.

1.2.3 Embedded systems

Embedded systems, widely used in AI applications, provide high-performance computing in a compact form. They benefit from the availability of development boards and can be programmed similarly to traditional computers, making them accessible to a broad community of developers. Despite these advantages, they tend to consume relatively high power, and in some cases, require custom hardware design to meet specific application needs.

1.2.4 Internet of Things

At the smallest scale, the Internet of Things (IoT) enables AI integration into pervasive, low-cost, battery-powered devices. These systems support wireless connectivity and often include sensing and actuating capabilities, making them essential for smart environments. However, IoT devices face limitations in computing power, energy efficiency, and memory capacity, which can complicate programming and constrain their ability to run advanced AI models.

CHAPTER 2

Hardware

2.1 Introduction

In embedded and edge AI systems, a typical setup includes sensors that capture data from the physical world, software that processes this data, and actuators that execute actions based on computational outcomes. All processing tasks rely on specialized hardware optimized for efficiency and performance.




Figure 2.1: AI systems

Embedded systems are computers designed to control and manage the electronics within various physical devices. Embedded software refers to the programs that run on these systems, enabling their functionality.

Unlike general-purpose computers such as laptops or smartphones, embedded systems are typically designed for a specific, dedicated task, ensuring optimized performance, reliability,

and energy efficiency for their intended application.

2.2 Architecture

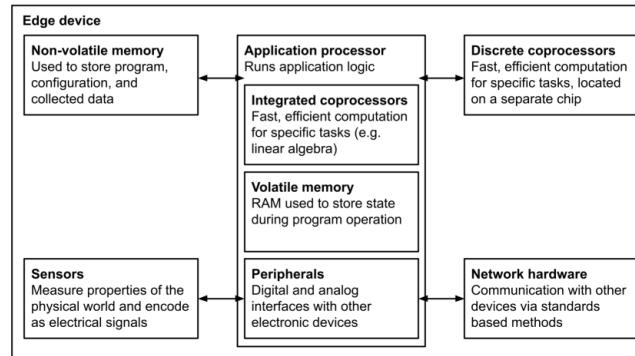


Figure 2.2: Hardware architecture

The hardware architecture of embedded and edge AI systems consists of several key components:

- *Non-volatile memory*: used to store programs, configurations, and collected data. Flash memory is typically used for this purpose, as it retains data even when the system is powered off. It is ideal for storing information that does not change frequently but is slow to read and extremely slow to write.
- *Application processor*: runs the application logic and manages program execution. It includes an integrated coprocessor for efficient computation of specific tasks, volatile memory (RAM) for storing the system state during operation, and various digital and analog peripherals that allow interaction with other electronic components.
- *Discrete coprocessors*: external chips designed for high-speed, efficient mathematical computations. They provide additional processing power for specialized AI workloads that require high performance.
- *Sensors*: measure physical-world properties and convert them into electrical signals for processing. They enable real-time data collection, which is essential for AI-driven decision-making.
- *Network hardware*: ensures communication with other devices using standardized protocols. Reliable connectivity is crucial for data exchange in distributed AI systems.

RAM is often the performance bottleneck in embedded and edge AI systems. It is very fast but consumes significant energy, making efficiency critical in power-sensitive applications. Since it is volatile, data is lost when power is turned off. RAM is also costly and takes up a large physical footprint, impacting the overall design of embedded AI devices.

2.3 Sensors and signals

Sensors are used to acquire measurements from the environment or from human interactions. They generate continuous streams of data, which can be used for various AI-driven applications.

In addition to sensor-generated data, other sources such as digital device logs, network packets, and radio transmissions can also provide valuable information. Sensors can output data in different formats depending on their purpose and design.

Data storage Data values can be stored in various formats, depending on precision and memory constraints. Boolean values (1 bit) represent binary states with two possible values. An 8-bit integer can store up to 256 distinct values, while a 16-bit integer extends this range to 65,536 possible values. A 32-bit floating point number can represent a wide range of values with up to seven decimal places, reaching a maximum of approximately 3.4×10^{38} . Quantization techniques help optimize memory usage by reducing the required storage for each value while maintaining sufficient precision for AI computations.

2.3.1 Data

Time series data A time series is a sequence of data points recorded in chronological order:

$$X = (x_1, x_2, \dots, x_N)$$

Essentially, it represents observations collected at consistent time intervals. Key factors to consider include:

- *Sampling period*: the time gap between consecutive data points.
- *Bit depth* (n): the number of bits used to represent each value.
- *Memory usage*: each sample requires n bits of storage.

Audio data Audio data is a specific type of time series, representing sound wave oscillations as they travel through air. Key parameters are:

- *Sampling rate* (Hz): the number of samples taken per second.
- *Quantization* (bit depth): the precision of each sample.
- *Signal duration* (s): the total length of the recording.
- *Number of channels*: mono (single channel) or stereo (two channels).

Memory consumption is calculated as:

$$\text{length} \times \text{sampling} \times \text{bit} \times \text{channel}$$

Image data Images capture visual information as a grid of pixels, where each pixel represents a specific property of the scene. Key characteristics are:

- *Dimensions* ($W \times H$): width and height of the image.
- *Bit depth* (N): the number of bits used to store each pixel.
- *Number of channels*: typically 1 (grayscale) or 3 (RGB).

Memory usage is given by:

$$W \times K \times N \times \text{channels}$$

Video data Videos are sequences of images displayed rapidly to create motion. They share the same structure as images but add an extra dimension: time. Critical parameters:

- *Resolution* ($W \times H$): width and height of each frame.
- *Bit depth* (N): bits per pixel.
- *Number of channels*: defines color representation.
- *Frame rate* (fps): number of frames per second.
- *Duration* (s): total length of the video.

Memory requirements are determined by:

$$W \times K \times N \times \text{channels} \times \text{frame rate} \times \text{length}$$

2.3.2 Sensors

There are thousands of different types of sensors available, each designed to capture specific kinds of data. In the context of embedded and edge AI, sensor technologies can be categorized into six main families:

1. *Acoustic and vibration*: detects sound and mechanical vibrations.
2. *Visual and scene*: captures images, video, and environmental light data.
3. *Motion and position*: measures movement, acceleration, and spatial positioning.
4. *Force and tactile*: detects pressure, touch, and force.
5. *Optical, electromagnetic, and radiation*: measures light, radio waves, and radiation levels.
6. *Environmental and chemical*: monitors temperature, humidity, gases, and other environmental factors.

Acoustic and vibration Detecting vibrations is a crucial capability in embedded and edge AI. These sensors allow systems to perceive movement, structural vibrations, and even communication signals from humans and animals at a distance. Acoustic sensors measure vibrations traveling through different media: air (microphones), water (hydrophones), and ground (geophones and seismometers). Since acoustic data is distributed across different frequencies, the sampling frequency plays a key role in ensuring accurate representation for a given application. These sensors typically produce audio data as their output.

Visual and scene Visual sensors capture information about the environment without direct contact. These range from tiny, low-power cameras to high-resolution multi-megapixel sensors. Key characteristics of image sensors: color channels, spectral response (infrared sensors), pixel size, resolution, and frame rate. The output of these sensors can be 2D or 3D images or video data, depending on the application.

Motion and position Motion and position sensors track movement and spatial positioning in various ways:

- *Tilt sensors*: simple mechanical switches that detect orientation changes.
- *Accelerometers*: measure acceleration along one or more axes.
- *Gyroscopes*: detect rotational movement.
- *Time-of-flight sensors*: emit light or radio waves to measure distances to objects.
- *Real-time locating systems*: use multiple transceivers placed around a space to track object positions.
- *Global Positioning System*: uses satellites to determine an object's precise location.

These sensors typically generate time-series data, tracking movement and positioning over time.

Force and tactile These sensors help users interact with devices, understand fluid and gas flow, or measure mechanical strain on objects:

- *Buttons and switches*: provide a simple on/off signal.
- *Capacitive touch sensors*: detect how much a conductive object touches a surface.
- *Strain gauges*: measure how much an object is being deformed.
- *Load cells*: determine the amount of force or weight applied to an object.
- *Flow sensors*: track the movement of liquids and gases.
- *Pressure sensors*: measure pressure in gases or liquids, whether in the environment or within a system.

Optical, electromagnetic, and radiation These sensors detect electromagnetic radiation, magnetic fields, and electrical properties:

- *Photosensors*: detect light across different wavelengths, from visible to infrared and ultraviolet.
- *Color sensors*: measure surface colors to help recognize different materials or objects.
- *Spectroscopy sensors*: analyze how light interacts with materials to determine their composition, useful in AI-driven material analysis.
- *Magnetometers*: measure the strength and direction of magnetic fields, like a digital compass.
- *Inductive proximity sensors*: use electromagnetic fields to detect nearby metal objects, commonly used in vehicle detection for traffic monitoring.
- *Electromagnetic field meters*: measure the intensity of electromagnetic fields, useful in industrial safety assessments.
- *Current and voltage sensors*: monitor electrical parameters, often represented as time-series data.

Environmental, biological, and chemical These sensors track environmental conditions, biological signals, and chemical presence:

- *Temperature sensors*: measure heat levels in the surroundings or specific systems.
- *Gas sensors*: detect gases like humidity, carbon dioxide, or air quality pollutants.
- *Particulate matter sensors*: monitor air pollution levels by measuring fine particles.
- *Biosignal sensors*: record electrical activity in the body.
- *Chemical sensors*: detect the presence or concentration of specific chemicals.

Most of these measurements are typically recorded as time-series data, allowing for trend analysis and predictive insights.

2.4 Microprocessor

A microprocessor is a general-purpose processor responsible for running embedded applications. Microcontrollers form the core of many modern, pervasive computing applications. These tiny, cost-effective computers are designed for specific tasks, making them ideal for embedded systems.

Unlike traditional computers, microcontrollers do not require an operating system. Instead, they run firmware (software that is directly executed on the hardware). This firmware includes low-level instructions to manage peripherals and system functions. The defining feature of microcontrollers is that they integrate all essential components into a single silicon chip.

2.4.1 Microcontroller

Microcontrollers have a fixed hardware architecture built around a central processing unit. The CPU manages a range of peripherals, providing both digital and analog functionality. Smaller devices typically include both volatile (RAM) and non-volatile (Flash/EEPROM) memory on the chip, while more powerful processors may require external memory. Programming is commonly done using low-level languages like Assembly or high-level languages like C.

Microcontrollers are the preferred choice for embedded systems because they integrate essential components like memory and peripherals, reducing the need for additional circuitry. This allows for compact, power-efficient designs that are crucial in space-constrained applications. A microcontroller typically consists of:

- A microprocessor.
- Program memory (Flash).
- Data memory (RAM).
- Various on-chip peripherals such as timers, serial communication ports, GPIO pins, counters, and ADCs (analog-to-digital converters).

Modern microcontrollers have significantly improved in both performance and efficiency:

- Higher clock speeds and 32-bit architectures for better computational power.

- SIMD support for parallel computing, beneficial in signal processing and machine learning.
- Larger RAM to handle more complex tasks—critical since memory is often a limiting factor.
- Improved energy efficiency, making them suitable for battery-powered applications.

Microprocessor and microcontroller A microprocessor is a more advanced computing unit, typically requiring external memory to fetch and execute program instructions. It offers greater processing power compared to microcontrollers but lacks integrated peripherals.




Figure 2.3: Microprocessor and microcontroller

2.4.2 System on chip

A System on Chip (SoC) integrates all the core functionalities of a traditional computing system into a single chip. This includes the CPU, memory, input/output interfaces, and sometimes specialized accelerators for AI, graphics, or signal processing. They run full operating systems and support standard development tools and libraries. However, SoCs have lower energy efficiency compared to microcontrollers, making them less ideal for ultra-low-power applications, and also more costly.

2.4.3 Comparison

The table below compares different processing devices based on their ability to handle various types of data.

Device	Time Series	Audio	Images	Video
<i>Low-end MCU</i>	Limited	None	None	None
<i>High-end MCU</i>	Full	Full	Low resolution	Limited
<i>High-end MCU with accelerator</i>	Full	Full	Full	Limited
<i>SoC</i>	Full	Full	Full	Full
<i>SoC with accelerator</i>	Full	Full	Full	Full
<i>Edge server</i>	Full	Full	Full	Full
<i>Cloud</i>	Full	Full	Full	Full

CHAPTER 3

Software

3.1 Introduction

Artificial Intelligence systems in embedded and edge computing are designed to process data efficiently and make real-time decisions.




Figure 3.1: AI systems

The first stage is preprocessing and feature extraction, where incoming data streams from sensors are processed. This step involves segmenting the data into meaningful windows, removing noise, and extracting relevant features to enhance accuracy.

Once the data is preprocessed, machine learning and deep learning algorithms are applied to analyze and interpret it. These algorithms recognize patterns, make predictions, and generate insights based on the extracted features.

Finally, in the postprocessing and decision-making stage, the system interprets the AI model's output and translates it into meaningful actions. This could involve making decisions, triggering automated responses, or providing feedback to users.

3.2 Data preprocessing and feature extraction

The goal of data preprocessing and feature extraction is to convert raw signals into structured data that can be processed by Machine Learning and Deep Learning algorithms. This process involves three key steps:

1. *Segmenting the raw signals*: breaking the continuous signal into smaller, manageable chunks of data.
2. *Processing the data*: applying digital signal processing techniques to reduce noise and enhance the most relevant parts of the signal.
3. *Feature extraction*: identifying and extracting meaningful patterns or characteristics from the processed signal chunks to be used in model training.

3.2.1 Data segmentation

Sensors generate continuous streams of data, which must be divided into smaller segments (windows) for processing. Each window represents a chunk of data that is analyzed by an algorithm to produce meaningful results.

Definition (*Latency*). Latency refers to the time required to process a single chunk of data.

Latency plays a crucial role in determining how efficiently an embedded system can process data. Lower latency allows for a higher number of processed chunks per unit of time. However, there is a trade-off:

- Larger windows increase latency but provide more information, often leading to improved accuracy.
- Smaller windows reduce latency but may capture less useful data, potentially affecting performance.

Windows can be:

- Overlapping, ensuring no information is lost from the signal.
- Non-overlapping, which may be more computationally efficient but risks missing important details.

The choice of algorithm significantly impacts both computational requirements (latency) and memory usage, making it a key design consideration.

Frame rate The frame rate defines how frequently the system can acquire and process data, similar to how frame rate applies to image and video streaming.

High latency can hinder real-time analysis by preventing the system from processing new incoming data while still handling previous chunks, potentially leading to data loss. Optimizing both latency and frame rate is essential for effective data segmentation and real-time performance.

3.2.2 Data processing

Data preprocessing using digital processing algorithms typically involves three key steps:

1. *Reconstruction of missing data:*

- *Global filling methods:* filling in missing data based on patterns and trends observed across the entire dataset.
- *Local filling methods:* estimating missing values using nearby data points. Examples include forward fill, moving averages, and local interpolation techniques.
- *Deletion of affected time periods:* in some cases, time periods with missing data are removed entirely to avoid inaccuracies in analysis.

2. *Resampling:*

- *Time series resampling:* Handling time series data with varying sampling frequencies. In upsampling we increase the sampling rate by replicating or interpolating between timestamps. In downsampling we reduce the sampling rate through subsampling. Be cautious of aliasing when changing the sampling frequency.
- *Image resampling:* adjusting the spatial resolution (pixels per image). In downsampling we reduce the image resolution by decreasing the number of pixels. In interpolation we increase the image resolution by adding pixels based on existing data.
- *Shape modification for images:* In cropping we trim parts of the image. In resizing we adjust the image dimensions while maintaining or altering aspect ratios.

3. *Filtering:* we can use different types of filter:

- *Low-pass filter:* retains low frequencies, removing high-frequency noise. Pay attention to the cutoff frequency and frequency response.
- *High-pass filter:* retains high frequencies, removing low-frequency components. The cutoff frequency and response are important here as well.
- *Band-pass filter:* keeps a specific range of frequencies, removing those outside the band.

3.2.3 Feature extraction

Feature extraction can be applied across various domains and types of data:

- *Time domain:* features like mean, PCA eigenvalues, amplitude, signal-to-noise ratio (SNR), peak decay, and energy can be extracted.
- *Frequency domain:* features such as maximum amplitude, dominant frequency, and peak variance are commonly extracted.
- *Images:* features like edges, corners, blobs, and ridges (curves) are often detected. In embedded systems, tools for feature detection include OpenCV, a widely used library for image processing and feature detection, especially in System on Chips, and OpenMV, a library optimized for high-end microcontroller units, which focuses on efficient image processing and feature extraction.

Sensor fusion In sensor fusion, the goal is to combine data from multiple sensors rather than relying on a single sensor. Each sensor provides unique perspectives and data, and by combining these diverse inputs, we create a more comprehensive and accurate representation of the environment or system.

The features extracted from each sensor are merged to enhance the robustness, accuracy, and reliability of the analysis. These fused features are then used in machine learning or deep learning models for further processing, ultimately enabling more precise and informed decision-making.

Normalization and standardization For more efficient training, data should be normalized or standardized. Features with different scales can negatively impact the performance of machine learning models, potentially leading to underfitting.

- *Minmax normalization*: scales data to a range of $[0, 1]$:

$$x' = \frac{x - x_{\min}}{x_{\max} - x_{\min}}$$

- *Standardization*: centers data around zero (mean), but does not limit it to a specific range:

$$x' = \frac{x - \mu}{\sigma}$$

Here, μ is the mean and σ is the standard deviation.

3.3 Machine and Deep Learning algorithms

Functionality Machine and Deep Learning algorithms can be categorized based on their functionality:

- *Classification*: assigns input data to predefined categories, such as identifying spam emails or recognizing handwritten digits.
- *Regression*: predicts continuous values based on input data, such as forecasting stock prices or estimating house prices.
- *Object detection*: identifies and locates objects within an image or video, commonly used in autonomous driving and security surveillance.
- *Segmentation*: divides an image into meaningful regions, such as medical image analysis or self-driving car lane detection.
- *Anomaly detection*: identifies rare or unusual patterns in data, useful in fraud detection and predictive maintenance.
- *Prediction*: forecasts future outcomes based on historical data, such as weather prediction or demand forecasting.
- *Feature reduction*: reduces the dimensionality of data while preserving essential information, improving efficiency in Machine Learning models.

Implementation Machine and deep learning algorithms can be categorized based on their implementation approach:

- *Conditional logic*: rule-based systems that use explicit if-then conditions to make decisions, commonly found in expert systems and traditional automation. The advantages of conditional logic are: determinism, efficiency, and there is no need for training data.
- *Machine Learning*: algorithms that learn patterns from data and make predictions without explicit programming.
- *Deep Learning*: a subset of machine learning that utilizes multi-layered neural networks to model complex patterns.

3.3.1 Tiny Machine Learning

Tiny Machine Learning focuses on bringing machine learning and deep learning capabilities to small, low-power devices, particularly in the Internet of Things. This enables real-time data processing directly on the device, reducing reliance on cloud computing and improving efficiency, privacy, and responsiveness.

The key steps in deploying TinyML applications are:

1. *Set up the hardware*: choose a suitable microcontroller or edge device with sufficient processing power and energy efficiency.
2. *Install the software*: set up the necessary development tools, frameworks, and libraries.
3. *Collect data*: gather relevant sensor data to train the model.
4. *Train the model*: develop and optimize a lightweight machine learning model tailored for low-power devices.
5. *Build the application*: integrate the trained model into an application that interacts with the device.
6. *Optimize and compile*: convert the application into a format suitable for deployment on the target hardware.
7. *Deploy to the device*: flash the compiled application onto the microcontroller or embedded system.
8. *Test and monitor*: evaluate the performance of the deployed model in real-world conditions and refine as needed.

Level	Category
6	End device training and inference
<i>5</i>	Edge training and inference
<i>4</i>	Cloud-edge co-training and inference
3	Cloud training, on-device inference
<i>2</i>	Cloud training, edge inference
<i>1</i>	Cloud training, cloud and edge inference
	Cloud training and inference

Tiny Machine Learning

4.1 Machine Learning

In traditional programming, we provide a computer with both data and a program, and it produces an output based on the instructions we've given. In contrast, with Machine Learning, we feed the computer data and corresponding outputs. The computer then trains itself by adjusting its parameters to minimize a loss function, ultimately producing a trained model that can make predictions.




Figure 4.1: Traditional programming and Machine Learning

An essential aspect of Machine Learning is choosing the right model for the task. We aim to select a model that minimizes the error when classifying the training data. Models can belong to various families, such as linear or polynomial models, neural networks, and more.

4.1.1 Supervised Learning

Let's consider a simple two-dimensional problem. Designing a classifier requires us to identify a function that separates the labeled data points. Some key aspects to focus on include:

- Linear or nonlinear approaches.
- The number of points available or the total number of points.
- The choice of technique for designing the classifier.

Regression In nonlinear regression, given a set of noisy data points (x_i, y_i) , the goal is to reconstruct the underlying function. This is expressed as:

$$y(x, \mathbf{w}) = \sum_{j=0}^M w_j x^j$$

To train the model, we minimize the Residual Sum of Squares with the following formula:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(x_n, \mathbf{w}) - t_n)^2$$

Here, $y(x_n, \mathbf{w})$ is the model's output for the input sample x_n , x_n represents the input data, and t_n is the target value (the expected output for a given input).

4.1.2 Training

The training error isn't always reliable because it only measures the performance of the model on the data used for training. To get a better understanding of how well the model generalizes, we need to split the dataset into a training set and a test set, ensuring that the model isn't tested on the same data it was trained on.

Typically, the test error will be higher than the training error, as the model is being evaluated on unseen data. The amount of data we have plays a crucial role in determining the final quality of the model (more data generally leads to better results).

Additionally, we can introduce a validation set. This allows us to adjust the complexity of the model based on the available data, helping to avoid overfitting or underfitting the model to the training data.

Model complexity The complexity of a model, denoted by M , plays a crucial role in its performance. If the model is too complex for the available data, it will overfit (it learns the training data too well, including noise, which leads to poor generalization and weak performance on new data). On the other hand, if the model is too simple, it won't capture all the important patterns in the data, leading to high error even on the training set. This is known as underfitting. To achieve the best performance on both training and testing data, we must find an optimal model complexity that balances these two extremes.

Number of samples Additionally, the number of samples, N significantly impacts the model's reliability. If we have too few samples, the model may fail to generalize properly, making it difficult to learn the correct patterns.

Error The total error in a learning model is composed of three main components:

$$\text{error} = \text{approximation error} + \text{estimation error} + \text{inherent error}$$

Here, we have:

- *Approximation error*: this occurs when the learning algorithm does not find the best possible model within the chosen hypothesis space. In other words, it's the difference between the best theoretical model within our chosen framework and the actual model we obtain due to limitations in data or training.

- *Estimation error*: this arises when the chosen model family (or hypothesis space) does not perfectly match the true data-generating process. It represents the error introduced by selecting a model class that may not include the optimal solution.
- *Inherent error*: this error is unavoidable as it stems from the complexity and noise in the data itself. It depends on the fundamental nature of the learning problem and can only be reduced by improving data quality or redefining the problem.

Reducing both approximation and estimation errors is key to improving model performance, while inherent error is often a fundamental limitation of the problem itself.

4.1.3 Validation

To properly evaluate a machine learning model, we divide the dataset into three parts:

- *Training set*: used to train the model by adjusting its parameters.
- *Validation set*: used to tune hyperparameters and prevent overfitting.
- *Test set*: used to assess the final model's performance on unseen data.

When evaluating a classifier, we use the confusion matrix, which summarizes the model's performance by comparing predicted and actual class labels:

	Predicted positive	Predicted negative
Actual positive	TP (True Positive)	FN (False Negative)
Actual negative	FP (False Positive)	TN (True Negative)

Table 4.1: Confusion matrix

Here, the total number of samples is:

$$N = TP + TN + FP + FN$$

From this matrix, we can derive key performance metrics:

- *Accuracy*: measures overall correctness but can be misleading for imbalanced datasets:

$$\text{accuracy} = \frac{TP + TN}{N}$$

- *Receiver Operating Characteristics*: plots the True Positive Rate against the False Positive Rate. The classifier's performance is represented as a point on this curve. Changing the classification threshold alters this position. A diagonal line corresponds to random guessing, and any performance below this line implies systematically incorrect predictions.

Model validation To validate a model we can use various techniques:

- *Apparent Error Rate* or *resubstitution*: the entire dataset Z_N is used for both training and error estimation. However, this method often underestimates the true error due to overfitting.
- *Sample Partitioning*: the dataset Z_N is randomly split into two disjoint subsets: S_D , used for training, and S_E , used for evaluation.
- *Leaving-One-Out*: improves upon this by iterating over the dataset N times. Each time, one sample is left out for evaluation while the remaining $N - 1$ samples are used for training. The final estimate is obtained by averaging all N iterations.
- *k-fold Cross Validation*: the dataset is randomly divided into k disjoint subsets of equal size. For each subset, the remaining $k - 1$ subsets form the training set S_D , while the reserved subset is used as the evaluation set S_E . The process repeats k times, and the results are averaged. This approach generalizes LOO and reduces variance when $k \ll N$.

By carefully selecting the validation method, we can ensure a more reliable assessment of the model's performance while minimizing bias and variance.

4.1.4 Neural Networks

Neural networks are mathematical models designed to capture relationships in space, time, and the state of individual neurons. Each neuron processes input data through a weighted sum:

$$a_t = \sum_{i=1}^m x^i w^i$$

This scalar product measures the affinity between input values and weights. Before activation, it represents the neuron's raw output. An activation function is then applied to introduce non-linearity and determine the final neuron output.

A typical neural network consists of an input layer, one or more hidden layers, and an output layer. Training a neural network involves minimizing a loss function. For a feedforward neural network, a common loss function is the Residual Sum of Squares:

$$\mathcal{L}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (y(\mathbf{x}_n, \mathbf{w}) - t_n)^2$$

Since this function lacks a closed-form solution due to its non-linearity, optimization techniques such as Stochastic Gradient Descent (SGD) are used to iteratively update weights:

$$\mathbf{W}^{(t+1)} = \mathbf{W}^{(t)} - \eta \nabla \mathcal{L}(\mathbf{W}^{(t)})$$

Here, η represents the learning rate, controlling the step size in the direction of the gradient $\nabla \mathcal{L}(\mathbf{W}^{(t)})$. The goal is to find the best possible local minimum, as computing the global minimum is often infeasible.

Theorem 4.1.1. *A Feed Forward Neural Network with a single hidden layer containing a finite number of neurons can approximate any continuous function defined on compact subsets.*

This theorem implies that, in theory, a properly configured neural network can achieve zero approximation error for any target function. However, in practice, finding the exact network that achieves this perfect approximation is extremely challenging.

Memory demand The memory required by a Feedforward Neural Network depends on the number of layers and the number of connections within each layer. Every connection between neurons has an associated weight, and each neuron also has a bias, which is another weight that needs to be stored. For a given layer, the total memory demand for weights can be calculated using the following formula:

$$\text{weights} = \sum_{l=1}^L = (N_l \cdot N_{l-1} + N_l)$$

Here, L is the total number of layers (excluding the input layer), N_l is the number of neurons in layer l , N_{l-1} is the number of neurons in the previous layer. This formula gives the total number of weights that must be stored in memory. If each weight is stored using a specific precision, the total memory demand can be computed by multiplying the number of weights by the storage size per weight.

4.2 Deep Learning

Deep learning is a powerful subset of machine learning that focuses on automatically learning the most relevant features from raw data. Unlike traditional machine learning methods, which rely on manually extracted features, deep learning models take raw input and learn hierarchical representations directly. This capability makes deep learning highly effective for complex tasks but also introduces challenges in terms of model complexity and interpretability. While deep learning often outperforms traditional machine learning, it is not always the best choice (especially when domain expertise allows for efficient manual feature extraction). Additionally, deep learning models are computationally more demanding and require significantly more resources to train and deploy.

A typical convolutional neural network (CNN) consists of two main types of layers: convolutional layers and fully connected layers.

4.2.1 Convolutional layer

Each convolutional layer typically includes several key components:

- *Convolutional filters* (trainable): these apply spatial operations to the input, with parameters such as filter size, number of filters, and stride.
- *Activation functions*: nonlinear transformations introduce nonlinearity into the model, enabling it to learn complex relationships.
- *Normalization* (trainable) : optional techniques like batch normalization improve training stability and performance.
- *Pooling*: reduces spatial dimensions, helping to control overfitting and computational costs.

Consider an RGB image with dimensions $H \times W \times C$, where H and W represent the height and width of the image, and C denotes the number of channels (3 for RGB, 1 for gray scale). Convolutional filters have dimensions $R \times S \times C$, where R and S are the filter's spatial dimensions, and C matches the number of channels in the input image.

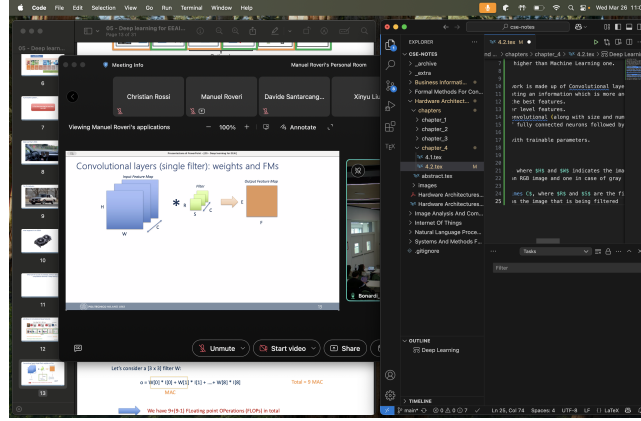


Figure 4.2: Convolutional layer

The output of applying a filter to the image has dimensions $E \times F$, with the number of channels equal to the number of filters used. The exact size of the output depends on the filter dimensions, the stride (U), and whether padding is applied. Specifically, with no padding we have :

$$\begin{cases} E = \frac{H-R+U}{U} \\ F = \frac{W-S+U}{U} \end{cases}$$

Each filter has $R \times S \times C + 1$ parameters (including the bias term). Even with a single filter, the number of parameters can grow rapidly, especially for large images.

Computational demand The computational cost of a convolutional layer is primarily determined by the number of multiply-and-accumulate (MAC) operations. For a single 3×3 filter, there are 9 MACs (one for each weight) and a total of $9 + (9 - 1) = 17$ floating-point operations (FLOPs). In general:

$$1 \text{ MAC} \approx 2 \text{ FLOPS}$$

For a convolutional layer with M filters, the total number of MAC operations is:

$$\text{MAC} = E \times F \times R \times S \times C \times M$$

Here, where $E \times F$ is the output size, $R \times S$ is the filter size, C is the number of input channels, and M is the number of filters.

4.2.2 Fully connected layer

Fully connected layers consist of densely connected neurons followed by activation functions. These layers are responsible for combining the high-level features extracted by the convolutional layers to make predictions.

In a fully connected layer, each neuron receives input from all neurons in the previous layer. Given an input layer with H neurons and a fully connected layer with W neurons, the total number of parameters is $H \times W + W$, where the additional W accounts for the bias terms.

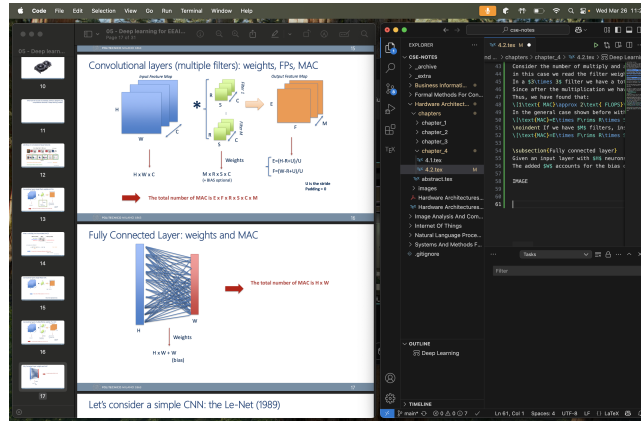


Figure 4.3: Fully connected layer

The computational cost is proportional to the number of MAC operations:

$$\text{MAC} = H \times W$$

4.2.3 Memory occupation

Convolutional layers generally occupy less memory than fully connected layers due to their sparse connectivity and parameter sharing. However, they tend to require more MAC operations, making them computationally expensive. Each parameter occupies approximately 4 bytes, so the total memory footprint of a neural network is determined by the number of parameters multiplied by this value. Additionally:

- The firmware for the device typically requires around 10 kB.
- Input, output, and intermediate feature maps must also be stored in memory. Their size is given by:

$$H \times W \times C \times 4$$

This can be a significant bottleneck, especially for large inputs or deep networks.

Modern architectures often reduce the number of fully connected layers while increasing the depth and efficiency of convolutional layers.

4.2.4 Tiny Deep Learning

To address the challenges posed by the size and complexity of neural networks, we can explore several strategies to make them more efficient and suitable for resource-constrained environments:

- *Redesigning CNN architectures*: develop new, lightweight layers that reduce the overall computational burden while maintaining performance. This involves creating compact network structures that achieve high efficiency without compromising too much on accuracy.
- *Approximate computing*: embrace techniques that trade off a small amount of precision for significant gains in computational speed and memory usage. By allowing minor inaccuracies in calculations, we can drastically improve the efficiency of both training and inference processes.

- *Code optimization for embedded systems*: fine-tune the implementation of deep learning models to better suit the hardware constraints of embedded devices. This includes optimizing memory usage, leveraging hardware-specific instructions, and minimizing latency to ensure smooth operation on edge devices.

4.3 Tiny Deep Learning architectures

Deep learning models designed for resource-constrained devices, such as edge devices, must balance computational efficiency with accuracy. These architectures are optimized to reduce memory usage, computation overhead, and energy consumption while maintaining competitive performance. In this context, three prominent families of architectures have emerged: SqueezeNet, MobileNet, and EfficientNet.

4.3.1 SqueezeNet

Introduced in 2016, SqueezeNet was the first convolutional neural network (CNN) architecture explicitly designed for deployment on edge devices. It achieves a remarkable reduction in model size (up to 50x to 500x fewer parameters than AlexNet) while maintaining comparable accuracy. This makes SqueezeNet highly suitable for distributed training environments, where communication overhead is proportional to model size. Additionally, its smaller footprint reduces the cost of exporting updated models to clients.

The objectives of SqueezeNet are to design a CNN with fewer parameters but equivalent accuracy to AlexNet. To achieve this goal, SqueezeNet employs several innovative strategies:

1. *Replace 3×3 filters with 1×1 filters*: using 1×1 filters instead of 3×3 filters reduces the number of parameters by a factor of 9. This also decreases the number of MACs. However, not all 3×3 filters are replaced, as they are essential for capturing spatial patterns.
2. *Reduce input channels to the 3×3 filters*: to further minimize the number of parameters, SqueezeNet introduces squeeze layers that reduce the number of input channels to 3×3 filters. This results in fewer 3×3 filters overall. This is done using squeeze layers. The reduction in parameters and MACs is proportional to $9 \times C_{\text{removed}}$, where C_{removed} is the number of removed channels.
3. *Downsample late in the network*: delaying downsampling preserves larger activation maps, which can improve accuracy despite using fewer weights. While this increases memory usage, it prevents the loss of critical features in early stages due to overly aggressive dimensionality reduction.

Fire module At the heart of SqueezeNet is the fire module, a novel convolutional layer designed to drastically reduce memory and computation demands while preserving accuracy. Each Fire module consists of two stages:

1. *Squeeze layer*: applies only 1×1 convolutional filters. Reduces the number of input channels to s , a hyperparameter chosen during design. The output dimensions is $H \times W \times s$.

2. *Expand layer*: takes the reduced activation map as input. Applies e 3×3 filters and e 1×1 , producing an output with $2e$ channels. Need $s < 2e$ to maintain computational efficiency. The output dimensions is $H \times W \times 2e$.

4.3.1.1 Architecture

The final example architecture is composed by:

1. *Initial convolutional layer*: processes the input image and prepares it for subsequent Fire modules.
2. *Eight fire modules*: each module progressively increases the number of filters to capture more complex features. Max-pooling is applied after Fire modules 4 and 8 to downsample the feature maps.
3. *Global Average Pooling and softmax*: instead of fully connected layers, SqueezeNet uses GAP followed by a softmax layer for classification. This eliminates the need for dense layers, further reducing the parameter count.

Further techniques, such as sparsity (pruning unnecessary weights) and quantization (reducing precision), can be applied to SqueezeNet to enhance efficiency without significant accuracy loss. The base SqueezeNet model achieves AlexNet-level accuracy while requiring far fewer parameters and MACs.

4.3.1.2 Variants

Several variants of SqueezeNet have been proposed to improve accuracy with minimal additional complexity:

1. *Simple bypass architecture*: adds skip connections around specific fire modules. These connections allow the modules to learn residual functions, improving gradient flow and accuracy. No extra parameters are introduced, making this variant highly efficient.
2. *Complex bypass architecture*: incorporates 1×1 convolutional layers within the skip connections. The number of filters in these convolutions equals the number of output channels, introducing additional trainable parameters. While slightly more accurate, this variant increases the model's size and computational cost.

Among the variants, the simple bypass architecture provides the best trade-off, offering modest accuracy improvements without adding extra weights.

4.3.2 MobileNet

MobileNet is a family of efficient neural network architectures designed specifically for mobile and embedded vision applications. It achieves lightweight computation through the use of depth-wise separable convolutions, a streamlined approach that significantly reduces the number of parameters and computational cost while maintaining competitive accuracy. This architecture has become a cornerstone for deploying deep learning models on resource-constrained devices.

Depth-wise scalable convolution The key innovation in MobileNet is the use of depth-wise separable convolutions, which decompose a standard convolution into two simpler operations:

1. *Depth-wise convolution*: applies a single $R \times S$ filter to each input channel independently. Unlike a standard convolution, this step does not combine information across channels. The output dimensions is $H \times W \times C$, where C is the number of input channels. The MAC operations are: $\text{MAC} = R \times S \times C \times E \times F$. The number of parameters is: $\text{weights} = R \times S \times C$.
2. *Point-wise convolution*: combines the outputs of the depth-wise convolution across channels using M 1×1 filters. This step aggregates channel-wise information, producing the final output feature map. The MAC operations are: $\text{MAC} = C \times M \times E \times F$. The number of parameters is: $\text{weights} = \text{MAC} = C \times M$.

Convolution	MAC ops	#Weights
Standard	$R \times S \times C \times E \times F \times M$	$R \times S \times C \times M$
Depth-wise	$R \times S \times C \times E \times F + C \times M \times E \times F$	$R \times S \times C + C \times M$

By decoupling spatial and channel-wise computations, depth-wise separable convolutions achieve significant reductions in both MAC operations and the number of weights compared to standard convolutions.

4.3.2.1 Architecture

Each block in MobileNet follows a consistent structure:

1. *Depth-wise convolution*: applies a 3×3 filter to each input channel independently. Reduces computational cost by avoiding cross-channel interactions at this stage.
2. *Batch normalization*: normalizes the outputs of the depth-wise convolution to stabilize training.
3. *ReLU activation*: introduces non-linearity to the model.
4. *Point-wise convolution*: combines information across channels using 1×1 convolutions.
5. *Batch normalization*: normalizes the outputs of the point-wise convolution.
6. *ReLU activation*: applies another non-linear transformation.

Regularization and augmentation MobileNet requires less regularization and data augmentation compared to larger models like VGG or GoogleNet. This is because smaller models are less prone to overfitting due to their reduced capacity. Despite this, MobileNet achieves comparable or better accuracy than many state-of-the-art models, including GoogleNet, while being significantly more efficient.

Performance MobileNet's accuracy is slightly higher than GoogleNet and comparable to VGG, but with far fewer parameters and computations. While MobileNet has more parameters than SqueezeNet, it achieves higher accuracy, making it a better choice for applications where performance is critical. The use of depth-wise separable convolutions makes MobileNet highly efficient, enabling real-time inference on mobile and embedded devices.

4.3.3 EfficientNet

EfficientNet represents a breakthrough in the design of efficient convolutional neural networks (CNNs) by introducing a novel method for scaling up models. Unlike traditional approaches that focus on increasing either depth, width, or resolution individually, EfficientNet employs compound scaling, which scales all three dimensions simultaneously while maintaining a balanced relationship between them. Scaling a CNN typically involves:

- *Increasing depth*: adding more layers to capture complex features.
- *Increasing width*: expanding the number of filters in each layer to improve representational capacity.
- *Increasing resolution*: using higher-resolution input images to capture finer details.

4.3.3.1 Compound scaling

The key idea behind compound scaling is to scale depth ($d = \alpha^\phi$), width ($w = \beta^\phi$), and resolution ($r = \gamma^\phi$) with constant ratios:

$$\alpha\beta^2\gamma^2 \approx 2 \quad \alpha, \beta, \gamma \geq 1$$

Here, α , β , and γ are constants that control how each dimension is scaled, and γ is a user-specified coefficient that determines the extent of scaling (how many more resources are available). EfficientNet's compound scaling is implemented in two steps:

1. *Determine scaling coefficients* (α, β, γ): Start with a baseline network and assume twice the available computational resources (double ϕ). Perform a small grid search to find optimal values of α , β , and γ , subject to the constraint $\alpha\beta^2\gamma^2 \approx 2$. Fix these coefficients as constants for subsequent scaling.
2. *Scale the baseline network*: Use the fixed scaling coefficients (α, β, γ) to scale the baseline network to different sizes by varying ϕ . Each variant (B0 to B7) corresponds to progressively larger models (with larger ϕ), achieving higher accuracy at the cost of increased computational resources.