

Internet Of Things

Christian Rossi

Academic Year 2024-2025

Abstract

The course provides an overview of the four main components of IoT systems: sensors, communication technologies, management platforms, and data processing and storage platforms for sensor data. In the first part, the course covers the characteristics of the hardware components of sensor nodes (microcontrollers/microprocessors, memory, sensors, and communication devices). It then delves into communication technologies used in IoT systems, distinguishing between short-range solutions (ZigBee, 6LoWPAN) and long-range solutions (LoRaWAN, NB-IoT). Finally, the course focuses on application-level protocols for IoT systems (COAP, MQTT) and the analysis of IoT management platforms. The course includes hands-on development activities and is delivered through flipped classroom and/or blended learning formats.

Contents

1	Introduction	1
1.1	Internet	1
1.1.1	Internet of Things	1
1.1.2	Industrial IoT	2
1.1.3	Architecture	2
1.2	Hardware	2
1.2.1	Processor	3
1.2.2	Sensor	4
1.2.3	Power consumption	5
1.3	Communication	7
1.3.1	Key Performance Indicators	7
1.3.2	Communication technology	7
1.3.3	Communication protocol	8
2	Application layer	9
2.1	Introduction	9
2.1.1	Taxonomy	9
2.2	Hyper Text Transfer Protocol	9
2.2.1	Communication	9
2.3	Constrained Application Protocol	10
2.3.1	Packets	11
2.3.2	Observation	12
2.3.3	Block transfer	12
2.3.4	Resource discovery	13
2.4	Message Queuing Telemetry Transport	13
2.4.1	Publisher	13
2.4.2	Subscriber	14
2.4.3	Session	14
2.4.4	Messages	14
2.4.5	MQTT for Sensor Networks	15
3	Physical layer	16
3.1	Introduction	16
3.1.1	Ethernet	17
3.1.2	Wireless	18
3.2	Low Power Wide Area Network	18
3.2.1	Messages	20

3.2.2	LoraWAN uplink	21
3.3	Narrowband Internet of Things	21
3.3.1	Summary	23
4	Network layer	24
4.1	Introduction	24
4.2	Zigbee	25
4.2.1	Physical layer	26
4.2.2	Medium Access Control layer	26
4.2.3	Network layer	28

CHAPTER 1

Introduction

1.1 Internet

Definition (*Internet*). The Internet is a global network that connects various types of networks, enabling communication and data exchange.

Traditionally, the internet was primarily used for fixed, stationary clients accessing well-defined services. However, modern internet usage has shifted significantly with the rise of mobile clients. These mobile devices, often equipped with sensing and actuating capabilities, are no longer just consumers of information and services.

Technological advancements Several breakthroughs have paved the way for the rapid growth of the Internet of Things (IoT). The miniaturization of hardware has enabled the development of compact yet powerful smart devices. At the same time, improvements in energy solutions have enhanced the efficiency and autonomy of these devices. Increased mobility has further expanded the reach and functionality of IoT applications.

In parallel, communication protocols have evolved to support low-power wireless technologies, ensuring efficient and reliable connectivity. The widespread adoption of cloud computing has also played a crucial role, providing scalable architectures and vast processing power. Additionally, the rise of Artificial Intelligence has unlocked new possibilities for intelligent data analysis, automation, and decision-making within IoT ecosystems.

1.1.1 Internet of Things

Definition (*Internet of Things*). Internet of Things (IoT) is a worldwide network of uniquely addressable interconnected objects, based on standard communication.

The IoT is based on:

- *Smart objects*: devices embedded with sensors, actuators, and connectivity.
- *Data*: continuous collection and processing of information.
- *Pervasiveness*: seamless integration into everyday life.

- *Seamless communication*: reliable and efficient interaction between devices, networks, and services.

The IoT primarily consists of connected low-cost endpoints, such as consumer devices and everyday smart objects, which focus on accessibility and widespread adoption.

1.1.2 Industrial IoT

Definition (*Industrial IoT*). The IIoT refers to a network of interconnected sensors, instruments, and devices integrated with industrial computing applications, including manufacturing, energy management, and automation.

The IIoT consists of connected industrial assets that are typically medium to high-cost. These devices are more expensive but also more responsive.

Cybersecurity is a central concern in the IIoT, where even minor disruptions can have severe consequences. Unlike consumer IoT, IIoT systems must operate with continuous availability, robustness, and resiliency, ensuring that industrial processes remain uninterrupted. IIoT environments often coexist with a significant amount of legacy operational technologies. These legacy systems, designed for reliability rather than cybersecurity, introduce additional challenges in securing industrial networks.

While usability and user experience are critical in consumer IoT, they are not primary concerns in IIoT. Instead, the focus is on system integrity, fault tolerance, and maintaining operational continuity in complex industrial ecosystems.

1.1.3 Architecture

The IoT endpoints require strong security and reliability. These devices are not just about connectivity; they depend on a combination of smart objects, reliable connectivity, data collection, and advanced analytics to function properly. The security of IoT endpoints is critical. Ensuring reliability ensures that these devices can perform their tasks without interruption.




Figure 1.1: IoT architecture

1.2 Hardware

Definition (*Sensor node*). A sensor node (or mote) is a device with capable of sensing external phenomena and process, store and communicate the data collected by the sensors.

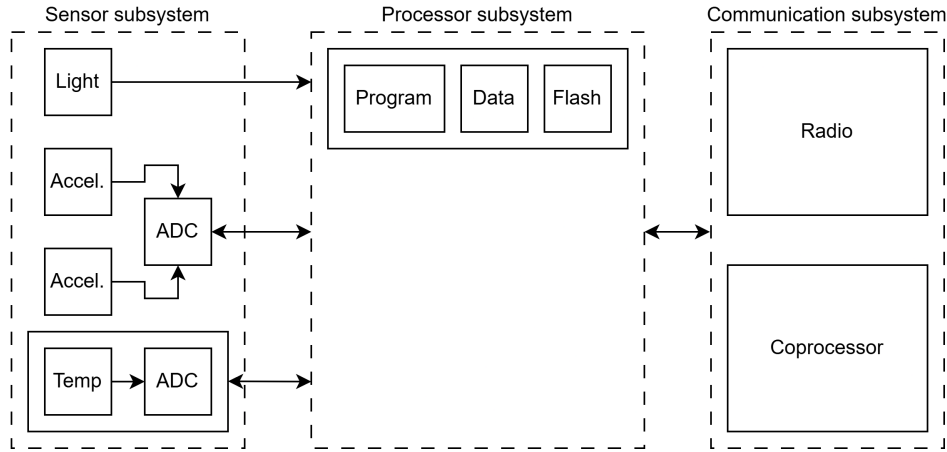


Figure 1.2: Sensor node architecture

Definition (*Actuator*). An actuator is a device capable of acting on processes in order to modify conditions based on the input data.

1.2.1 Processor

The processor subsystem of a sensor node is often designed based on the SHARC architecture.

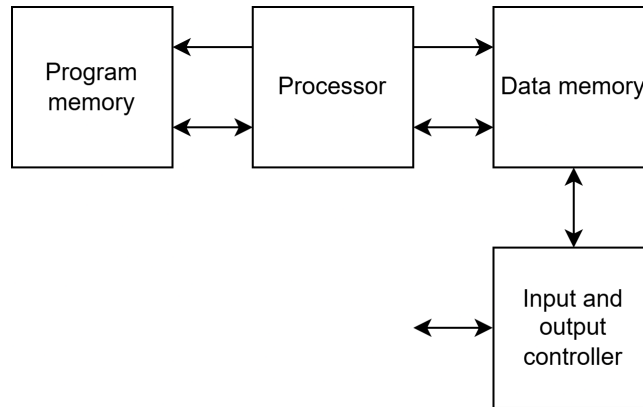


Figure 1.3: Processor SHARC architecture

A microcontroller is generally used as the processor in sensor nodes.

Definition (*Microcontroller*). A microcontroller is a single integrated circuit designed for a specific application.

A microcontroller is usually composed of a Central Processing Unit and a clock generator. It is usually equipped with RAM, flash memory and an EEPROM. It is connected with a serial BUS, I/O interfaces and analogic and digital converters. While microcontrollers are flexible and low-cost, they can compromise speed in certain use cases.

Definition (*Digital Signal Processor*). A Digital Signal Processor (DSP) is a specialized microprocessor optimized for processing discrete signals using digital filters.

DSPs excel at performing complex mathematical operations with extremely high efficiency. While they are well-suited for data-intensive operations, they are less flexible than microcontrollers.

Definition (*Application Specific Integrated Circuit*). An Application Specific Integrated Circuit (ASIC) is a custom integrated circuit tailored for a specific application.

ASICs offer high speed and can be tailored for specific tasks, but they come with high development costs and limited flexibility once designed.

Definition (*Field Programmable Gate Array*). A Field Programmable Gate Array (FPGA) has a high level architecture that allows for some degree of reconfigurability after manufacturing.

FPGAs offer high-speed performance, supporting parallel programming, and moderate reconfigurability. However, they are more complex and costly than microcontrollers or DSPs.

1.2.2 Sensor

Sensors samples the data from the real worlds and traduce them in digital format




Figure 1.4: Analog to digital converter

The Nyquist analog to digital converter involves reading a time-continuous signal at specific points in time. The sampling rate, or bandwidth, is the inverse of the sampling interval:

$$f_s = \frac{1}{T}$$

The key idea is that if the sampling frequency is properly set, the original signal can be losslessly reconstructed from its samples.

Theorem 1.2.1 (Nyquist theorem). *Given the signal bandwidth B , we have that the sampling frequency must be chosen as:*

$$f_s = 2B$$

Quantization is used to approximate the input voltage V_{in} in a digital codeword. An ideal quantizer maps input to output with the smallest variation in the input causing a change in the codeword.




Figure 1.5: Quantization

The resolution refers to the smallest input variation that causes a change in the codeword:

$$\text{LBS} = \frac{V_{FS}}{2^n}$$

Here, V_{FS} is the full-scale voltage and n is the number of bits of resolution. The quantization error occurs when the output voltage either overestimates or underestimates the input voltage. This error decreases as the resolution increases.

1.2.3 Power consumption

The absence of cables in wireless sensors and actuators means no wired power or connectivity. A sensor node typically operates with a limited power source, and its lifetime directly depends on the battery lifetime. The goal is to maximize the energy provided while minimizing the cost, volume, weight, and recharge requirements. There are two main types of batteries used:

- Primary batteries, which are not rechargeable.
- Secondary batteries, which are rechargeable, but only make sense when paired with some form of energy harvesting.

Power cycle The power cycle of an IoT device consists of sleep and active states (wake-up and work). During the sleep state, power consumption is minimal, with only essential components running.




Figure 1.6: IoT device life cycle

The average power consumption is then defined as:

$$P_{\text{avg}} = f_{\text{sleep}}P_{\text{sleep}} + f_{\text{wake up}}P_{\text{wake up}} + f_{\text{work}}P_{\text{work}}$$

Here, f_{sleep} , $f_{\text{wake up}}$, and f_{work} are the fractions of time spent in sleep, wake-up, and work states, respectively. Then, the lifetime of the device is given by:

$$\text{lifetime} = \frac{\text{energy stored}}{P_{\text{avg}} - P_{\text{gen}}}$$

Here, P_{gen} is the power generated.

1.2.3.1 Data transmission

When data needs to be sent, the device first wakes up and then performs the actual transmission. The total energy consumption for this process is given by:

$$E_{\text{tx}} = P_{\text{tx}}(T_{\text{wu}} + T_{\text{tx}}) + P_0T_{\text{tx}}$$

Here, P_{tx} is the power consumed by the transmitter, P_0 is the output power of the transmitter, T_{tx} is the time taken to transmit a packet, and T_{wu} is the wake-up time.

1.2.3.2 Data reception

When the device needs to receive data, it first wakes up and then performs the reception. The total energy consumption in this case is:

$$E_{\text{rx}} = P_{\text{rx}}(T_{\text{uw}} + T_{\text{rx}})$$

Here, P_{rx} is the power consumed by the receiver, T_{rx} is the time taken to receive a packet, and T_{uw} is the wake-up time.

Sensitivity Each receiver is characterized by a sensitivity parameter, which is the minimum input signal power required for the receiver to demodulate the data correctly. Knowing this sensitivity, the required emitted power at the transmitter can be determined by inverting the propagation law of the communication channel.

1.2.3.3 Emitted power

The emitted power is often a tunable parameter, and it is generally considered good practice to set it to the lowest value that still allows for reliable reception. The quality of the reception process is typically measured using metrics like:

- *Bit Error Rate*: the fraction of bits that are incorrectly received.
- *Packet Error Rate*: the fraction of packets that are not received correctly.

The relationship between BER and PER, for a packet of length l with independent errors, is given by:

$$\text{PER} = 1 - (1 - \text{BER})^l$$

Both BER and PER are influenced by the level of noise in the transmission and reception channels, which in turn is determined by the transmitted and received power. The Signal-to-Interference-plus-Noise Ratio is a key factor in determining this quality, and is calculated as:

$$\text{SINR} = 10 \log \left(\frac{P_{\text{recv}}}{N_0 + \sum_{i=1}^k I_i} \right)$$

Here, N_0 is the thermal noise (KTB), P_{recv} is the received power, and I_i are the interference contributions from other signals. Given the SINR and the specific modulation of the channel, BER can be computed.

1.2.3.4 Processing power

The power dissipation of the CPU is due to several factors, including:

$$P_p = P_{\text{dyn}} + P_{\text{sc}} + P_{\text{leak}}$$

Here, $P_{\text{dyn}} = CfV^2$ is the power consumed by the work done, P_{sc} is the power lost due to short circuits, and P_{leak} is the power lost due to leakage. Local data processing is crucial in minimizing power consumption, especially in multi-hop networks.

1.2.3.5 Sensor power

The power consumption due to sensing is highly dependent on the type of sensor used. A rough model for the power consumption of an analog to digital converter can be expressed as:

$$P_s \approx f_s 2^n$$

Here, f_s is the sampling frequency, and n is the resolution of the analog to digital converter in bits.

1.3 Communication

In IoT, various devices and systems need to exchange information seamlessly. Industrial environments present unique challenges for communication networks such as timeliness, determinism, and reliability.

Definition (*Cyclic traffic*). Cyclic traffic involves periodic transmissions, ensuring continuous process monitoring.

Definition (*Acyclic traffic*). Acyclic traffic occurs in response to unpredictable events.

Definition (*Multimedia traffic*). Multimedia traffic includes images, video streams, and other data-intensive transmissions.

Definition (*Backhand traffic*). Backhand traffic aggregates data flows from multiple sources, consolidating information for centralized processing and analysis.

1.3.1 Key Performance Indicators

Definition (*Throughput*). Throughput refers to the rate at which data can be transmitted over a network link, typically measured in bits per second or bytes per second.

The actual throughput depends on the type of link being used, as different technologies offer varying transmission capacities.

Definition (*Delivery time*). Delivery time represents the total time required to transfer a service data unit from the source to the destination.

Measured in seconds, it is influenced by multiple factors, including transmission time, propagation time, and the execution time of network protocols.

Definition (*Minimum cycle time*). Minimum Cycle Time defines the shortest duration needed to complete one full cycle within a control loop.

This metric is crucial in real-time systems, where consistent and predictable execution times are essential for maintaining stability and efficiency.

Definition (*Jitter*). Jitter measures the precision and reliability of periodic operations, particularly in time-sensitive applications.

Variability in timing can lead to performance issues, making jitter a critical factor in industrial and communication networks that require consistent and predictable timing.

1.3.2 Communication technology

Defining a communication technology involves several key aspects, starting with the physical infrastructure. IoT connectivity integrates multiple heterogeneous technologies to ensure seamless communication across diverse environments. The classification of communication technologies often depends on their range.




Figure 1.7: Connection range

Network topology also can vary a lot:

- *Ring topology*: structured flow of data, while linear chains connect devices sequentially.
- *Tree topology*: hierarchical communication, whereas star topology centralizes connections around a single node.
- *Mesh topology*: robust, decentralized connectivity, enhancing reliability and redundancy in complex systems.

1.3.3 Communication protocol

Communication protocols define the set of rules governing the exchange of information between two entities. In IoT communication networks, protocols facilitate packet-based data exchange. Most communication protocols follow a layered architecture. This structured approach enhances interoperability, modularity, and scalability in network communication. The layers are:

- *Application layer*: handles messages exchanged between applications, enabling services such as HTTP, MQTT, and CoAP.
- *Transport layer*: manages end-to-end communication through segmentation and reassembly, with protocols like TCP and UDP.
- *Network layer*: routes packets across different networks, using IP-based addressing and routing mechanisms.
- *Data link layer*: organizes data into frames, ensuring reliable point-to-point or point-to-multi point transmission.
- *Physical layer*: deals with the actual transmission of raw bits over physical media such as wired or wireless connections.

Application layer

2.1 Introduction

Definition (*Application*). An application refers to a program running on a host that can communicate with another program over a network.

However, communication technologies are highly fragmented. Different protocols and languages coexist, often creating compatibility challenges. Many field-level and factory-level technologies do not natively support Internet-based communication. A practical solution to this fragmentation is the adoption of service oriented architectures.

2.1.1 Taxonomy

The application layer protocols can be classified as:

- *Client and server* (HTTP and CoAP): a client requests data from a server.
- *Publish and subscribe* (MQTT): devices publish data to a central broker, which then distributes updates to subscribed clients.

2.2 Hyper Text Transfer Protocol

Web pages are composed of multiple resources which are uniquely identified by a Uniform Resource Locator (URL) or a Uniform Resource Identifier (URI). When a client requests a resource using a URL, the request is resolved through Domain Name Resolution (DNS), followed by the retrieval of the requested data from the corresponding web server.

On the web, resources are managed by servers, identified through URIs, and accessed synchronously by clients using a request and response paradigm. This model aligns with the REpresentational State Transfer (REST) architectural style.

2.2.1 Communication

HTTP follows a client-server architecture, where communication occurs as follows:

- The client sends an HTTP request for a specific resource, identified by its URI.

- The server processes the request and responds with the requested resource or an appropriate status message.

HTTP is a stateless protocol, meaning each request is independent, and the server does not retain memory of previous interactions.

Requests HTTP requests are ASCII-encoded and follow a format that includes various fields.

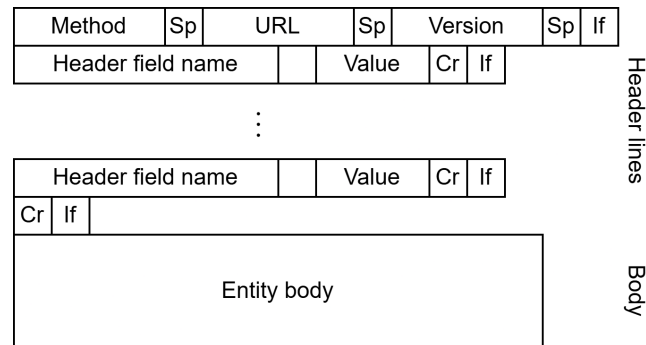


Figure 2.1: HTTP request

Responses HTTP responses follow a similar format, including a status code that indicates the outcome of the request:

- *Informational* (1xx): request received and processing continues.
- *Success* (2xx): the request was successfully processed.
- *Redirection* (3xx): further action is needed to complete the request.
- *Client-side error* (4xx): the request contains incorrect syntax or cannot be fulfilled.
- *Server-side error* (5xx): the server encountered an issue while processing the request.

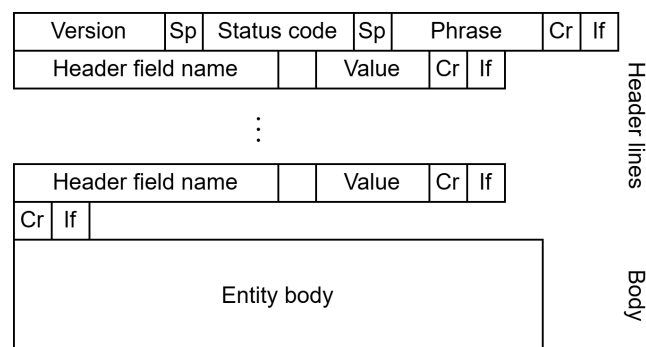


Figure 2.2: HTTP response

2.3 Constrained Application Protocol

Constrained Application Protocol (CoAP) is designed to enable web-based services in constrained wireless networks. These constraints include: low-power devices, limited memory

and processing capabilities, and energy-efficient, low-bandwidth networks. To overcome these challenges, CoAP redesigns web-based communication to suit constrained environments while maintaining compatibility with standard web technologies.

CoAP is an embedded web transfer protocol optimized for lightweight, efficient communication.

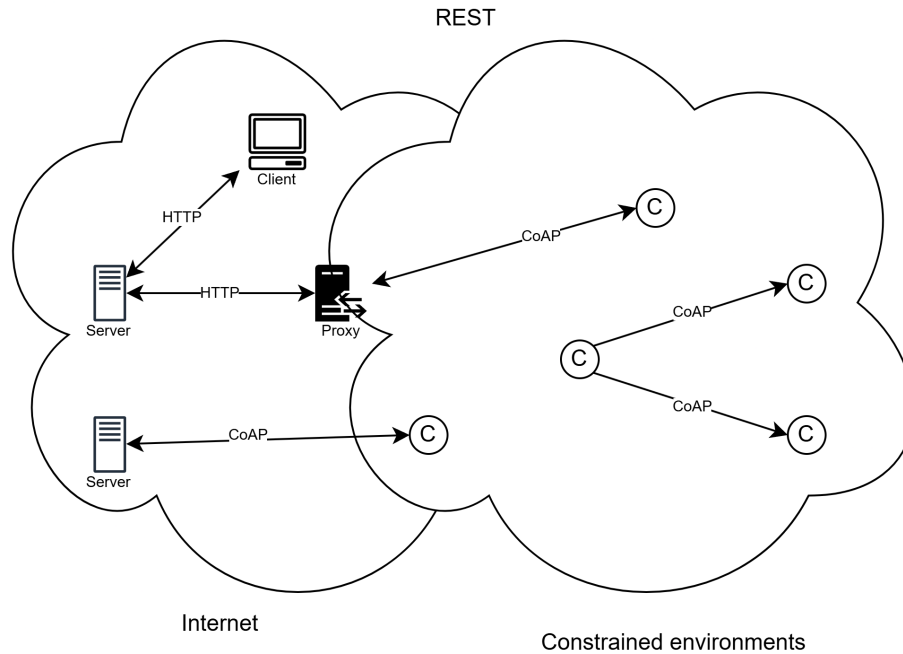


Figure 2.3: CoAP architecture

CoAP plays a critical role in IoT ecosystems, enabling lightweight, low-latency communication between resource-constrained devices while maintaining compatibility with traditional web technologies. Its UDP-based nature makes it well-suited for low-power wireless networks.

Proxy CoAP proxies act as intermediaries between clients and servers, forwarding requests and responses. They help with load balancing, security, and NAT traversal, improving network efficiency and scalability.

Caching Caching: CoAP supports caching responses to reduce repeated network requests. Clients and proxies can store responses and reuse them if the resource hasn't changed, saving bandwidth and improving speed.

2.3.1 Packets

CoAP primarily relies on UDP for lightweight communication, making it ideal for constrained networks. Message exchange occurs between endpoints with a compact 4-byte header, ensuring minimal overhead. Each message contains a 16-bit Message ID, enabling mechanisms for both reliable and unreliable transmission:

- *Confirmable messages*: require acknowledgment (**ACK**) or may trigger a reset message (**RST**) if lost. A stop-and-wait retransmission strategy with exponential backoff ensures reliability.

- *Non-confirmable messages*: sent without requiring acknowledgment, suitable for non-critical or frequent updates.
- *Duplicate detection*: both confirmable and non-confirmable messages use message IDs to identify and discard duplicates.

CoAP follows a RESTful architecture, embedding request and response semantics directly into its messages.

Ver	T	Tkl	Code	Message ID
Token				
Options				
11111111			Payload	

Figure 2.4: CoAP message

The CoAP 4-byte message header consists of several key fields such as version, message type, token length, code, and message ID.

Loss CoAP employs a stop-and-wait retransmission mechanism to handle packet loss in unreliable networks. This approach ensures that lost messages are resent while maintaining low overhead. When sending a confirmable message, the sender expects an acknowledgment or a reset message. If no response is received within a timeout period, the request is retransmitted. If no response is received after the initial timeout, and the transmission counter is less than a threshold the transmission counter is increased, the timeout value is doubled (exponential backoff), and the message is retransmitted. This process continues until one of the following conditions is met:

- An **ACK** is received, confirming successful delivery.
- An **RST** message is received, indicating rejection.
- The transmission counter exceeds **MAX_RETRANSMIT** (4 retransmissions).
- The total attempt duration surpasses **MAX_TRANSMIT_WAIT** (93 seconds), at which point the transmission is aborted.

2.3.2 Observation

In Wireless Sensor Network data is often periodic or triggered by specific events. CoAP introduces the concept of observation to address this. Observation allows clients to subscribe to resources and receive updates automatically whenever the resource's state changes.

CoAP clients use the **observe** option in their requests to subscribe to a resource. When the resource state changes, the server sends an update to the client. The client is notified asynchronously, receiving updates as they occur without needing to send explicit requests.

2.3.3 Block transfer

In networks using IPv6, large payloads are often fragmented at the lower layers. CoAP uses block transfer at the application layer to handle large messages by splitting them into smaller blocks, avoiding fragmentation at lower layers.

2.3.4 Resource discovery

CoAP supports resource discovery, allowing clients to discover and interact with available resources on CoAP servers. The goal is to enable clients to discover the links hosted by a CoAP server. The server returns resource details in a link-header format, which includes:

- *URL*: the resource location.
- *Relation*: describes the relationship between resources.
- *Type*: specifies the resource type.
- *Content type*: indicates the format of the resource.
- *Interface*: describes the access protocol.

2.4 Message Queuing Telemetry Transport

Message Queuing Telemetry Transport (MQTT) is a client-server, publish and subscribe messaging protocol designed for lightweight and efficient communication. In MQTT, clients do not communicate directly with each other. Instead, they publish messages to topics, and other clients subscribe to those topics. A single client can publish a message, and all clients subscribed to that topic will receive the message. Unlike CoAP's pull-based request and response model, MQTT uses a push model where updates are automatically sent to clients when new information is available.

Connection In MQTT, each client establishes a single connection (which supports push capabilities) to the MQTT broker for communication. MQTT can also work even through firewalls or NAT devices, thanks to its use of TCP and well-defined protocols. When a client connects to the broker, it sends a **CONNECT** message with several fields. When the broker responds to the connect message, it sends a **CONNACK** message which indicates whether the client's session was retained and the result of the connection attempt.

2.4.1 Publisher

In MQTT, the **PUBLISH** message is used by the publisher to send data to the broker. The Quality of Service (QoS) of a publisher can be:

- *QoS 0*: the message is delivered at most once, with no guarantee of delivery.
- *QoS 1*: the message is guaranteed to be delivered, but it may be delivered multiple times. The client stores the message and retransmits it until the broker acknowledges receipt. Once the broker receives the publish message, it sends a **PUBACK** message back to the client to acknowledge receipt.
- *QoS 2*: ensures that the message is delivered exactly once, and involves a 4-step handshake:
 1. *Publish reception* (broker): the MQTT broker processes the packet and sends a **PUBREC** message back.

2. *Reception* (client): upon receiving the PUBREC message, the client discards the original packet and sends a PUBREL message to the broker.
3. *Acknowledgment* (broker): the broker clears any current state and sends a PUBCOMP message to confirm the delivery of the message.
4. *Completion* (broker): the client receives the PUBCOMP message.

2.4.2 Subscriber

In MQTT, the SUBSCRIBE message is used by the client to subscribe to one or more topics. Once the broker receives the SUBSCRIBE message, it responds with a SUBACK message.

To unsubscribe from topics, the client sends an UNSUBSCRIBE message. When the broker processes the UNSUBSCRIBE request, it sends an UNSUBACK message.

2.4.3 Session

In MQTT, the session management behavior depends on the type of session:

- *Non-persistent session*: when a client disconnects, all client-related information is cleared from the broker. This means that when the client reconnects, it needs to re-subscribe to topics and will not receive messages sent during its disconnection.
- *Persistent session*: the client and the broker maintain the session state even if the client disconnects. This ensures that the client can resume its communication without losing its subscription information or missing messages. With a persistent session, even if the client is offline, the broker will hold the state and deliver any relevant messages when the client reconnects.

2.4.4 Messages

In MQTT, publishing and subscribing are asynchronous operations. Retained messages are PUBLISH messages where the **retained** flag is set to one. These messages are stored by the broker, and whenever a new client subscribes to the topic, the broker immediately sends the last retained message on that topic to the subscribing client. This ensures that subscribers receive the latest available message, even if no new message has been published.

Last will The Last Will and Testament (LWT) message in MQTT allows a client to notify other clients about an unexpected disconnect. When a client connects to the broker, it can specify a LWT message. The broker stores the LWT message and only sends it if it detects that the client has disconnected unexpectedly. When the client experiences a hard disconnection, the broker sends the stored LWT message to all subscribed clients on the specified topic. If the client disconnects gracefully, the broker will discard the LWT message.

Keep alive It is the client's responsibility to maintain an active MQTT connection. If no other interactions occur within the specified keep-alive interval, the client sends a ping request to the broker. The broker, upon receiving the ping, responds with a ping response to confirm the connection is still active. This mechanism ensures the connection remains stable and prevents unnecessary disconnections due to inactivity.

2.4.5 MQTT for Sensor Networks

MQTT for Sensor Networks (MQTT-SN) is a variation of the standard MQTT protocol designed to better suit the constraints of sensor networks and environments with limited resources. Here are the main differences compared to the standard MQTT protocol:

- Extended architecture with gateways and forwarders.
- New gateway discovery procedures and messages.
- Some messages are more.
- Extended keep alive procedures to support sleeping clients.

CHAPTER 3

Physical layer

3.1 Introduction

Fieldbus technology is primarily designed to interconnect control devices with field devices. It typically handles short messages exchanged frequently between these devices. One of the key characteristics of fieldbus systems is their strict requirements for low latency and deterministic behavior, which ensures that communication happens within a predictable timeframe.

A significant challenge in such systems arises when multiple devices need to share a single communication medium. There are two main strategies for managing this communication:

1. *Scheduled access*: this approach is akin to telling each device when it can transmit, ensuring that transmissions happen in a specific, sequential order, eliminating the risk of collisions. Centralized scheduling and polling schemes are commonly employed in this scenario, where a central authority determines the timing for each device's transmission.
2. *Random access*: this method allows devices to decide when to transmit, but they must be smart about avoiding collisions. While transmissions are less coordinated, leading to the potential for overlap, conflicts can be resolved through distributed procedures like introducing random delays before retransmissions, which helps in managing the collisions.

Scheduled access offers guaranteed performance, providing predictable delays and throughput. However, it requires coordination, such as through a central node or synchronization, which adds complexity to the system. On the other hand, random access is simpler to implement, as devices can opportunistically access communication resources. While this flexibility is advantageous, the downside is that it only offers statistical guarantees on performance. Additionally, under heavy traffic conditions, the performance can degrade due to collisions.

CAN bus The CAN bus operates with a shared physical bus, where every device connected to the bus receives all messages transmitted. This approach utilizes Carrier Sensing Multiple Access, which means that devices first sense the bus to determine if it is free. If the bus is clear, they proceed with their transmission; otherwise, they wait and try again later.

Bus arbitration To manage transmission priorities, each message on the bus has an associated priority. The priority is determined by the identifier field, with lower identifier values

signifying higher priority. Stations continuously monitor both their own transmissions and the bus's status. If a station hears a higher-priority transmission while it is transmitting, it will immediately cease transmission and allow the higher-priority message to proceed.

3.1.1 Ethernet

Initially, in 1976, Ethernet was based on a physical shared bus using coaxial cables. Between 1990 and 2000, Ethernet transitioned to star-like topologies with hubs and repeaters, using twisted pair cables. From 2000 to the present, Ethernet has moved towards fully switched and full-duplex topologies, supporting higher speeds.

Ethernet uses a variety of techniques to manage traffic and ensure efficient communication, including multiplexing, frame filtering, and multiple access protocols.

Ethernet frame An Ethernet frame consists of several components. The frame starts with a synchronization preamble, followed by a frame delimiter that marks the start of the frame.




Figure 3.1: Ethernet frame

The frame includes the following key fields:

- *48-bit addresses*: these uniquely identify the source and destination devices.
- *Type field*: this is a multiplexing field, used to indicate the type of data encapsulated in the frame.
- *Data field*: this contains the actual data being transmitted.
- *Frame check sequence*: a cyclic redundancy check used for error detection to ensure the integrity of the frame.

MAC address The MAC address is crucial for filtering and routing Ethernet frames. The first 3 bytes of the MAC address identify the manufacturer of the network interface card, while the last 3 bytes are unique to each interface. An all-ones address is used for broadcast communication, ensuring the message reaches all devices on the network.

CSMA/CD Ethernet networks originally used CSMA/CD as a method for managing access to the shared communication medium. In CSMA/CD, devices listen to the bus before transmitting to check if it is free. If the bus is clear, they proceed with transmission. If a collision occurs, all transmissions are aborted, and devices wait a random amount of time before attempting to retransmit. This collision detection process ensures that communication is efficient, although it can introduce delays.

Industry 4.0 ethernet In the context of Industry 4.0, Ethernet has been adapted for various use cases with different time-critical requirements. Ethernet in Industry 4.0 is classified into three classes based on the Maximum Cycle Time, which is a measure of the maximum time it takes to complete a cycle of communication:

- *Class A*: suitable for less time-sensitive industrial applications.

- *Class B*: designed for applications that require faster communication.
- *Class C*: this class supports the most demanding applications where extremely low-latency communication is essential.

Hub and switches Ethernet networks traditionally used hubs, where all devices connected to the hub shared the same bandwidth, leading to potential collisions. In a hub-based setup, each device listens to the shared medium and can transmit only when it senses that the bus is free.

Fully switched LANs use network switches, which effectively eliminate collisions by providing a dedicated path between devices. In a switched Ethernet network, each device communicates directly with the switch, and each transmission is isolated, meaning that multiple devices can transmit simultaneously without interference. This removes the need for CSMA/CD, as there are no shared communication channels to compete for.

3.1.2 Wireless

The evolution of wireless networks has been largely driven by the rapid development of smart objects, which are interconnected devices capable of communicating over networks. These smart objects, including everything from sensors to consumer electronics, have become a central part of modern communication infrastructures, especially with the rise of the IoT. Some possible wireless technologies are:

- *Mobile radio networks*: evolving Radio Access Networks and Core Networks to meet growing data demands.
- *Cellular IoT operators*: focused on low-power, wide-area connectivity for IoT devices.
- *Low power long range technology*: enables long-distance communication with minimal power consumption, ideal for remote IoT applications.
- *Capillary multi-hop networks*: distributed communication for short/medium-range with backhauling, extending coverage over multiple hops.

3.2 Low Power Wide Area Network

LoRaWAN is a low-power, wide-area network protocol designed for IoT applications. It operates on a cellular-like architecture but is association-less, meaning devices do not need to establish permanent connections with the network. The main components of LoRaWAN include end devices, gateways, and a network server.

End devices are typically the field devices that generate data. Gateways serve as intermediaries, receiving messages from the end devices and forwarding them to the network server. The network server is the central hub where most of the intelligence of the system resides. It is responsible for removing duplicate messages, managing acknowledgments, adjusting radio link parameters, and other essential tasks that ensure efficient communication.

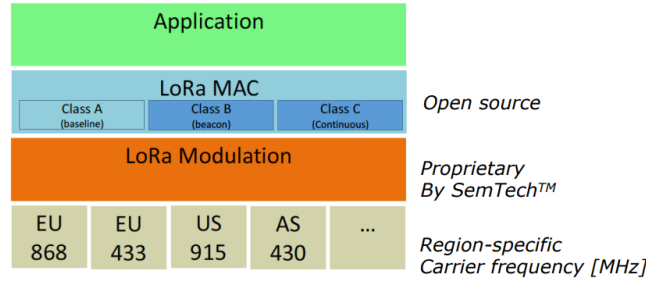


Figure 3.2: LoraWAN protocol stack

Modulation The LoRa™ modulation scheme is based on a proprietary chirp-based spread spectrum technique. This method involves modulating a signal onto a chirp signal at a higher data rate. The chip rate (in chips per second) and nominal bit rate (in bits per second) are related by the formulas:

$$R_C = BW \quad R_b = SF \frac{4BW}{2^{SF}(4 + CR)}$$

Here, BW is the reference bandwidth, SF is the spreading factor, and CR is the coding rate.

LoRaWAN balances two primary objectives: rate and reliability. The rate refers to how fast data can be transmitted, and reliability refers to how secure and stable the transmission is. For a given emitted power level, increasing the spreading factor improves sensitivity and reliability, while reducing SF increases the rate. In general:

- Lower spreading factors (SF) allow higher data rates.
- Higher spreading factors (SF) enhance reliability.

The minimum power level required for a given signal sensitivity is calculated by the formula:

$$P_{\min} = -174 + 10 \log_{10} BW + NF + SNR$$

Here, BW is the reference bandwidth, NF is the noise figure, and SNR is the signal-to-noise ratio.

End devices LoRaWAN defines three classes of end devices, each with different characteristics for handling downlink communications:

- *Class A*: this class allows for downlink messages to be received in slots that occur after the uplink transmission. It uses an ALOHA-like access method for uplink communication, where devices randomly transmit messages when required.
- *Class B*: in addition to the standard downlink slots, Class B devices have extra downlink slots that are scheduled at specific times, providing more predictable communication.
- *Class C*: almost continuously listen for downlink messages, making them the most responsive but also the highest in power consumption.

LoraWAN A In the European Union, the timing for receive windows in LoRaWAN Class A devices is standardized. If a preamble is detected during one of these receive windows, the radio receiver remains active until the downlink frame is demodulated. If a frame is detected and successfully demodulated during the first receive window, and the frame is intended for the end device (after address and Message Integrity Code checks), the device will not open the second receive window. This method ensures energy efficiency, as the device only stays awake when necessary to receive messages, and it does not waste energy on unnecessary communication attempts.

3.2.1 Messages

LoRaWAN messages are structured in multiple layers that serve distinct functions for communication and integrity verification.




Figure 3.3: LoraWAN messages

In particular, we have:

- *PHY Message structure:* begins with the Preamble, a sequence of bits used to synchronize the transmitter and receiver. Following the preamble, the Physical Header contains essential information regarding the message, such as its type and format. The PHDR_CRC and CRC fields are used to perform integrity checks, ensuring that the message hasn't been corrupted during transmission. The PHY Payload contains the actual data being sent in the message, or the payload, which varies depending on the type of communication.
- *MAC Message Header:* specifies certain parameters about the message itself. The MType field indicates the type of message being sent, distinguishing between various message types such as join requests or data uplinks. The RFU (Reserved for Future Use) field is there to ensure future compatibility as the protocol evolves. The Major field tells the receiver which message format is being used, making it possible for devices to communicate across different versions of the protocol.
- *MAC message payload:* follows the header and is where the end device-related information is stored. This section includes the DevAddr, a unique identifier assigned to the end device, ensuring that the message is directed to the correct device. The FCtrl, or frame control field, is used to manage important features like adaptive data rate (ADR), acknowledgments (ACK), and MAC commands. The FCnt is a counter that increments with each transmitted frame, ensuring proper sequencing of messages. The FOpts field contains additional frame options, often used for sending MAC commands that control device behavior or optimize network resources.

- *Frame Control Field*: critical part of the message, managing the transmission characteristics of the data. For example, the ADR (Adaptive Data Rate) function is used to optimize both the uplink and downlink data rates based on current network conditions, while the ADRAckReq flag requests an acknowledgment for ADR-related operations. The ACK field acknowledges the successful reception of a message, confirming reliable communication. The FPending field signals whether there are additional frames waiting to be sent to the device in a downlink message. Finally, the FOptsLen indicates the size of the FOpts field, allowing the receiver to interpret the data correctly.

3.2.2 LoraWAN uplink

LoRaWAN employs an ALOHA-like procedure to handle channel access and retransmissions. In this system, when a confirmed message is sent, if no acknowledgment (ACK) is received, the message is retransmitted after a random timeout. The timeout, referred to as ACK_TIMEOUT, is randomly chosen from a time interval between 1 and 3 seconds and begins once Receive_window2 concludes.

The ALOHA protocol does not require channel feedback; instead, it relies on the ACK to confirm receipt of the transmitted message. Time flows continuously within this protocol. When a message is ready for transmission, the first packet in the transmission queue is sent immediately. If an acknowledgment is not received, the message is retransmitted after a random delay, which corresponds to a randomly determined number of time slots.

In terms of performance, several assumptions are made for modeling the system. The assumption of stationarity means that the incoming and outgoing traffic rates, denoted as S_{in} and S_{out} , are equal. Traffic is distributed according to a Poisson process, with packet arrivals occurring randomly and independently over time. The packet arrival rate is denoted by λ , and the transmission duration is represented by T . This leads to the definition of G , which represents the total traffic on the channel and is calculated as $G = \lambda T$.

The probability of successful transmission \Pr_s is the likelihood that no other packet begins transmission during the conflict period of $2T$. This can be expressed as:

$$\Pr_s = \Pr(N(t - T, t + T) = 0) = e^{-2G}$$

The throughput of the system, which refers to the successful transmissions per unit of time, is given by:

$$S = GP_s = Ge^{-2G}$$

This equation highlights the relationship between traffic load G and the resulting throughput, showing that there is an optimal load for maximum throughput.

3.3 Narrowband Internet of Things

NB-IoT is a specialized technology designed to support the growing demand for efficient and scalable IoT deployments. It meets several critical requirements for IoT applications, including optimized signaling overhead, enhanced security based on 3GPP standards, extended battery life, and the ability to handle both IP and non-IP data, as well as SMS messages. Operating within licensed spectrum, NB-IoT benefits from a robust and interference-free connection, making it an ideal solution for various IoT use cases. Additionally, it can leverage existing mobile network infrastructure, simplifying integration into current systems.

Power saving mode A key feature of NB-IoT is its Power Saving Mode which is designed to reduce energy consumption even further than traditional idle modes. This mode allows devices to remain registered with the network without needing to re-establish packet data network connections, thus avoiding energy-intensive re-attachments. Devices in Power Saving Mode are not immediately reachable for mobile-terminated services, which helps conserve power without sacrificing network connectivity.

Random access procedure The random access procedure defines the process through which a device can initially connect to the network, ensuring an efficient use of available resources. Alongside this, the mechanism for message repetition ensures that critical messages are reliably delivered to devices, even in poor coverage areas, with the possibility of repeating messages up to 2048 times to maintain signal integrity.

Coverage enhancements Coverage in NB-IoT is adaptable to varying radio conditions through the use of Coverage Enhancement levels. In areas of good coverage, a device operates at CE-Level 0, requiring minimal signal boosting. As coverage weakens, the network can boost the signal, moving to CE-Level 1 and then CE-Level 2 for areas with poor signal reception, where the network increases the number of repetitions of downlink messages and uplink preambles to ensure reliable communication. This enhanced coverage capability significantly improves the network's reach, even in challenging environments, by increasing the maximum coupling loss from 144 dB to 164 dB.

Cost To reduce implementation costs, NB-IoT simplifies LTE features, focusing on lower-complexity user equipment and network configurations. For instance, NB-IoT devices use significantly smaller transport block sizes for both downlink and uplink communication. They support only single-stream transmissions and require only a single antenna. The reduced complexity also means that there is no need for turbo decoders in the UE, as only Turbo Block Code is used for downlink channels. Additionally, the system simplifies mobility measurements, with the device only needing to perform mobility measurements during idle mode, reducing operational overhead. This streamlined architecture also limits the system to half-duplex frequency-division duplexing operation and eliminates the need for parallel processing in the physical layer.

In-band operation In-band operation is another key feature of NB-IoT. By introducing NB-IoT carriers within existing LTE carriers, this approach allows for scalable deployments by adding more carriers as needed. However, this can introduce near-far interference with non-upgraded LTE base stations, so a network-wide deployment is recommended to avoid performance issues. While this in-band deployment may reduce the capacity and maximum speed of the LTE carrier, solutions such as boosting the NB-IoT output power or dynamically sharing baseband capacity with LTE can mitigate these effects.

Enhanced discontinuous reception One of the standout features for power-saving in NB-IoT is Enhanced Discontinuous Reception. This feature introduces longer sleep cycles, which drastically reduce the frequency at which devices must remain awake to monitor the network. Additionally, the eDRX feature enables a long paging cycle of up to 2.92 hours, which can be negotiated between the UE and the network, further optimizing battery life for devices that need to stay connected without frequent wake-ups.

3.3.1 Summary

NB-IoT represents a natural evolution of LTE networks, providing an efficient solution for cellular IoT. Notable improvements in Release 12 significantly reduced modem complexity, offering a 50% reduction compared to Category 1 devices and ensuring over 10 years of battery life for delay-tolerant traffic. Release 13 further optimized the system with a 75% reduction in modem complexity, a 10+ year battery life for other IoT use cases, and a 15-20 dB coverage enhancement. While narrower LTE system bandwidths, such as 200 kHz, could be introduced for further optimization, this requires substantial efforts compared to the simpler enhancements already provided.

As the world moves toward 5G, the evolution of NB-IoT continues. 5G promises to unlock even more industrial IoT applications through enhanced features like millimeter-wave access, massive MIMO, network densification, slicing, and mobile edge computing. These advancements will enable a wide range of applications, from immersive and virtual reality to autonomous vehicles and collaborative robotics, expanding the potential of IoT across various industries.

Network layer

4.1 Introduction

In the world of short-range communication, a variety of technologies and protocols exist, each designed to serve specific needs. While these solutions work well for small-scale applications and localized coverage (such as smart homes or industrial automation), they often struggle with interoperability.

Capillary networks Communication technologies today are highly fragmented, with different standards and protocols optimized for different use cases. While some technologies provide effective coverage for small areas, seamless integration across networks remains a challenge. This has led to the need for a new, unified communication stack that can bridge these gaps.

Taxonomy Several communication solutions are available, and they can generally be categorized based on:

- Proprietary or open standards.
- Application-specific or application-agnostic.
- IP-compliant or not IP-compliant.

Zigbee and 6LowPAN Among the many existing technologies, ZigBee and 6LowPAN are two of the most widely discussed. While they share similarities at the lower communication layers, their fundamental difference lies in their approach to IP connectivity:

- 6LowPAN extends IP to IoT devices, making it easier to integrate with existing internet infrastructure
- ZigBee, on the other hand, operates independently of IP, requiring additional gateways for internet connectivity

4.2 Zigbee

Zigbee is a widely used wireless communication technology designed for low-power, low-data-rate applications. Its main features include:

- *Low-cost hardware and software*: devices typically cost around \$2, making Zigbee an affordable option.
- *Short transmission range*: around 10 meters, which is suitable for local communication.
- *Low latency*: ensures quick response times.
- *High energy efficiency*: optimized for battery-powered devices, making it ideal for IoT applications.

History In the mid-1990s, IoT solutions were highly fragmented, with numerous proprietary protocols leading to major compatibility and interoperability challenges. This also drove up costs, making large-scale adoption difficult. To address this, IEEE launched Working Group 4 in 2001 to develop a standardized reference technology. This effort resulted in the IEEE 802.15.4 standard, which was officially published in March 2003.

Zigbee communication stack is composed as follows.




Figure 4.1: Zigbee communication stack

Devices Zigbee networks consist of different types of devices, each playing a specific role:

- *Full function device*: can send beacons to manage network communication, communicate with other FFDs, route network frames, and act as a PAN coordinator. Typically powered by an external power supply.
- *Reduced function device*: cannot route frames or communicate with other RFDs. Communicates only with an FFD. Designed for low-power, battery-operated applications.
- *PAN coordinator*: handles device association and de-association.

Network topologies Zigbee networks support three different topologies, each suited for specific use cases:

- *Star topology*: all devices connect to a central coordinator. Simple, but limited in scalability. Common in home automation and industrial monitoring.
- *Mesh topology*: devices communicate with multiple neighboring nodes. Highly reliable and scalable, as data can take multiple paths. Used in large-scale IoT deployments.
- *Cluster-tree topology*: an hybrid between star and mesh. Devices are organized in clusters, with a coordinator managing each cluster. Efficient for applications requiring structured hierarchies, such as industrial automation.

4.2.1 Physical layer

The physical layer (PHY) of Zigbee, based on IEEE 802.15.4, is responsible for managing radio communication and ensuring reliable data transmission. Its main functions include:

- *Radio transceiver control*: activates and deactivates the radio module as needed.
- *Energy Detection*: measures signal strength within the current channel, useful for scanning and channel selection.
- *Link Quality Indicator*: evaluates the quality of received packets.
- *Clear Channel Assessment*: determines if a channel is free before transmitting, enabling Carrier Sense Multiple Access with Collision Avoidance.
- *Channel frequency selection*: supports multiple frequency bands to reduce interference.
- *Data transmission and reception*: handles the actual exchange of data between devices

Zigbee operates across multiple frequency bands, offering flexibility based on regional regulations and application requirements: 868 MHz, 915 MHz, and 2.4 GHz.

Packet The physical layer defines the structure of transmitted data packets, ensuring synchronization and proper interpretation of information. The key components of a PDU (Packet Data Unit) are:

- *Preamble*: helps receivers synchronize with the incoming signal.
- *Start frame delimiter* (SFD): marks the beginning of a valid frame.
- *Frame length*: specifies the size of the PHY payload (in octets). For MAC data frames, the payload length ranges between 9 and 127 octets.

4.2.2 Medium Access Control layer

The MAC sublayer in IEEE 802.15.4 plays a crucial role in managing communication between devices in a Zigbee network. It provides key functionalities such as:

- *Beacon management*: synchronizes devices and schedules transmissions.

- *Channel access management*: determines how devices share the communication medium.
- *Guaranteed Time Slot management*: allocates dedicated transmission slots for time-sensitive data.
- *Frame validation*: ensures data integrity and discards invalid frames.
- *Acknowledged frame delivery*: confirms successful message reception.
- *Association and disassociation*: manages device connections within a network.
- *Security hooks*: provides mechanisms for implementing encryption and authentication.

Zigbee defines two operation modes to balance power efficiency and communication reliability:

1. *Beacon-enabled mode*: the PAN Coordinator periodically transmits beacons to synchronize devices. Commonly used in star topologies. Uses a combination of slotted CSMA-CA and scheduled transmissions to manage channel access.
2. *Non beacon-enabled mode*: devices communicate using unslotted CSMA-CA, meaning no predefined transmission schedule. Provides uncoordinated access, reducing overhead but increasing the risk of collisions.

Channel access mechanism Since multiple devices share the same frequency, Zigbee uses a mix of scheduled and random access to prevent interference:

- *Scheduled Access*: managed by the PAN Coordinator (only in beacon-enabled mode).
- *Random Access*: used for communication between RFDs or between RFD/FFD and the PAN Coordinator. Allowed in both operation modes.

CSMA-CA Zigbee devices rely on CSMA/CA to reduce transmission conflicts. Each device tracks key variables to manage communication: the number of backoffs, which counts transmission delays; the contention window, controlling how long the channel must remain clear before transmission; and the backoff exponent, determining the wait time before attempting access. The CSMA/CA process begins with the transmitting node waiting for a random backoff period. If the channel remains clear for the required CW duration, the device proceeds with transmission and awaits an acknowledgment. However, if the channel is busy, the node increases both BE and NB, repeating the process until either a successful transmission occurs or the retry limit is reached, at which point the packet is discarded. Additional factors influence this process. All transmissions, including ACKs, must complete within the Contention Access Period. While slotted CSMA/CA requires synchronization, unslotted CSMA/CA operates without it. Certain packets, such as beacons and ACKs, bypass CSMA entirely and transmit immediately. If a collision occurs—evident when no ACK is received—the device restarts the procedure.

Network formation To join a network, Zigbee devices scan for available PANs. Active scanning, used by full-function devices, involves sending a beacon request, prompting nearby PAN coordinators to respond with beacons. After scanning, the device selects a PAN ID and joins. Passive scanning, available to both FFDs and reduced-function devices, operates similarly but without actively requesting beacons. Instead, the device listens for beacon transmissions from coordinators.

Extensions Enhancements to the IEEE 802.15.4 standard have expanded Zigbee’s capabilities for specialized applications. IEEE 802.15.4e introduces slotted channel access for improved efficiency. IEEE 802.15.4g targets smart utility applications, such as smart meters, while IEEE 802.15.4k optimizes communication for critical infrastructure monitoring, including power grids and industrial control systems.

4.2.3 Network layer

The network layer provides the following key functionalities:

- *Configuring a new device*: this involves setting up the device and configuring the stack to operate as required.
- *Starting a network*: establishing a new network from scratch.
- *Joining, rejoining, and leaving a network*: devices can join, rejoin, or leave a network. ZigBee coordinators or routers can also request devices to leave the network.
- *Addressing*: ZigBee coordinators and routers have the responsibility of assigning unique addresses to devices joining the network.
- *Neighbor discovery*: the ability to discover and report information about the device’s one-hop neighbors.
- *Route discovery*: the process of finding and recording paths within the network, allowing messages to be routed efficiently.
- *Reception control*: the device can manage when the receiver is activated and for how long, supporting MAC synchronization and direct reception.
- *Routing*: this encompasses various routing mechanisms, such as unicast, broadcast, multicast, and many-to-one, to efficiently exchange data within the network.

In ZigBee, there are three types of devices:

- *ZB coordinator* (FFD): the central coordinator that initiates and manages the network.
- *ZB router* (FFD): routers that relay messages and allow other devices to join the network.
- *ZB end-device* (RFD or FFD): end devices, which may either be reduced-function devices (RFDs) or full-function devices (FFDs).

ZigBee routing integrates two primary protocols: ad-hoc on-demand distance vector and cluster tree algorithm.

4.2.3.1 Cluster tree algorithm

The Cluster Tree Algorithm defines how a ZigBee network organizes itself into a tree structure for efficient data routing. The most important operations are:

- *Tree initialization by FFD*: a ZigBee Full-Function Device (FFD) begins the tree formation by scanning available channels and selecting one with minimal interference. It sets its PAN identifier and assigns itself the network address 0, indicating its role as the coordinator.

- *Association of devices*: other devices can then join the network by associating with the coordinator through standard association procedures. These devices may include routers (only FFDs) and end-devices.
- *Address assignment*: during the association process, devices are assigned 16-bit short addresses. The parent device (either the PAN coordinator or a ZB router) assigns address ranges to its child devices. Each router receives a consecutive block of addresses based on its depth in the tree, while the first address in this block is used as the node's address, and the remaining addresses are available for its child devices. Address assignment follows a hierarchical structure. The size of the address range $A(d)$ assigned to a router node at depth is calculated using the following formula:

$$A(d) = \begin{cases} 1 + D_m + R_m & \text{if } d = L_m - 1 \\ 1 + D_m + R_m A(d+1) & \text{if } 0 \leq d < L_m - 1 \end{cases}$$

For each device, addresses are assigned in a range determined by its depth in the tree. Devices further from the coordinator receive a larger range of addresses for their child devices.

- *Coordinator constraints*: the ZigBee coordinator also defines key network parameters, such as the maximum number of routers (R_m) and end-devices (D_m) each router can support, as well as the maximum depth (L_m) of the tree.

Tree routing Routing within a ZigBee network follows the tree structure formed by the coordinator and its child devices. Messages are routed through the tree in the following manner:

- If the destination address belongs to one of the child end-devices, the message is routed directly to the destination.
- If the destination address corresponds to a child router, the message is sent to that router for further forwarding.
- If the destination address is not found among the children, the message is forwarded to the parent node for further routing.

This tree-based routing system is simple and works efficiently, although it may not always provide the most optimized routes. It is also quality-agnostic, meaning the routing decisions do not take into account the quality of the link.

4.2.3.2 Ad-hoc on demand distance vector

In contrast to tree routing, Ad-hoc On-demand Distance Vector (AODV) is a reactive routing protocol used when a device needs to send data to an unknown destination. Here's how AODV works:

1. *Route request (RREQ)*: when a device wants to communicate with a destination, it broadcasts a Route Request (RREQ) message.
2. *Flooding*: the RREQ message is flooded across the network, and each node that relays the request stores the address of the node that forwarded it, setting up a reverse path to the source.

3. *Route reply* (RREP): when the destination node receives the RREQ, it sends a Route Reply (RREP) message back to the source, traveling along the reverse path set up during the flooding process.

Routing tables are updated as nodes process these RREQ and RREP messages. A node's routing table contains the following entries:

- *Destination address*: the address of the destination.
- *Next-hop address*: the address of the next hop toward the destination.
- *Entry status*: whether the entry is active, in discovery, or inactive.

Additionally, a Routing Discovery Table (RDT) keeps track of RREQ messages, storing information like the source address, the sender's address, and the accumulated path cost.

Routing cost The cost of a path P is the sum of the individual costs of the links that make up the path. The cost for each link l is determined by factors like the packet reception rate p_l and follows a standard cost function, which suggests a value of 7 for a link with poor reception and a function based on the reception rate for links with better performance:

$$C(l) = \begin{cases} 7 \\ \min \left(7, \text{round} \left(\frac{1}{p_l^4} \right) \right) \end{cases}$$

Application profiles ZigBee also provides predefined application profiles that ensure interoperability across different devices and manufacturers. These profiles define the following:

- A common language for data exchange.
- A well-defined set of processing actions for devices.
- Device interoperability across manufacturers.
- Simplicity and reliability for end users.

Each profile is made up of:

- A set of devices needed for the application.
- A set of clusters that implement the required functionality.
- A set of attributes representing device state.
- A set of commands enabling communication between devices.